

# Beyond linear models

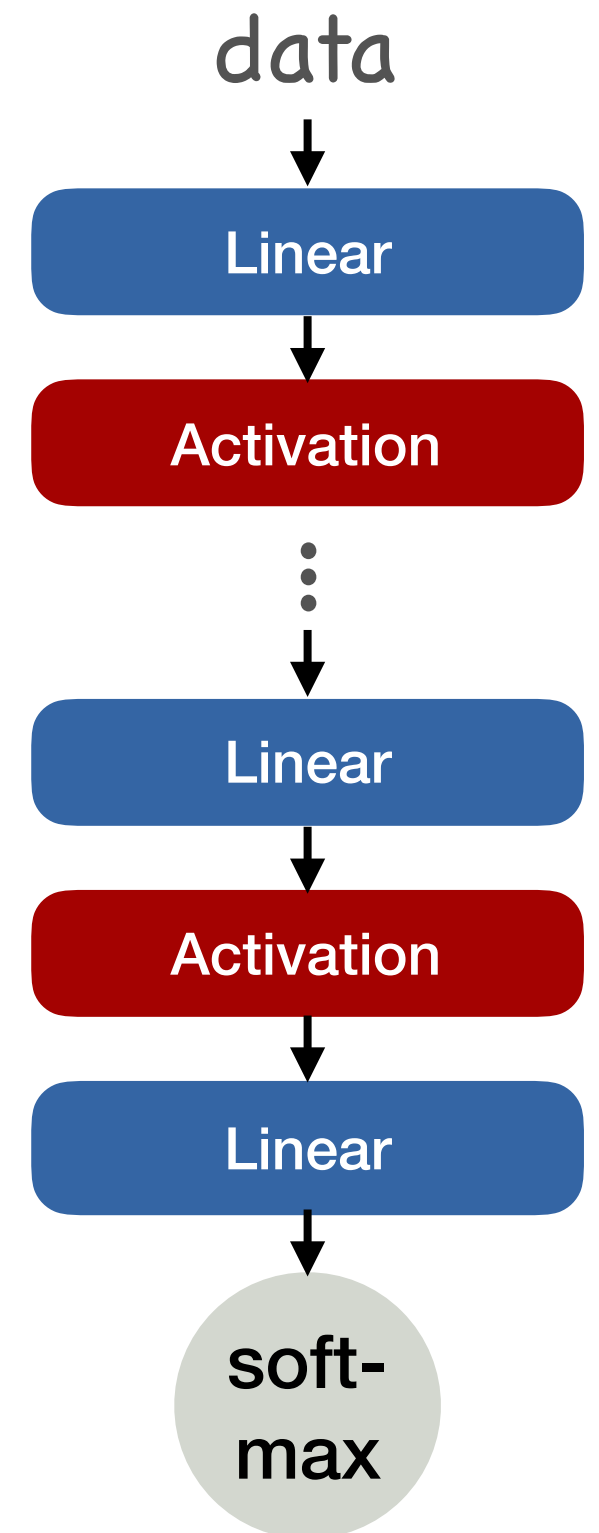
© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Linear models

- Limitations of linear models
- Linear to deep networks

# Deep networks


- Modular design
- Non-linearities and activation functions
- Outputs and losses
- Optimization




# Limitations of linear models

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Linear classifier




# Linear classifier



# A simple example

- Binary paw classification
- Dog paw or not




# A simple example






# Why does the linear model break?

- By linearity:
  - $\mathbf{w}^\top \mathbf{x}_1 + b > 0$
  - $\mathbf{w}^\top \mathbf{x}_2 + b > 0$
  - Then  $\mathbf{w}^\top \mathbf{x} + b > 0$   
for any  $\mathbf{x} = \alpha \mathbf{x}_1 + (1 - \alpha) \mathbf{x}_2$
- Cannot learn xor



# Does adding more linear layers help?

- No
- Combination of linear layers still linear
- $\mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2$   
 $= (\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$




# Non-linearities

© 2019 Philipp Krähenbühl and Chao-Yuan Wu


# Non-linearities

- Rectified Linear Unit
  - $\text{ReLU}(x) = \max(x, 0)$
- Non-linear and differentiable almost everywhere



# Deep networks

- Alternates linear and non-linear layers



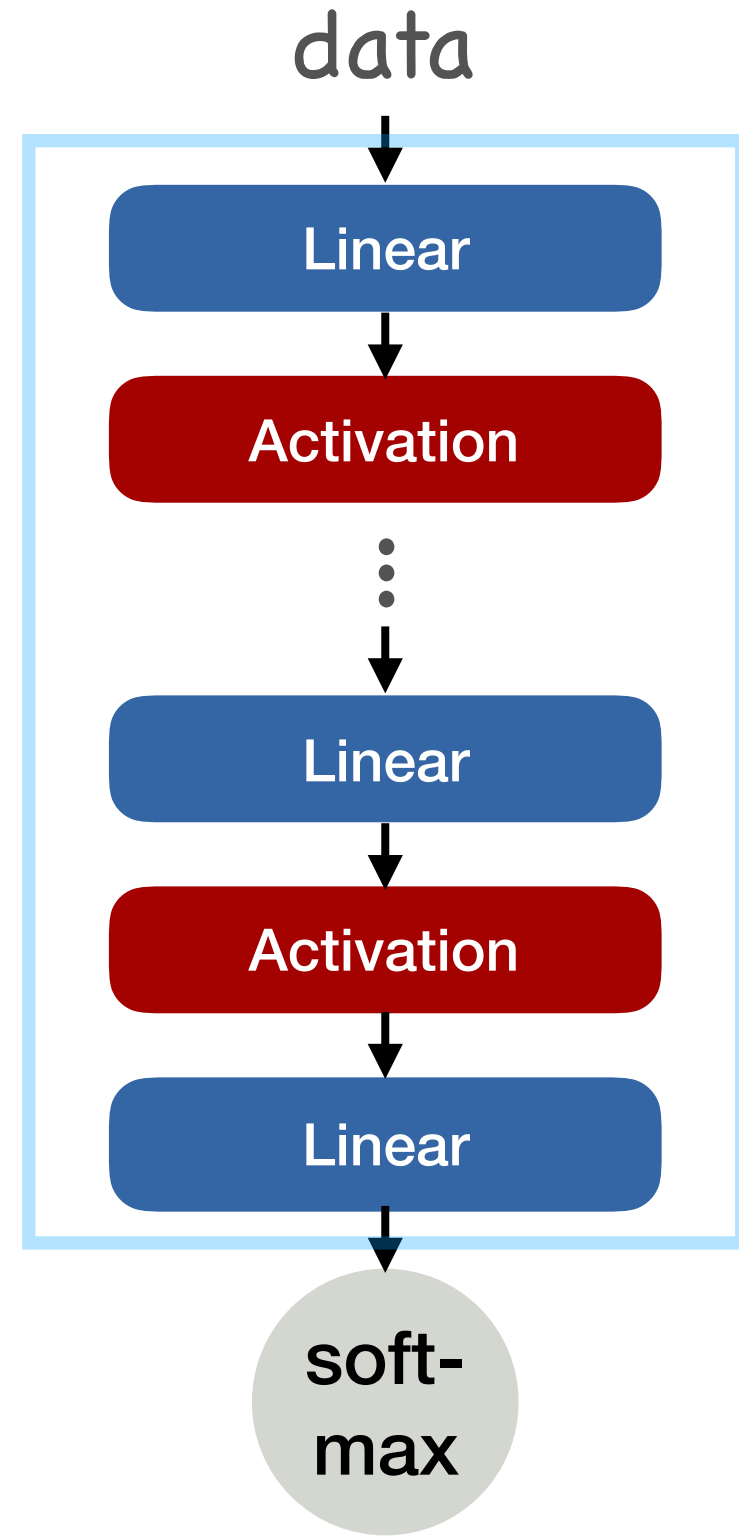
# A simple example

- “Shallow” network
- Dog paw or not?



# Deep networks

- Class of continuous functions  $f_{\theta} : \mathbf{X} \rightarrow \mathcal{O}$
- Parameters  $\theta$
- Can approximate any continuous function



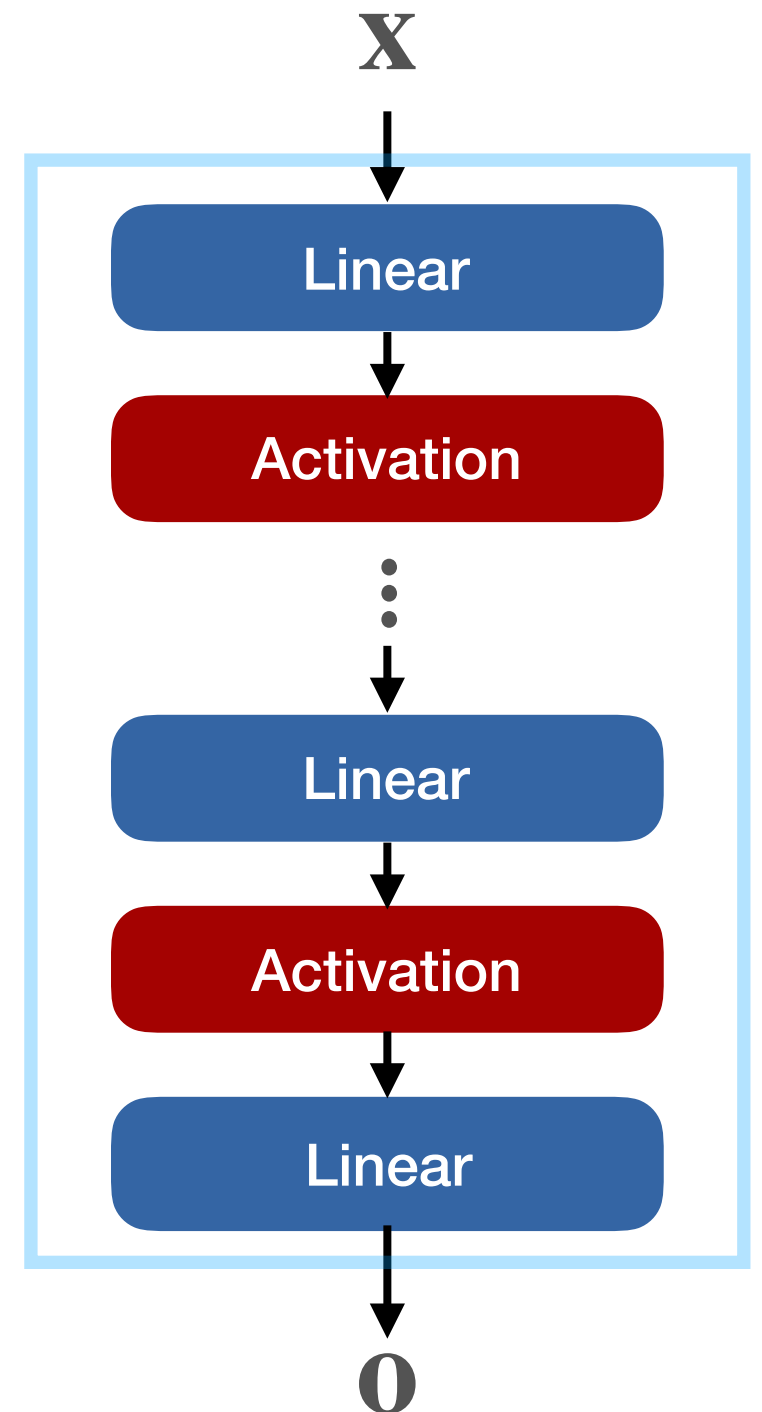
# Output representations

© 2019 Philipp Krähenbühl and Chao-Yuan Wu



# Inputs and outputs of networks

- Input:
  - Tensor  $\mathbf{x}$
- Output:
  - Tensor  $\mathbf{y}$



# Regression

- vanilla tensor  $\hat{\mathbf{y}} = \mathbf{0}$

# Positive regression

- Option 1: ReLU
  - $\hat{y} = \max(\mathbf{o}, 0)$
- Option 2: Soft ReLU
  - $\hat{y} = \log(1 + e^{\mathbf{o}})$

# Binary Classification

- Option 1: Thresholding

- $\hat{y} = \mathbf{1} > 0$

- Option 2: Logistic Regression

- $p(1) = \sigma(\mathbf{z})$

# General Classification

- Output more values, one per class
- Option 1: argmax
  - $\hat{y} = \operatorname{argmax}_i \mathbf{o}_i$
- Option 2: softmax
  - $p(y) = \operatorname{softmax}(\mathbf{o})_y$

# Output representations in practice


- Do not add into model
- Always output raw values

# Loss functions

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Inputs and outputs of networks

- Input:
  - Tensor  $\mathbf{x}$
- Output:
  - Tensor  $\mathbf{y}$





# Regression

- L1 loss

- $\ell = |\mathbf{y} - \mathbf{o}|$

- L2 loss

- $\ell = \|\mathbf{y} - \mathbf{o}\|^2$

# Classification

- Compute likelihood
  - $p(1) = \sigma(o)$
  - $\mathbf{p} = \text{softmax}(\mathbf{o})$
- Cross entropy / -Log likelihood
  - $-\log(p(y))$

# Classification losses in practice

- $\sigma(o) = 0$  for  $o \rightarrow -50$ 
  - $\log(\sigma(o)) = \log(0)$   
is undefined
- Combine  $\log$  and  $\sigma$ 
  - BCEWithLogitsLoss
  - CrossEntropyLoss

## 05


January 23, 2024

```
[1]: %pylab inline
import torch
# Making sure we can find the data loader
import sys
sys.path.append('.')
sys.path.append('../..')
from data import load
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.  
Populating the interactive namespace from numpy and matplotlib

```
[2]: # Let's load the dataset
train_data, train_label = load.get_dogs_and_cats_data(resize=(32,32),
    ↪n_images=100)
input_size = 32*32*3
to_image = load.to_image_transform()
```

```
[3]: figure(figsize=(9,9))
# Plot the first 9 images (all cats)
for i, (data, label) in enumerate(zip(train_data[:9], train_label[:9])):
    subplot(3,3,1+i)
    imshow(to_image(data))
    title('label = %d'%label)
    axis('off')
```



```
[4]: class Network1(torch.nn.Module):
      def __init__(self, n_hidden=100):
          super().__init__()
          self.linear1 = torch.nn.Linear(input_size, n_hidden)
          self.activation = torch.nn.ReLU()
          self.linear2 = torch.nn.Linear(n_hidden, 1)

      def forward(self, x):
          return self.linear2(self.activation(self.linear1(x.view(x.size(0), 1)
          ↪-1))))
```

```
[5]: # Create the network
net1 = Network1(100)
# Run an image through it
print( net1(train_data).view(-1).detach().numpy() )
```

```
[-0.11666452 -0.01218061 -0.02647986 -0.05582439 -0.06884713 -0.07426661
 -0.11988138 -0.1508754 -0.09473266 -0.10555119 -0.11205442 -0.07451358
 -0.0114538 -0.03645098 -0.06230658 -0.05653633 -0.20482528 -0.10789528
 -0.09619182 -0.07970785 -0.12600176 -0.05568714  0.00681441 -0.14392285
 -0.12056908 -0.12222722 -0.10236578 -0.24464132 -0.04859121 -0.20569515
 -0.09635501 -0.09910933 -0.19064035 -0.03423431 -0.08156552 -0.15408464
 -0.02703031 -0.14437844 -0.12504634 -0.1360234 -0.06810249 -0.0230343
 -0.20682594 -0.05603386 -0.14723194 -0.14903176 -0.04260117 -0.06173945
 -0.0771127 -0.07101672 -0.13693443 -0.16763532 -0.12506317 -0.0840079
 -0.08369774 -0.00672981  0.00092466 -0.07512563  0.00588925 -0.10326148
  0.01207101 -0.14386982 -0.09098267 -0.18191543 -0.12514006  0.01794229
 -0.17029686 -0.19610392 -0.11749545 -0.10202523 -0.11068309 -0.0528039
 -0.06975032 -0.06040952 -0.03295252 -0.06511676 -0.13189809 -0.06583469
 -0.04055259 -0.0561346 -0.02101819 -0.03989616 -0.05980999 -0.17139488
 -0.02495858 -0.12706935 -0.20838268 -0.06545367 -0.03899799 -0.07631201
 -0.11093491 -0.02930016 -0.20559758 -0.05399808 -0.06877933 -0.12470949
 -0.08101448 -0.09563463 -0.08646178 -0.12490054]
```

```
[6]: class Network2(torch.nn.Module):
    def __init__(self, *hidden_size):
        super().__init__()
        layers = []
        # Add the hidden layers
        n_in = input_size
        for n_out in hidden_size:
            layers.append(torch.nn.Linear(n_in, n_out))
            layers.append(torch.nn.ReLU())
            n_in = n_out

        # Add the classifier
        layers.append(torch.nn.Linear(n_out, 1))
        self.network = torch.nn.Sequential(*layers)

    def forward(self, x):
        return self.network(x.view(x.size(0), -1))
```

```
[7]: # Create the network
net2 = Network2(100, 50, 50)
# Run an image through it
print( net2(train_data).view(-1).detach().numpy() )
```

```
[-0.03342889 -0.0351971 -0.02820556 -0.06019168 -0.04007806 -0.05432146
 -0.04470311 -0.03901586 -0.04551048 -0.03731422 -0.05093319 -0.03082472]
```

```

-0.0259193 -0.04776499 -0.04250351 -0.029782 -0.03178665 -0.03900372
-0.04654014 -0.03617655 -0.03003443 -0.04539562 -0.02983748 -0.03075305
-0.03199217 -0.0243466 -0.054005 -0.03918456 -0.03946443 -0.05627764
-0.04166208 -0.02576187 -0.03961597 -0.04083673 -0.03986252 -0.03892076
-0.0456367 -0.03987778 -0.06198189 -0.03728481 -0.03155395 -0.036805
-0.03400616 -0.04274697 -0.04948325 -0.05085205 -0.04161504 -0.01854287
-0.04013046 -0.06409714 -0.04915904 -0.02758344 -0.04848149 -0.03329031
-0.04119885 -0.03161262 -0.03294888 -0.04454168 -0.02716818 -0.05270016
-0.01403332 -0.04991054 -0.03975805 -0.05376447 -0.04529271 -0.0293345
-0.04886403 -0.04442658 -0.0296757 -0.03936881 -0.04778569 -0.05143041
-0.03218606 -0.03332945 -0.0312345 -0.03794158 -0.04582901 -0.06189429
-0.04525457 -0.03291155 -0.03706281 -0.02100243 -0.03566442 -0.05207755
-0.04284387 -0.0400481 -0.03666504 -0.03544376 -0.04095592 -0.03590231
-0.06017258 -0.03795452 -0.05677029 -0.05759883 -0.05293075 -0.03998635
-0.02990015 -0.03219899 -0.02292 -0.02041762]


```

```

[8]: plot( net2(train_data[train_label==0]).view(-1).detach().numpy(), '*',
↪label='cats')
plot( net2(train_data[train_label==1]).view(-1).detach().numpy(), '*',
↪label='dogs')
hlines(net2(train_data).detach().numpy().mean(), 0, 50)
legend()

```

[8]: <matplotlib.legend.Legend at 0x7fc99558c8b0>



[ ]:



# Optimization of deep networks

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Data


- Input:  $\{\mathbf{x}_0, \dots, \mathbf{x}_{N-1}\}$
- Label:  $\{\mathbf{y}_0, \dots, \mathbf{y}_{N-1}\}$
- Dataset:  $D = \{(\mathbf{x}_0, \mathbf{y}_0), \dots, (\mathbf{x}_{N-1}, \mathbf{y}_{N-1})\}$



⋮

# Model

- Deep network  $f: (\mathbf{x}, \theta) \rightarrow \mathbf{o}$
- Layers of computation
- Parameters  $\theta$
- Differentiable computation graph



# Loss

- Differentiable  $\ell(\mathbf{o}, \mathbf{y})$

- Regression

- Distance norm

$$\ell(\mathbf{o}, \mathbf{y}) = \|\mathbf{o} - \mathbf{y}\|$$

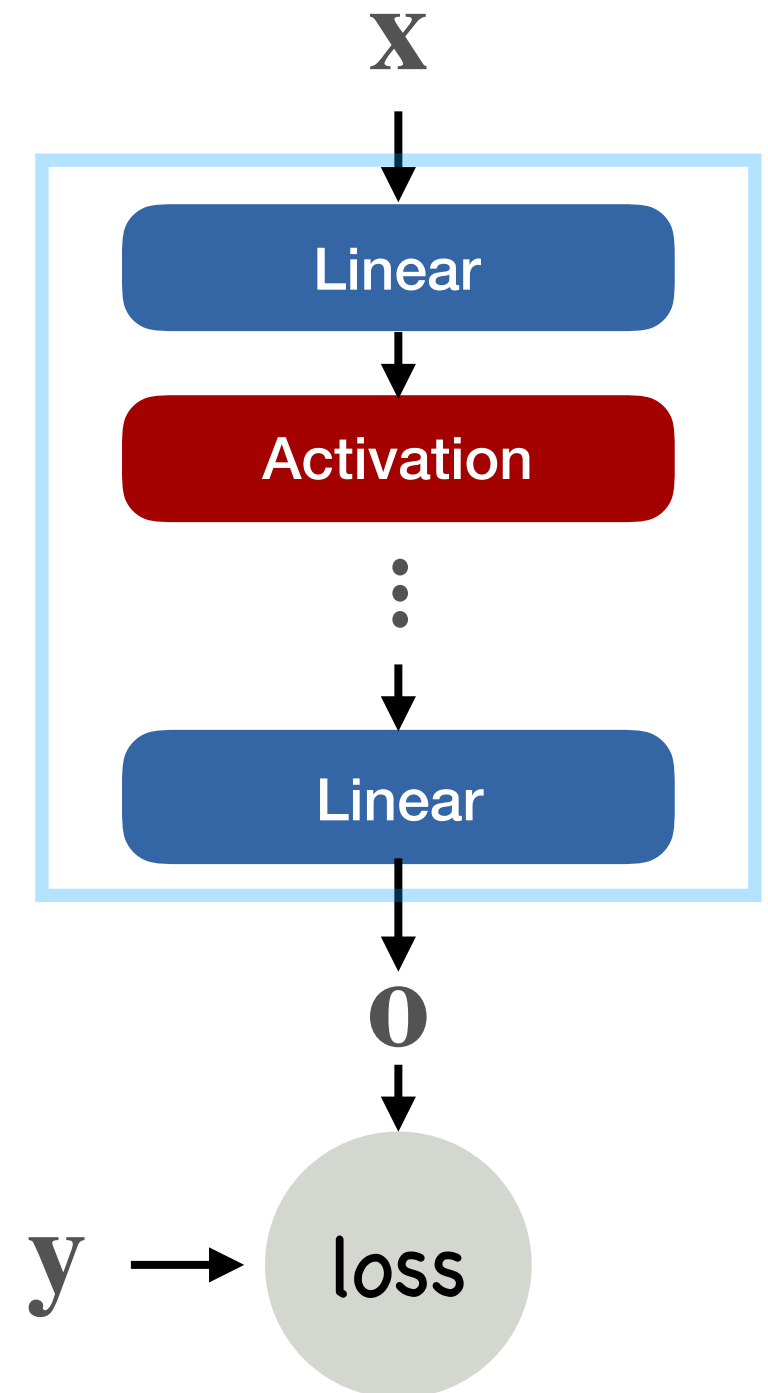
- Classification

- Cross Entropy

$$\ell(\mathbf{o}, y) = -\log p(y)$$

- Over training dataset

- $L(\theta) = \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim D}[\ell(f(\mathbf{x}, \theta), \mathbf{y})]$




# Optimization

- Minimize  $L(\theta)$

# Gradient Descent

- Repeat until convergence:

- $\theta := \theta - \epsilon \frac{dL(\theta)}{d\theta}$



# Issue with Gradient Descent

- Slow to compute gradient

- $$\frac{dL(\theta)}{d\theta} = \mathbb{E}_{\mathbf{x}, \mathbf{y} \in D} \left[ \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right]$$

# Stochastic Gradient Descent

© 2019 Philipp Krähenbühl and Chao-Yuan Wu



# Gradient Descent

- Repeat until convergence:

- $\theta := \theta - \epsilon \frac{dL(\theta)}{d\theta}$

# Gradient Descent

- Repeat until convergence:
  - $\theta_0 := \theta$
  - for  $\mathbf{x}, \mathbf{y} \sim D$  :
    - $\theta := \theta - \epsilon \frac{d\mathcal{L}(f(\mathbf{x}, \theta_0), \mathbf{y})}{d\theta_0}$

# Stochastic Gradient Descent

- Repeat until convergence:
- for  $\mathbf{x}, \mathbf{y} \sim D$  :
  - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$

# Terminology


- Repeat until convergence:
- for  $\mathbf{x}, \mathbf{y} \sim D$  :
  - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$

# Practical SGD


- For  $n$  epochs:
  - for  $\mathbf{x}, \mathbf{y} \sim D$  :
    - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$

# Learning curves

GD



SGD



# The Variance of SGD

$$\frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \neq \frac{dL(\theta)}{d\theta}$$

- Variance

- $$\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim D} \left[ \left( \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} - \frac{dL(\theta)}{d\theta} \right)^2 \right]$$
$$= \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim D} \left[ \left( \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right)^2 \right] - \left( \frac{dL(\theta)}{d\theta} \right)^2$$

# Mini-batches

© 2019 Philipp Krähenbühl and Chao-Yuan Wu



# Stochastic Gradient Descent

- For  $n$  epochs:
  - for  $\mathbf{x}, \mathbf{y} \sim D$  :
    - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$

# Stochastic Gradient Descent

- For  $n$  epochs:
  - for  $i$  in  $0, \dots, |D| - 1$ 
    - $\mathbf{x}, \mathbf{y} := D_i$
    - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$

# Mini-batches

- For  $n$  epochs:
  - Split dataset  $D$  into  $m$  mini-batches  $B_0, \dots, B_{m-1}$  of size BS
  - for each batch  $B_i$ 
    - $\theta := \theta - \epsilon \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim B_i} \left[ \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right]$

# Variance of mini-batches

- Variance of SGD

- $$\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim D} \left[ \left( \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right)^2 \right] - \left( \frac{dL(\theta)}{d\theta} \right)^2$$

- Variance of SGD with mini-batches

- $$\mathbb{E}_{B_i} \left[ \left( \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim B_i} \left[ \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right] \right)^2 \right] - \left( \frac{dL(\theta)}{d\theta} \right)^2$$

Always use mini-batches

# Variance of mini-batches

Jensen's inequality

$$\left( \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim B_i} \left[ \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right] \right)^2 \leq \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim B_i} \left[ \left( \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right)^2 \right]$$

# Momentum

- © 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Stochastic Gradient Descent

- For  $n$  epochs:
  - for  $\mathbf{x}, y \sim D$  :
    - $\theta := \theta - \epsilon \frac{d\ell(f(\mathbf{x}, \theta), y)}{d\theta}$



# Momentum


- $\mathbf{v} := 0$
- For  $n$  epochs:
  - for  $\mathbf{x}, \mathbf{y} \sim D$ 
    - $\mathbf{v} := \rho \mathbf{v} + \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta}$
    - $\theta := \theta - \epsilon \mathbf{v}$

# Variance reduction of Momentum

- Variance of SGD

$$\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim D} \left[ \left( \frac{d\ell(f(\mathbf{x}, \theta), \mathbf{y})}{d\theta} \right)^2 \right] - \left( \frac{dL(\theta)}{d\theta} \right)^2$$

# Momentum



January 23, 2024

```
[1]: %pylab inline
import torch
# Making sure we can find the data loader
import sys
sys.path.append('.')
sys.path.append('../..')
from data import load
device = torch.device('cuda') if torch.cuda.is_available() else torch.
    ↪device('cpu')
print('device = ', device)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.  
 Populating the interactive namespace from numpy and matplotlib  
 device = cuda

```
[2]: # Let's load the dataset
train_data, train_label = load.get_dogs_and_cats_data(resize=(32,32))
valid_data, valid_label = load.get_dogs_and_cats_data(split='valid',
    ↪resize=(32,32))
input_size = 32*32*3
to_image = load.to_image_transform()

train_data, train_label = train_data.to(device), train_label.to(device)
valid_data, valid_label = valid_data.to(device), valid_label.to(device)
```

```
[3]: class Network2(torch.nn.Module):
    def __init__(self, *hidden_size):
        super().__init__()
        layers = []
        # Add the hidden layers
        n_in = input_size
        for n_out in hidden_size:
            layers.append(torch.nn.Linear(n_in, n_out))
            layers.append(torch.nn.ReLU())
            n_in = n_out

        # Add the classifier
```

```

        layers.append(torch.nn.Linear(n_out, 1))
        self.network = torch.nn.Sequential(*layers)

    def forward(self, x):
        return self.network(x.view(x.size(0), -1)).view(-1)

```

```

[4]: %load_ext tensorboard
import tempfile
log_dir = tempfile.mkdtemp()
%tensorboard --logdir {log_dir} --reload_interval 1 --bind_all

```

<IPython.core.display.HTML object>

```

[5]: import torch.utils.tensorboard as tb
n_epochs = 100
batch_size = 128

train_logger = tb.SummaryWriter(log_dir+'/deepnet1/train', flush_secs=1)
valid_logger = tb.SummaryWriter(log_dir+'/deepnet1/valid', flush_secs=1)

# Create the network
net2 = Network2(100,50,50).to(device)

# Create the optimizer
optimizer = torch.optim.SGD(net2.parameters(), lr=0.01, momentum=0.9,
    ↪weight_decay=1e-4)

# Create the loss
loss = torch.nn.BCEWithLogitsLoss()

# Start training
global_step = 0
for epoch in range(n_epochs):
    # Shuffle the data
    permutation = torch.randperm(train_data.size(0))

    # Iterate
    train_accuracy = []
    for it in range(0, len(permutation)-batch_size+1, batch_size):
        batch_samples = permutation[it:it+batch_size]
        batch_data, batch_label = train_data[batch_samples],
    ↪train_label[batch_samples]

        # Compute the loss
        o = net2(batch_data)
        loss_val = loss(o, batch_label.float())

```

```

train_logger.add_scalar('train/loss', loss_val, global_step=global_step)
# Compute the accuracy
train_accuracy.extend(((o > 0).long() == batch_label).cpu().detach().
↪numpy())

optimizer.zero_grad()
loss_val.backward()
optimizer.step()

# Increase the global step
global_step += 1

# Evaluate the model
valid_pred = net2(valid_data) > 0
valid_accuracy = float((valid_pred.long() == valid_label).float().mean())

train_logger.add_scalar('train/accuracy', np.mean(train_accuracy), ↵
↪global_step=global_step)
valid_logger.add_scalar('valid/accuracy', valid_accuracy, ↵
↪global_step=global_step)

```

```

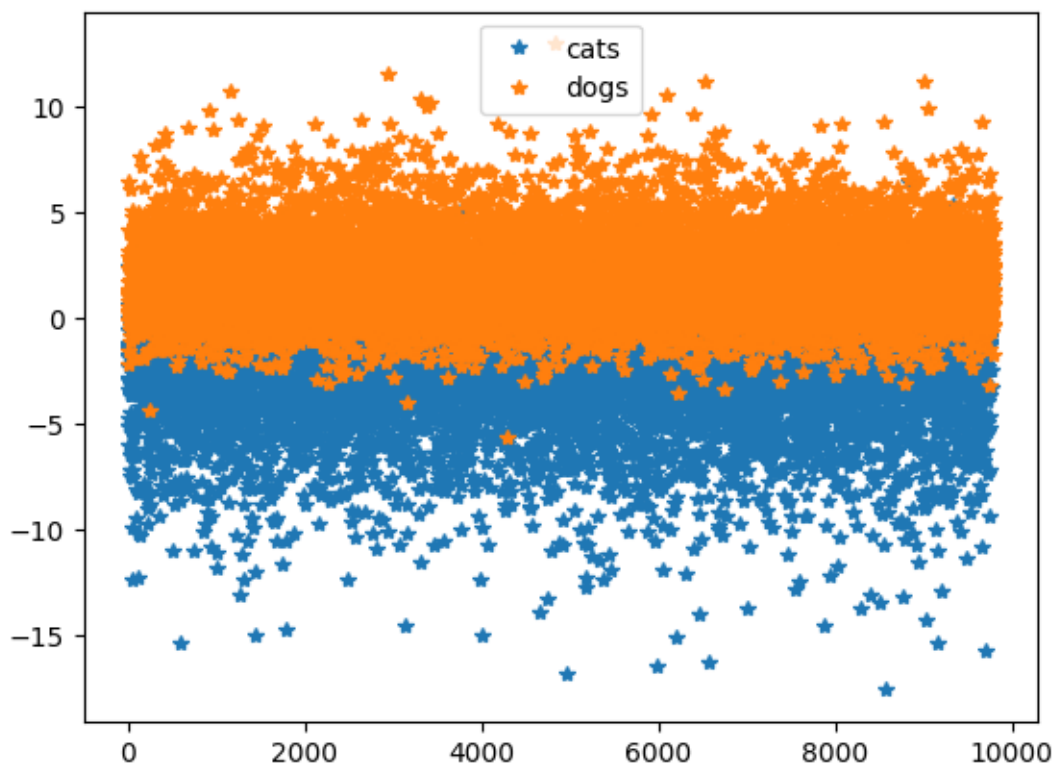
[6]: plot( net2(train_data[train_label==0]).view(-1).cpu().detach().numpy(), '*', ↵
↪label='cats')
plot( net2(train_data[train_label==1]).view(-1).cpu().detach().numpy(), '*', ↵
↪label='dogs')
legend()

```

```


[6]: <matplotlib.legend.Legend at 0x7f3ba1400310>

```



```
[7]: plot( net2(valid_data[valid_label==0]).view(-1).cpu().detach().numpy(), '*',  
         ↪label='cats')  
plot( net2(valid_data[valid_label==1]).view(-1).cpu().detach().numpy(), '*',  
     ↪label='dogs')  
legend()
```

```
[7]: <matplotlib.legend.Legend at 0x7f3ba147a3e0>
```



[ ]:




# What is a layer?

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

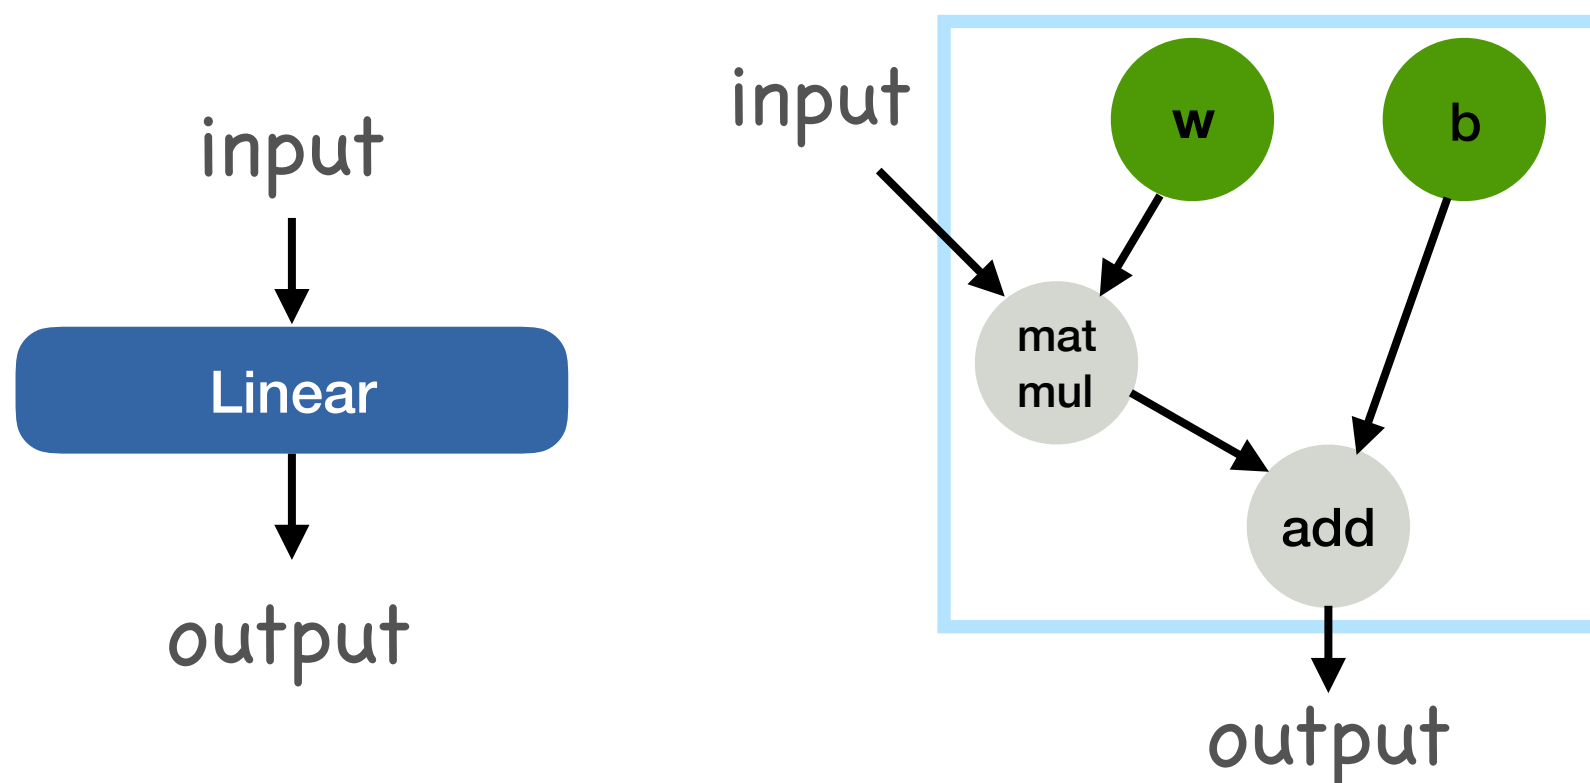
# Examples of layers

- Linear
- ReLU
- Loss functions





# What is a layer?

- Largest computational unit that remains unchanged throughout different architectures




# How many layers does a deep network have?

- We only count linear layers



# Layer naming




# Activation functions

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Non-linearities

- Allow a deep network to model arbitrary differentiable functions



# Zoo of activation functions

ReLU

Leaky ReLU

PReLU

ELU

tanh


Sigmoid



Maxout



# Sigmoid


- $\frac{1}{1 + e^{-x}}$
- Same as tanh
- Do not use





$f(x)$	
$\frac{df(x)}{dx}$	

# ReLU


- $\max(x, 0)$



$f(x)$	
$\frac{df(x)}{dx}$	


# Dead ReLUs



- Prevent dead ReLUs:
- Initialize Network carefully
- Decrease the learning rate



# Leaky ReLU



- $\max(x, \alpha x)$
- For  $0 < \alpha < 1$
- Called PReLU if  $\alpha$  is learned



$f(x)$	
$\frac{df(x)}{dx}$	


# ELU

- $$\begin{cases} x & \text{for } x \geq 0 \\ \alpha(e^x - 1) & \text{for } x < 0 \end{cases}$$



# Maxout

- $\max(x_1, x_2)$



# Which activation to choose?

- ReLU
  - Carefully initialize
  - Small enough learning rate
- If ReLU fails, try:
  - Leaky ReLU / PReLU
- Avoid sigmoid and tanh

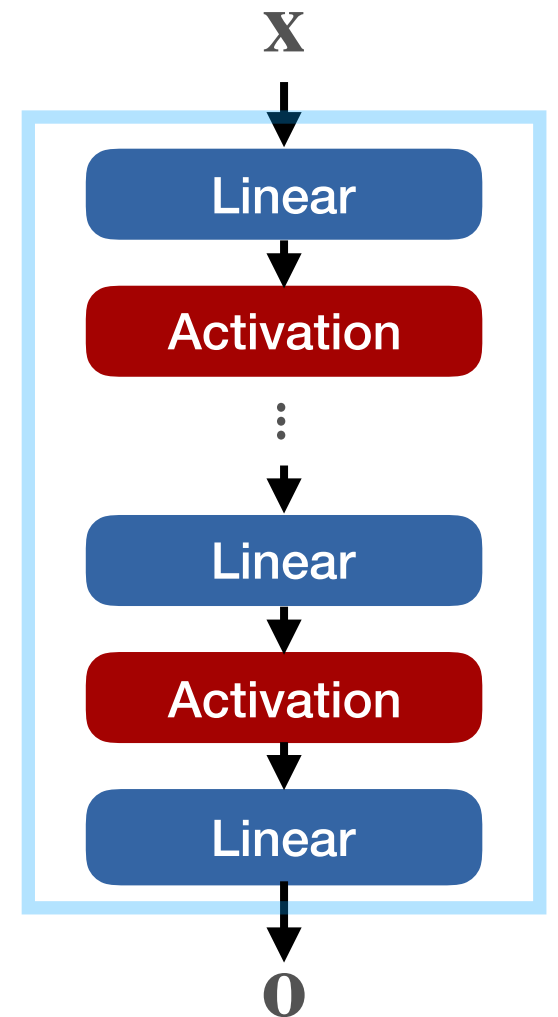

# Hyper-parameters

© 2019 Philipp Krähenbühl and Chao-Yuan Wu



# Hyper-parameters

- Any parameters set by hand and not learned by SGD



# Summary, a practical guide to deep network design

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

# Loss functions

Continuous output?

N

Y

Classification

Regression




2 classes

>2 classes

BCEWithLogitsLoss

CrossEntropyLoss

L1Loss / MSELoss



# Activation functions

- Try ReLU
- If ReLU fails, try:
  - Leaky ReLU / PReLU

# SGD

- Select batch size
- Tune learning rate
- momentum = 0.9

