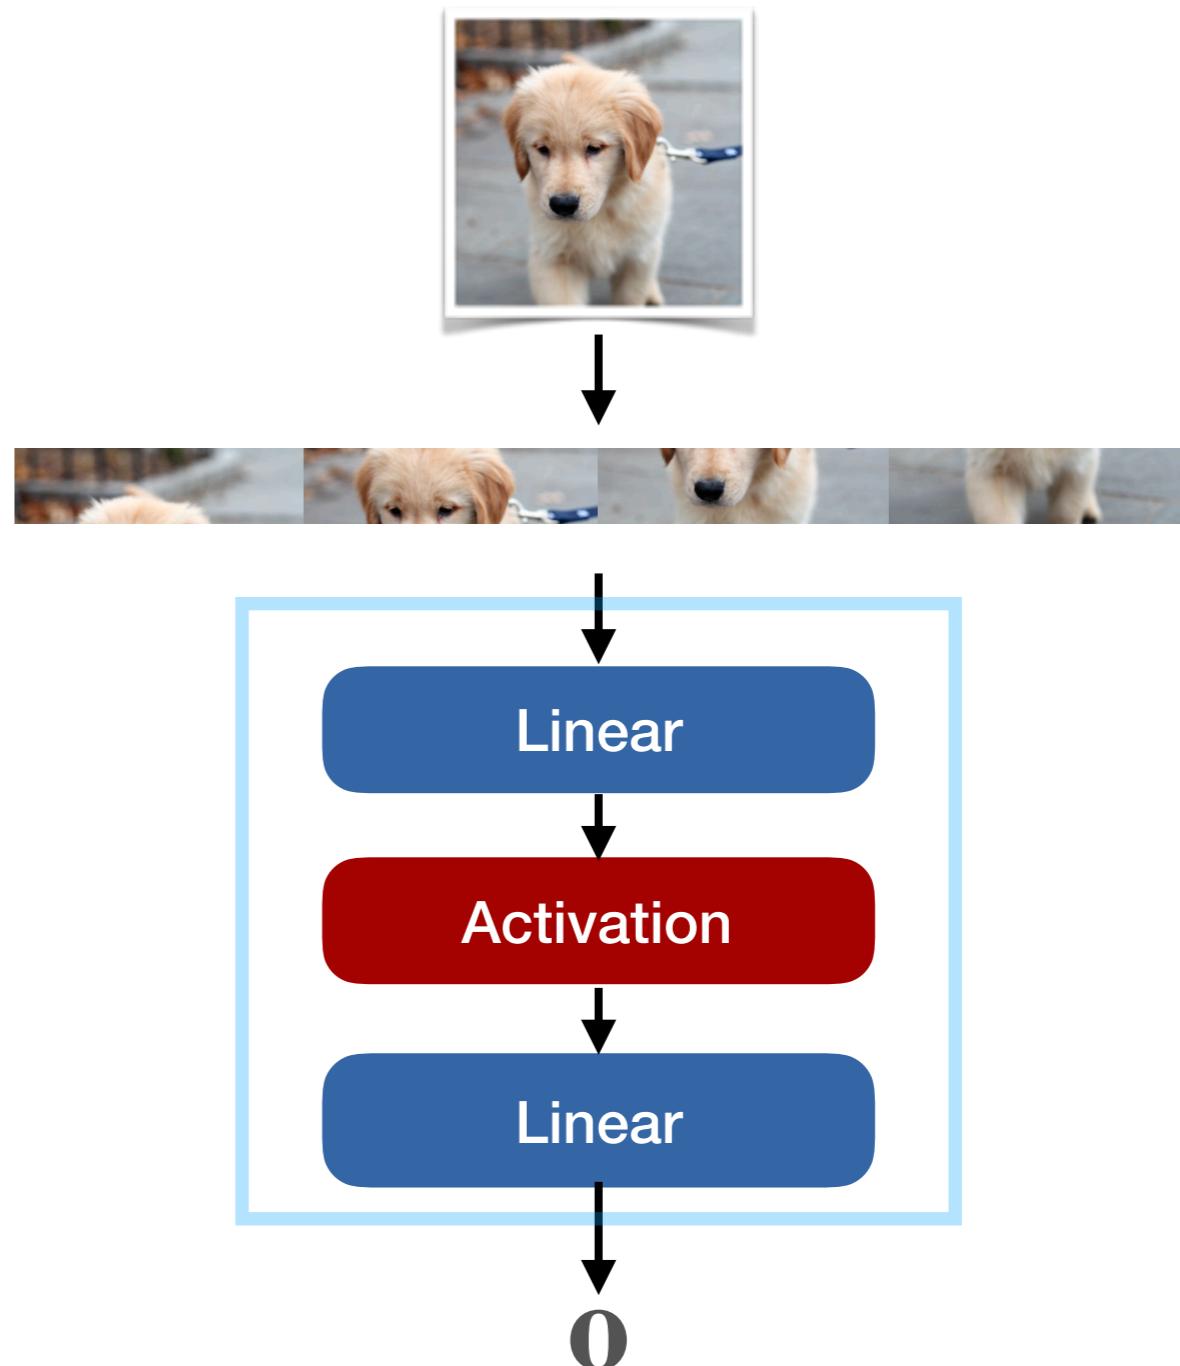


Images and structure

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

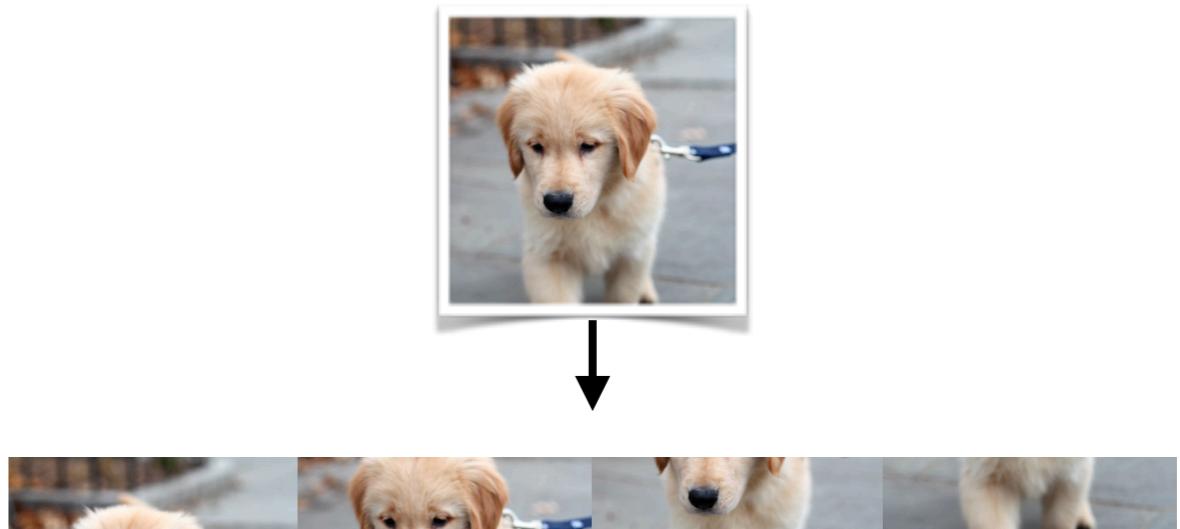
Images and structure

- Fully-connected networks

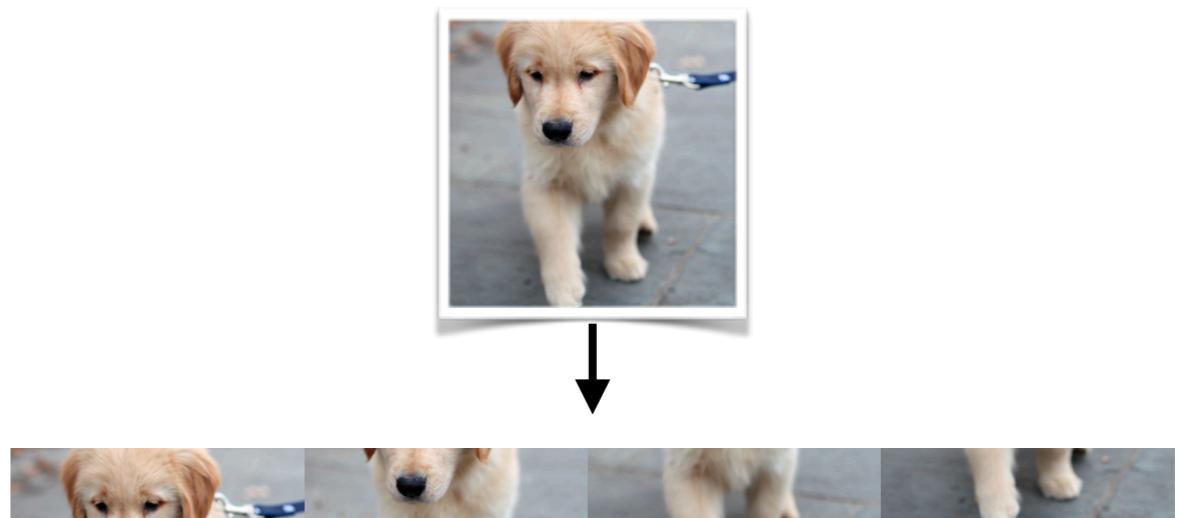


Images and structure

- Visual patterns are shift, rotation and scale invariant



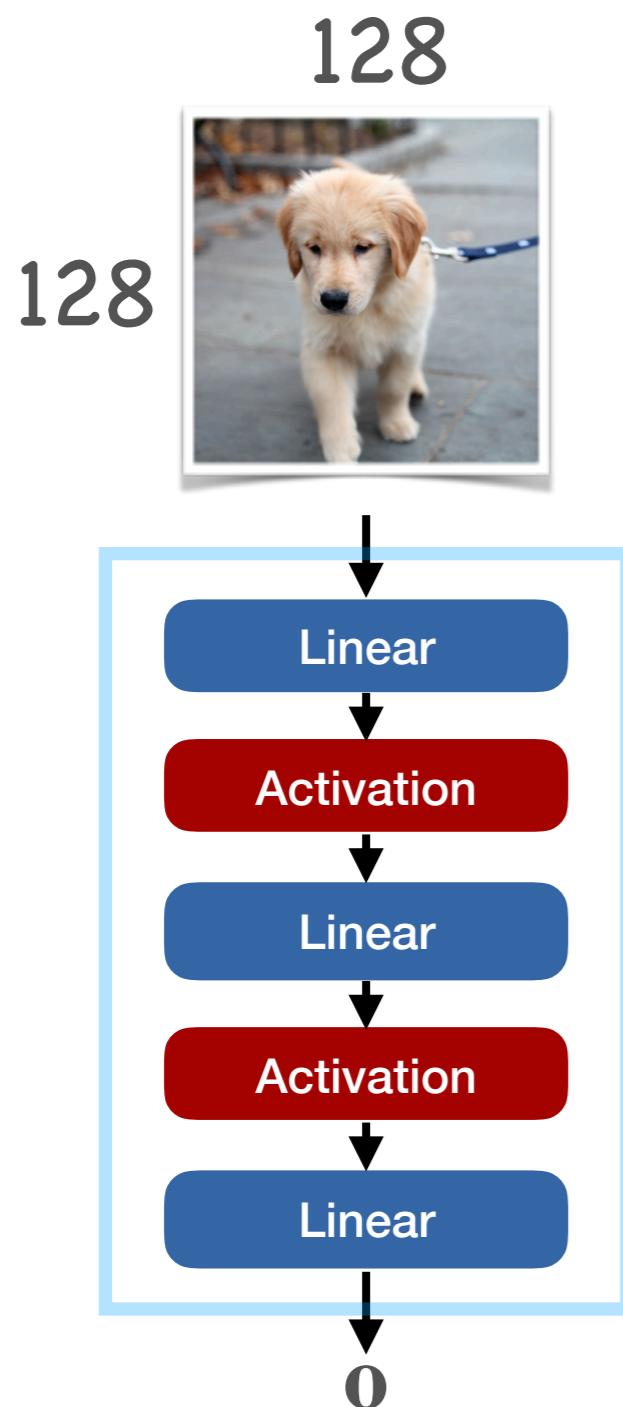
- Fully connected networks are not



High dimensional inputs

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

High dimensional inputs

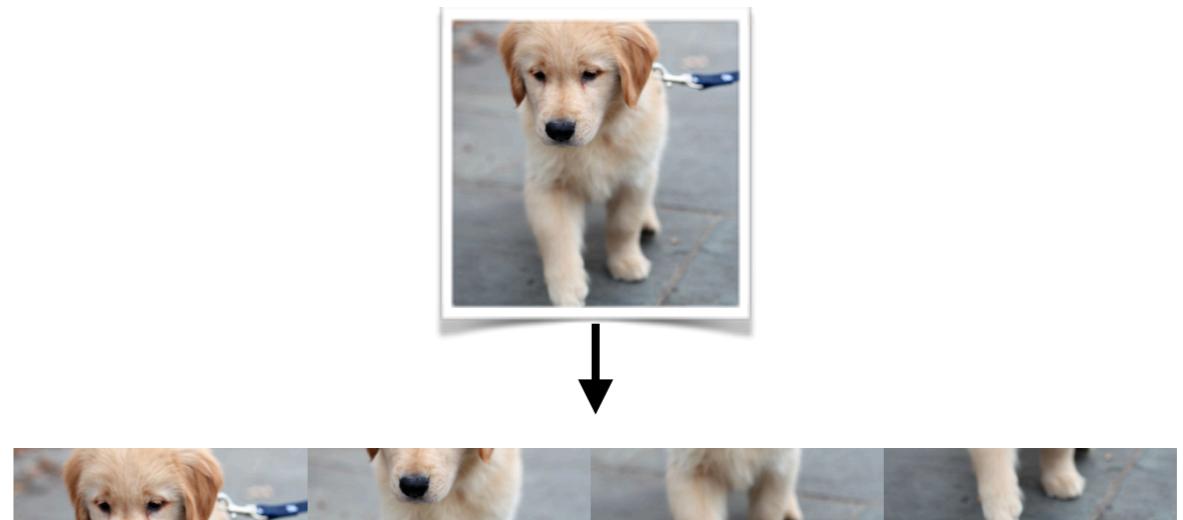
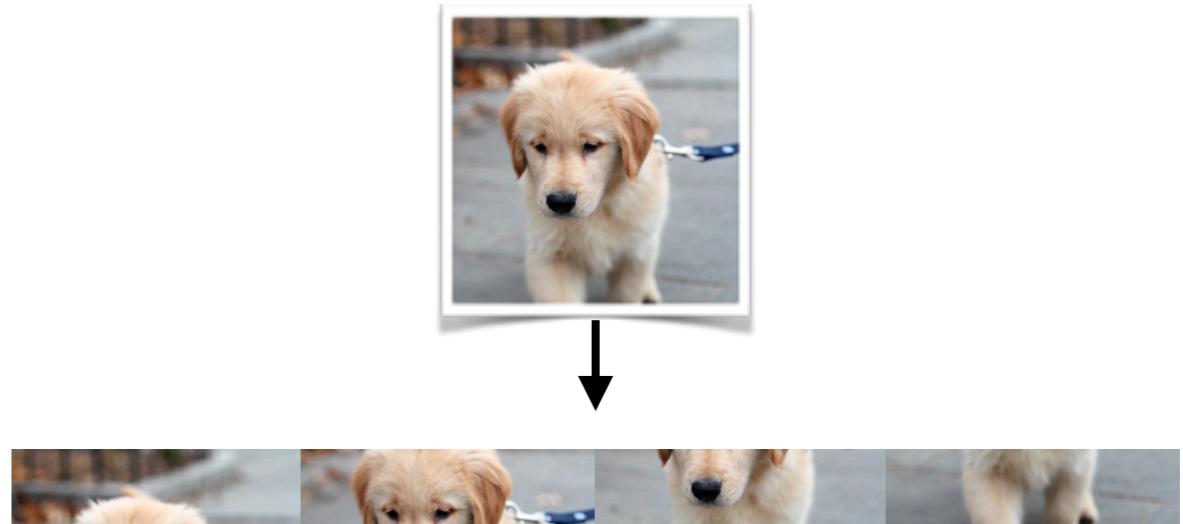


Convolutions

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Images and structure

- Fully connected networks are not shift invariant

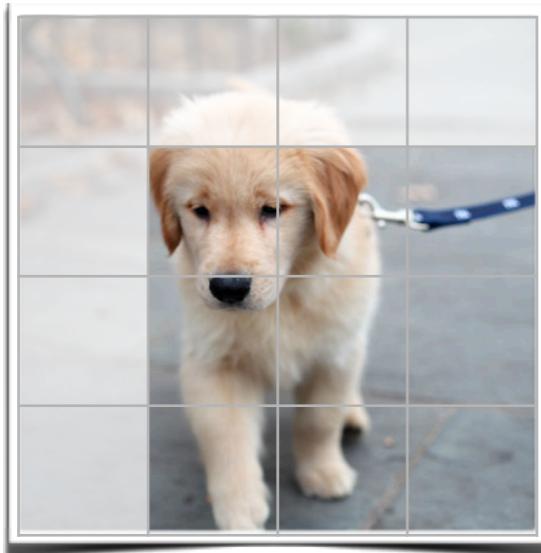


Finding shift-invariant patterns



Convolutions

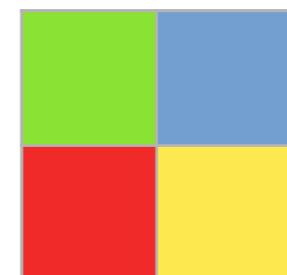
- “Sliding” linear transformation



*

a	b	c
d	e	f
g	h	i

=



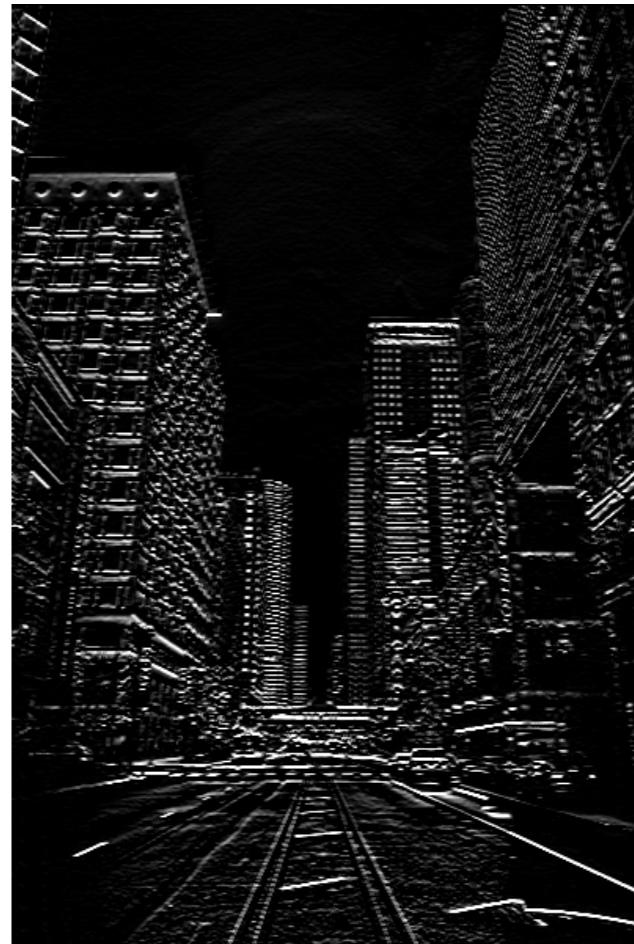
Examples of convolutions



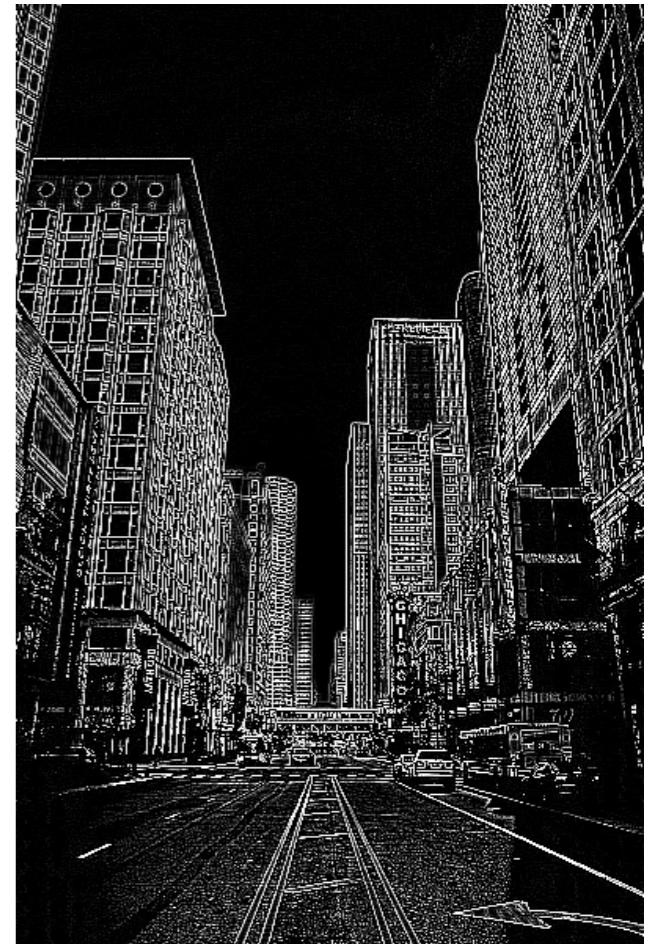
Original



Vertical edges



Horizontal edges



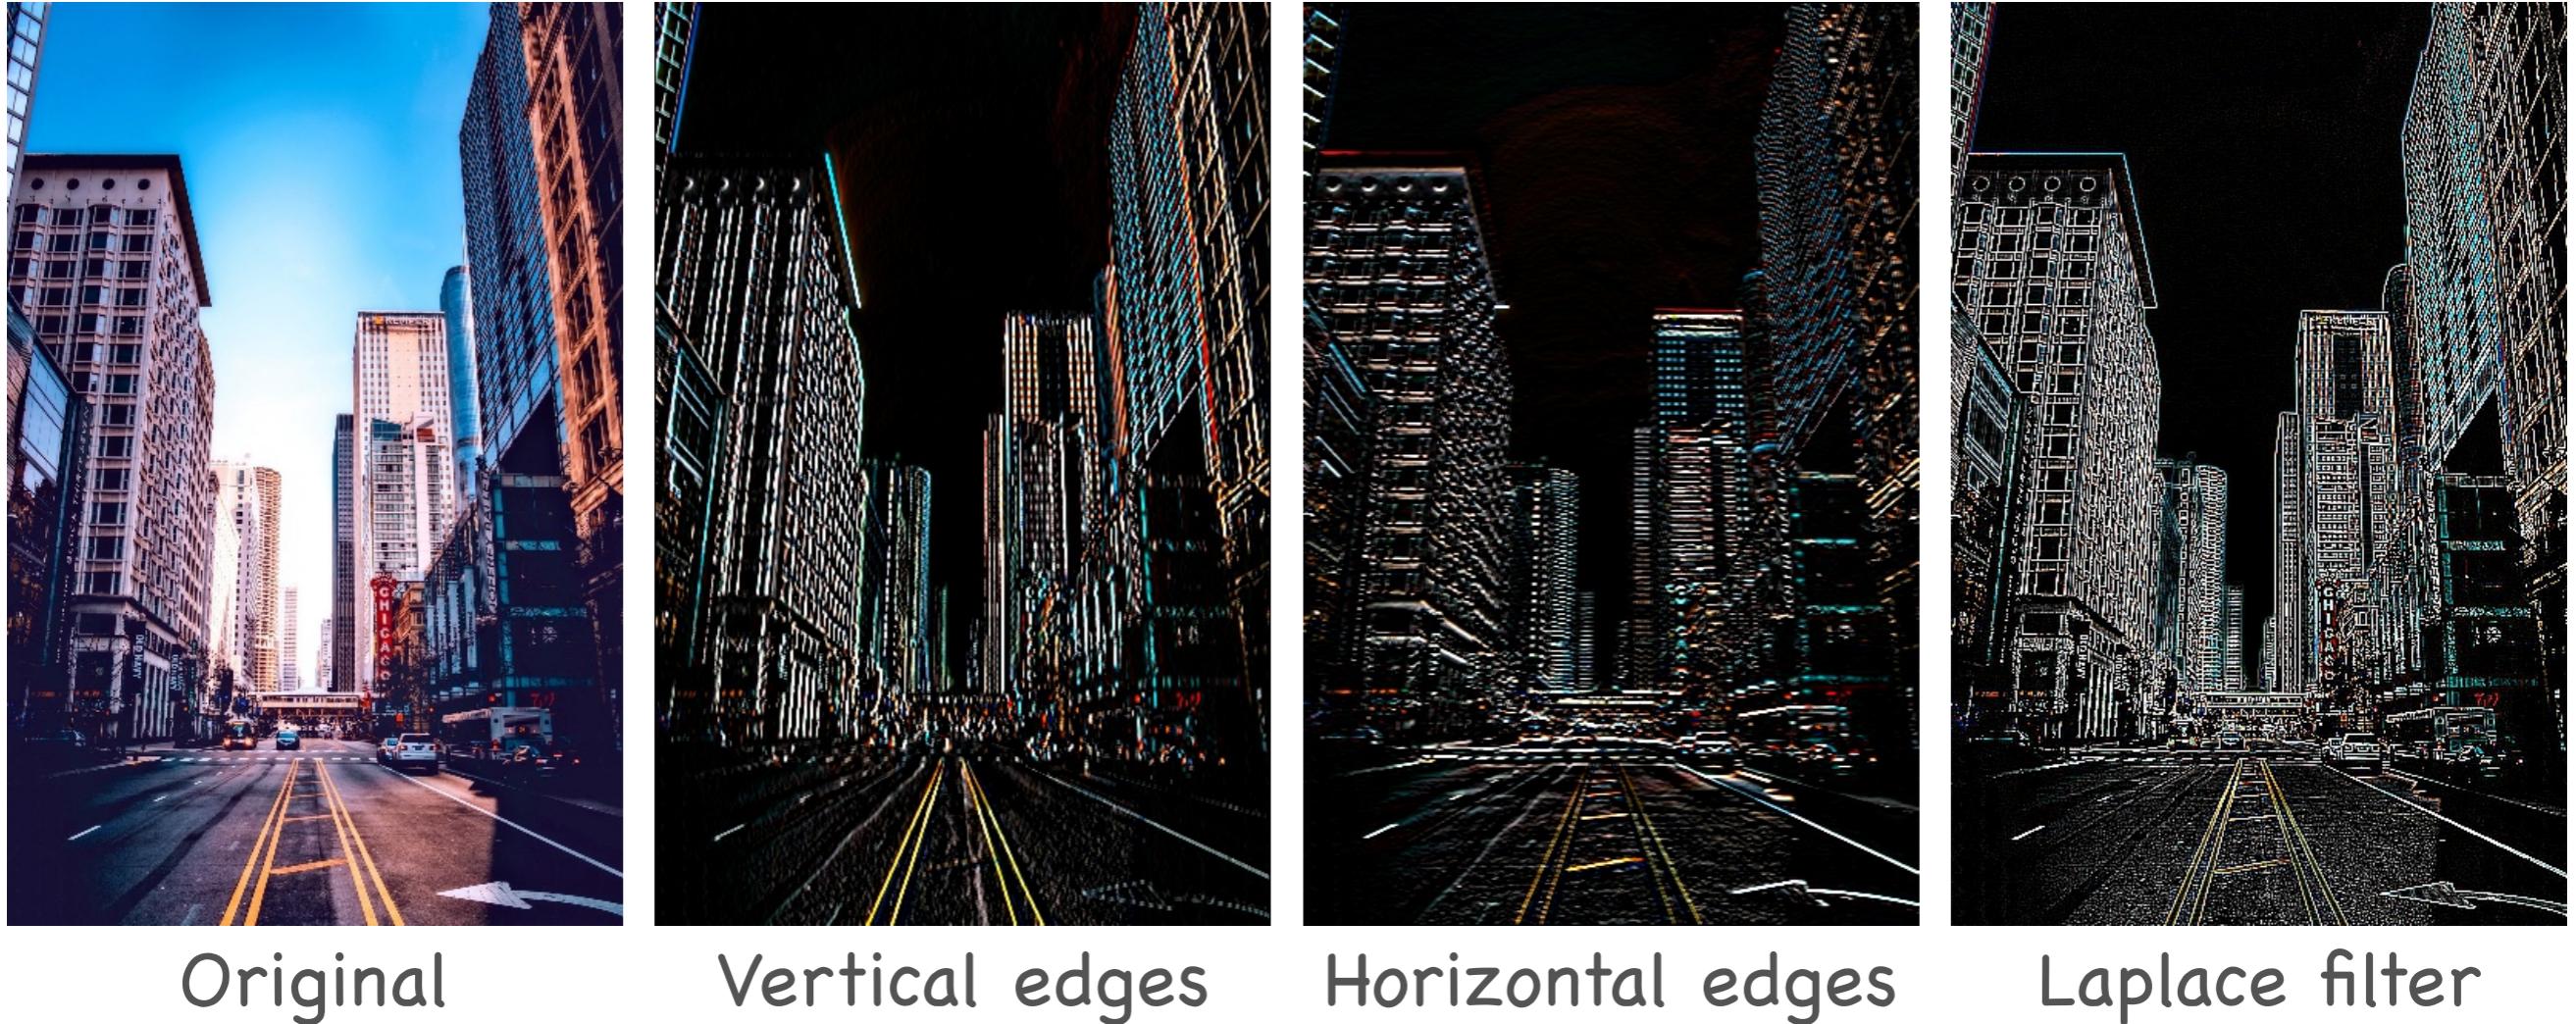
Laplace filter

-1	0	1
-1	0	1
-1	0	1

-1	-1	-1
0	0	0
1	1	1

-1	-1	-1
-1	8	-1
-1	-1	-1

Convolutions on multiple channels



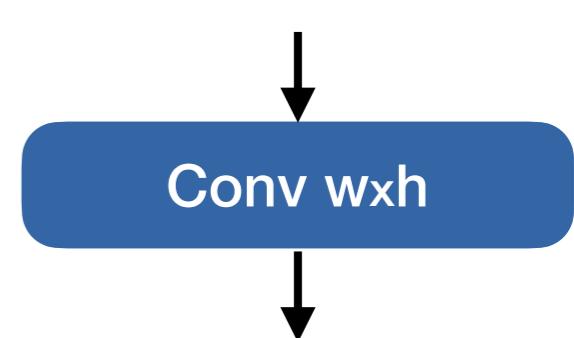
$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$

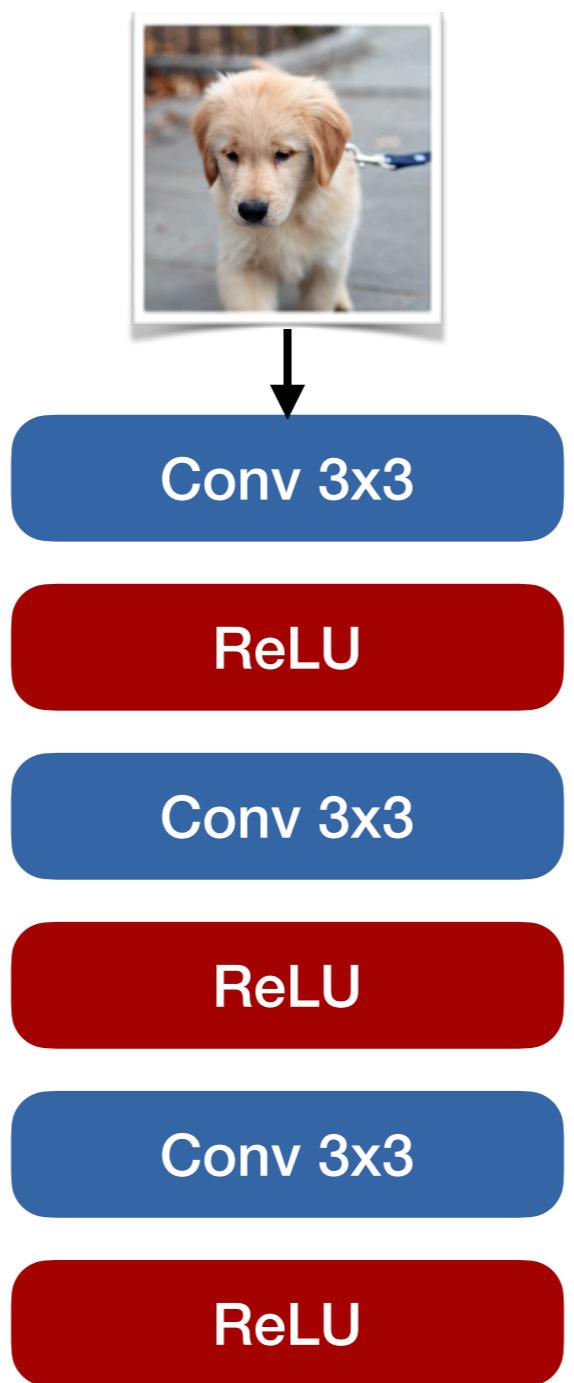
Formal definition

- Input: $\mathbf{X} \in \mathbb{R}^{H \times W \times C_1}$
- Kernel: $\mathbf{w} \in \mathbb{R}^{h \times w \times C_1 \times C_2}$
- Bias: $\mathbf{b} \in \mathbb{R}^{C_2}$
- Output: $\mathbf{Z} \in \mathbb{R}^{(H-h+1) \times (W-w+1) \times C_2}$



$$\mathbf{Z}_{a,b,c} = \mathbf{b}_c + \sum_{i=0}^h \sum_{j=0}^w \sum_{k=0}^{C_1} \mathbf{X}_{a+i,b+j,k} w_{i,j,k,c}$$

Stacking multiple layers



Convolution as a linear layer

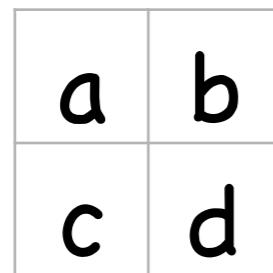
input: $3 \times 3 \times 1$



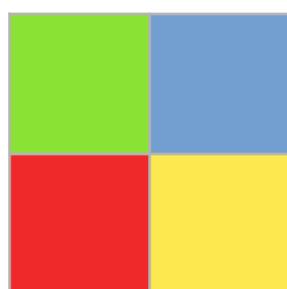
input: 9



kernel: $2 \times 2 \times 1 \times 1$



output: $2 \times 2 \times 1$



weight: 4×9

a	b		c	d				
		a	b		c	d		
				a	b	c	d	
					a	b	c	d

output: 4



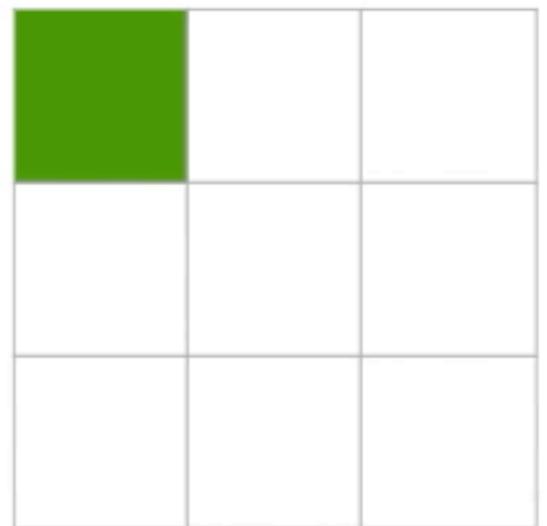
Special case: 1×1 convolution

- Pixel-wise linear transformation
- Kernel: $1 \times 1 \times C_1 \times C_2$



*

W =



January 23, 2024

```
[1]: %pylab inline
import torch
import sys
sys.path.append('..')
sys.path.append('../..')
from data import load
train_data, train_label = load.get_dogs_and_cats_data(resize=(32,32),
n_images=10)
device = torch.device('cuda') if torch.cuda.is_available() else torch.
device('cpu')
print('device = ', device)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
Populating the interactive namespace from numpy and matplotlib
device = cuda

```
[2]: class ConvNet(torch.nn.Module):
    def __init__(self, layers=[], n_input_channels=3, kernel_size=3):
        super().__init__()
        L = []
        c = n_input_channels
        for l in layers:
            L.append(torch.nn.Conv2d(c, l, kernel_size))
            L.append(torch.nn.ReLU())
            c = l
        L.append(torch.nn.Conv2d(c, 1, kernel_size=1))
        self.layers = torch.nn.Sequential(*L)

    def forward(self, x):
        return self.layers(x).mean(dim=[1,2,3])

net = ConvNet([32,64])
```

```
[3]: print( train_data[:1].shape )
print( net(train_data[:1]).shape )
```

```
torch.Size([1, 3, 32, 32])
torch.Size([1])
```

```
[4]: net2 = ConvNet([32, 64, 128])
print( net2(train_data[:1]).shape )
```

```
torch.Size([1])
```

```
[5]: %load_ext tensorboard
import tempfile
log_dir = tempfile.mkdtemp()
%tensorboard --logdir {log_dir} --reload_interval 1
```

```
<IPython.core.display.HTML object>
```

```
[6]: from util import train
train.train(net2, batch_size=128, resize=(32,32), log_dir=log_dir, device=device, n_epochs=100)
```

```
WARNING:root:loading dataset
```

```
WARNING:root:loading done
```

```
0%|          | 0/100 [00:00<?, ?it/s]
```

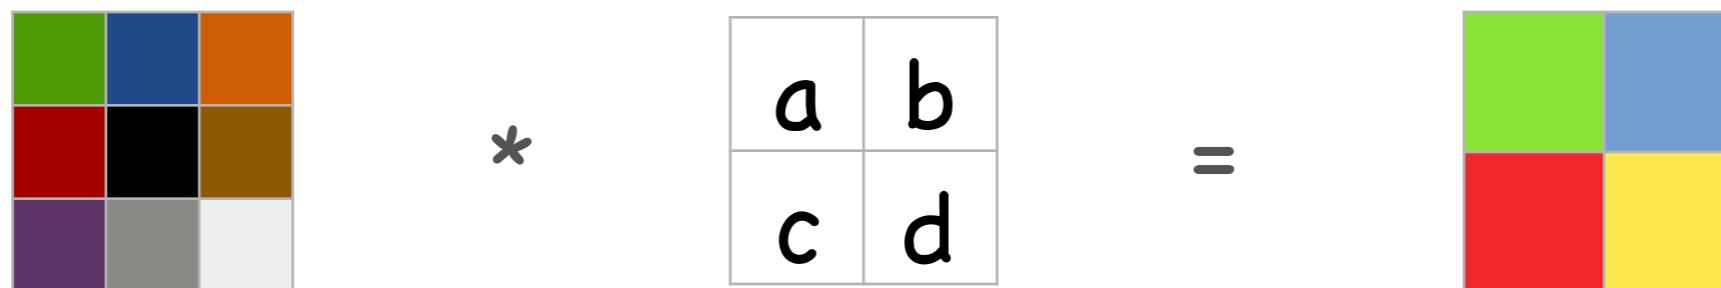
```
[ ]:
```

Convolutional operators and their structure

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Output size

- Input: $\mathbf{X} \in \mathbb{R}^{H \times W \times C_1}$
- Kernel: $\mathbf{w} \in \mathbb{R}^{h \times w \times C_1 \times C_2}$
- Output: $\mathbf{Z} \in \mathbb{R}^{(H-h+1) \times (W-w+1) \times C_2}$



Padding

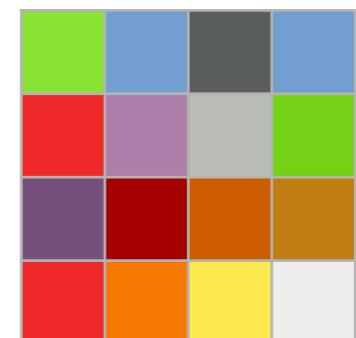
- Add p_w, p_h zeros in each dimension
- Input: $X \in \mathbb{R}^{H \times W \times C_1}$
- Kernel: $w \in \mathbb{R}^{h \times w \times C_1 \times C_2}$
- Output: $Z \in \mathbb{R}^{(H-h+2p_h+1) \times (W-w+2p_w+1) \times C_2}$

0	0	0	0	0	0
0	green	blue	orange		0
0	red	black	brown		0
0	purple	gray			0
0	0	0	0	0	0

*

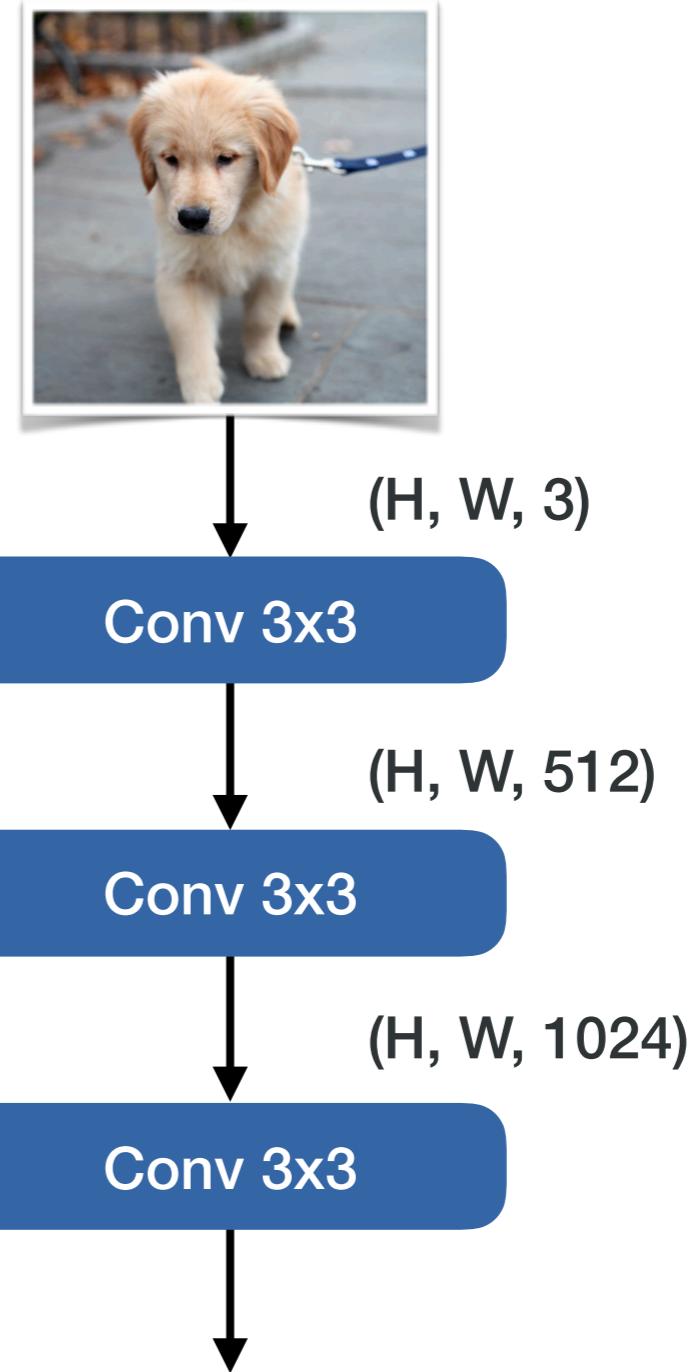
a	b
c	d

=



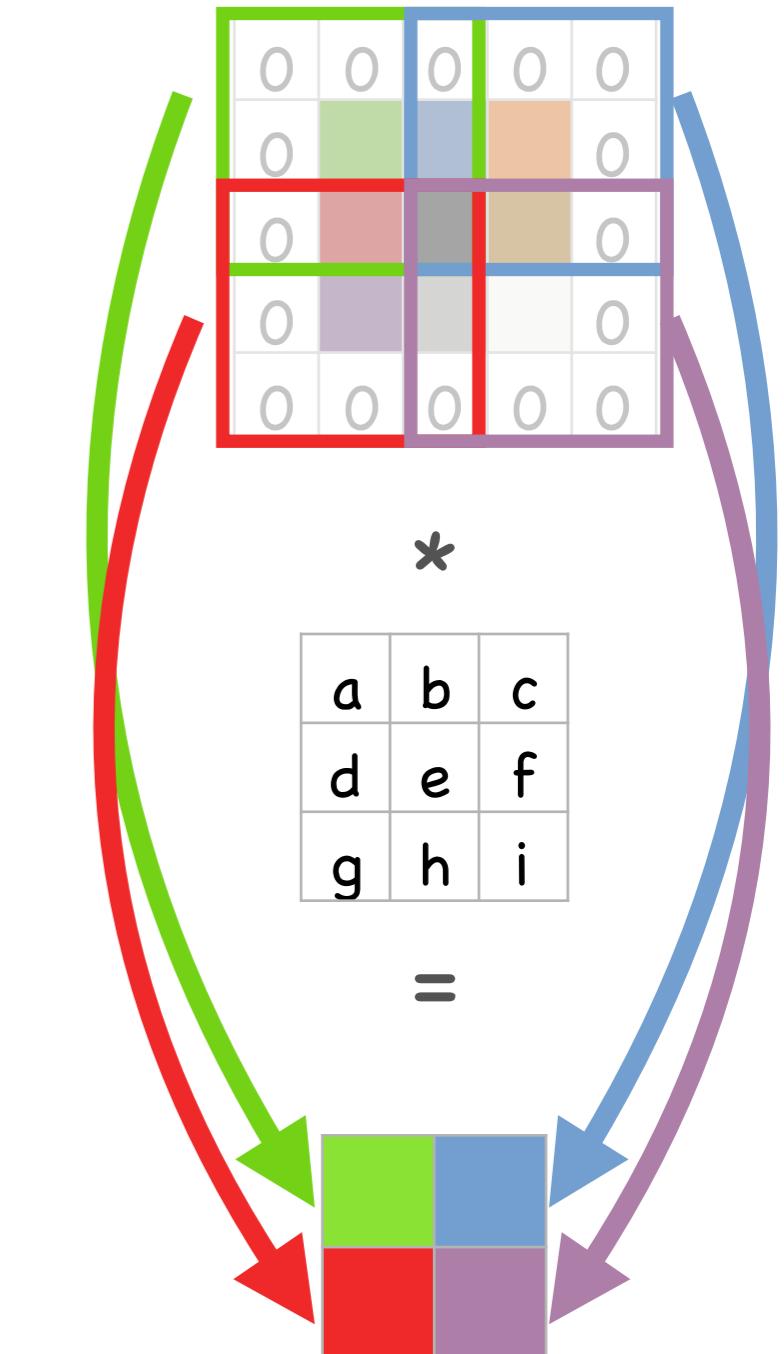
Output resolution

- High output resolution
 - Slow computation

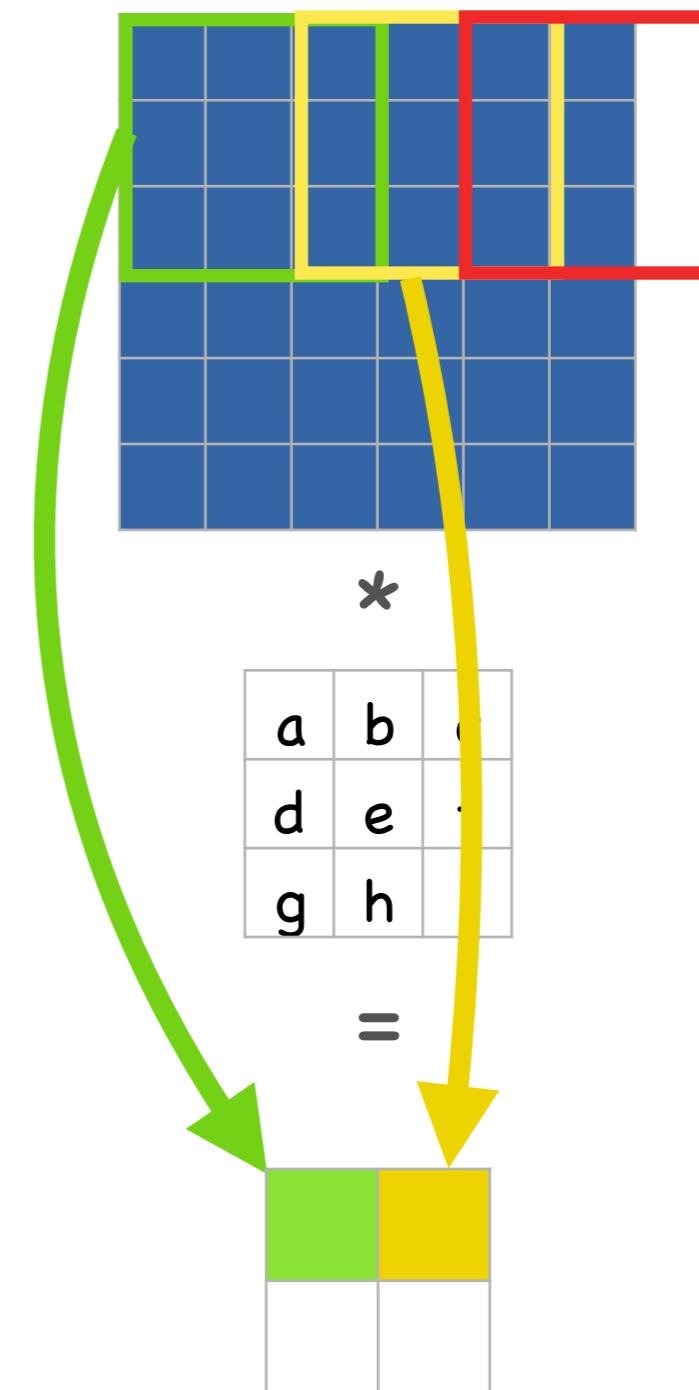
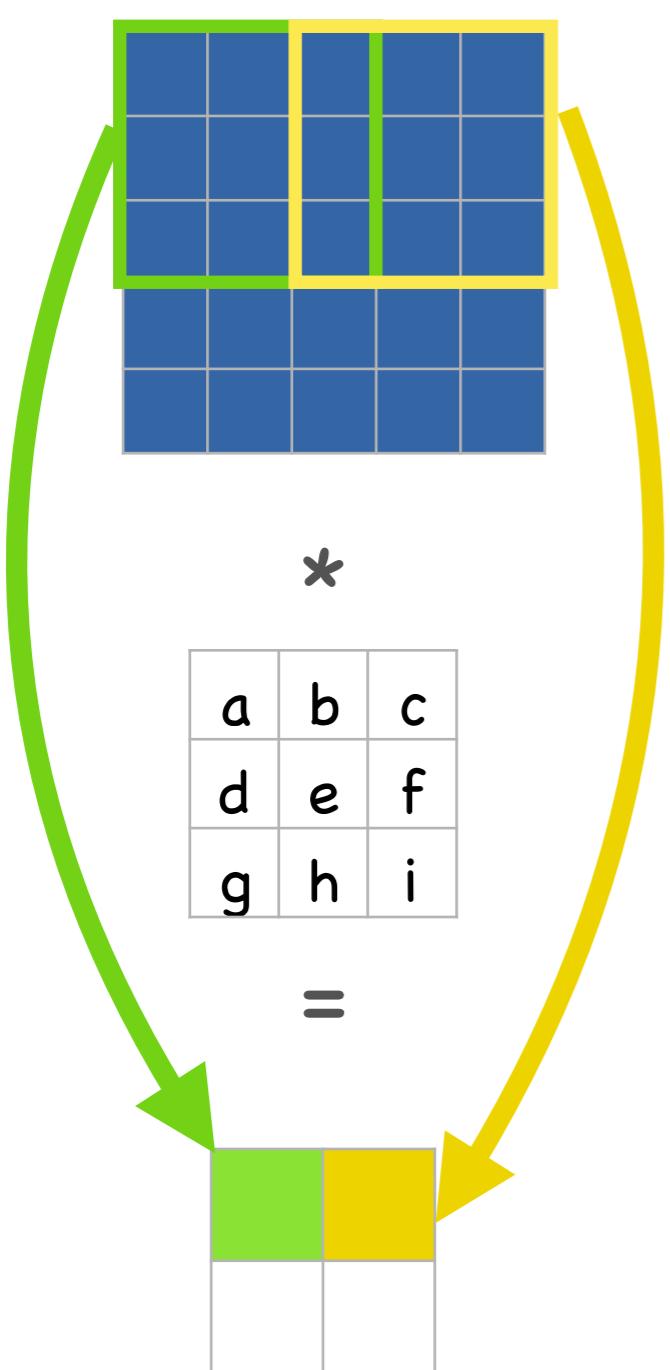


Striding

- Only compute every n-th output: s_w, s_h
- Input: $\mathbf{X} \in \mathbb{R}^{H \times W \times C_1}$
- Kernel: $\mathbf{w} \in \mathbb{R}^{h \times w \times C_1 \times C_2}$
- Output:
$$\mathbf{z} \in \mathbb{R}^{\left(\frac{H-h+2p_h}{s_h}+1\right) \times \left(\frac{W-w+2p_w}{s_w}+1\right) \times C_2}$$

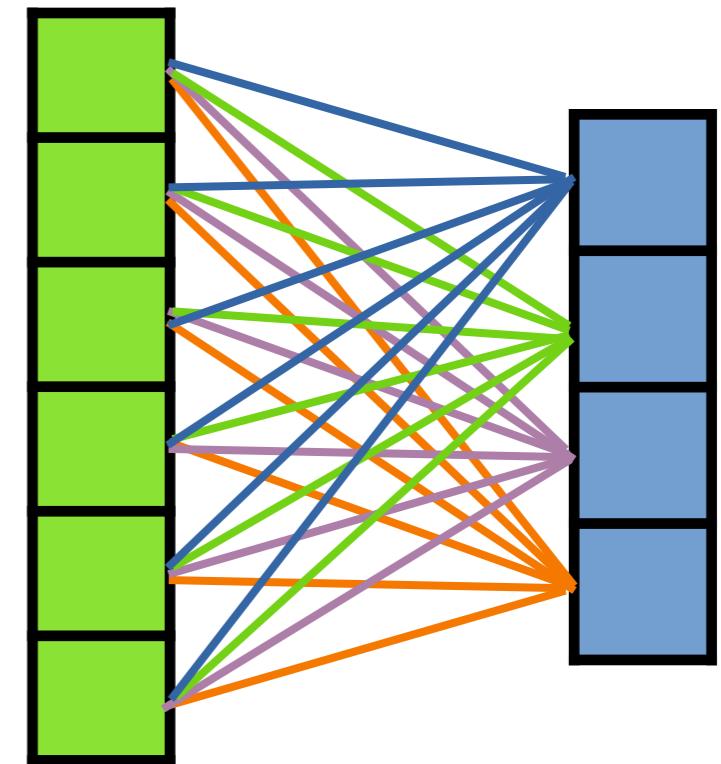


Output size with striding



Parameters

- Every input channel C_1 is connected to every output channel C_2

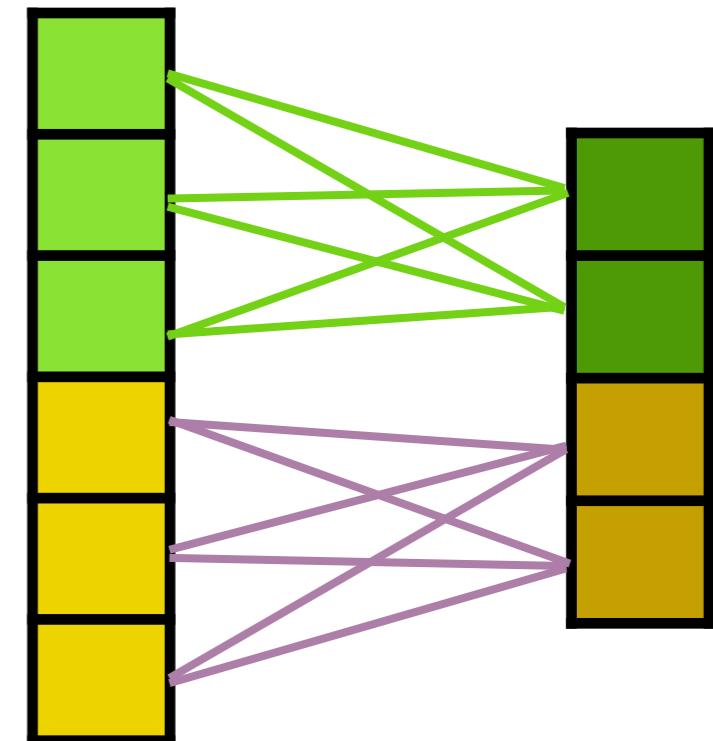


$$C_1 = 6$$

$$C_2 = 4$$

Grouping

- Split channels into g groups
- Reduce parameters and computation by factor g



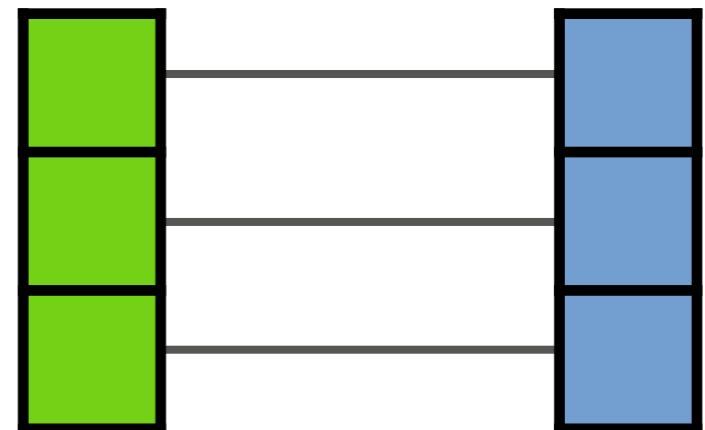
$$C_1 = 6$$

$$C_2 = 4$$

Depthwise convolution

- Special grouping

- $C_1 = g$
- $C_2 = g$

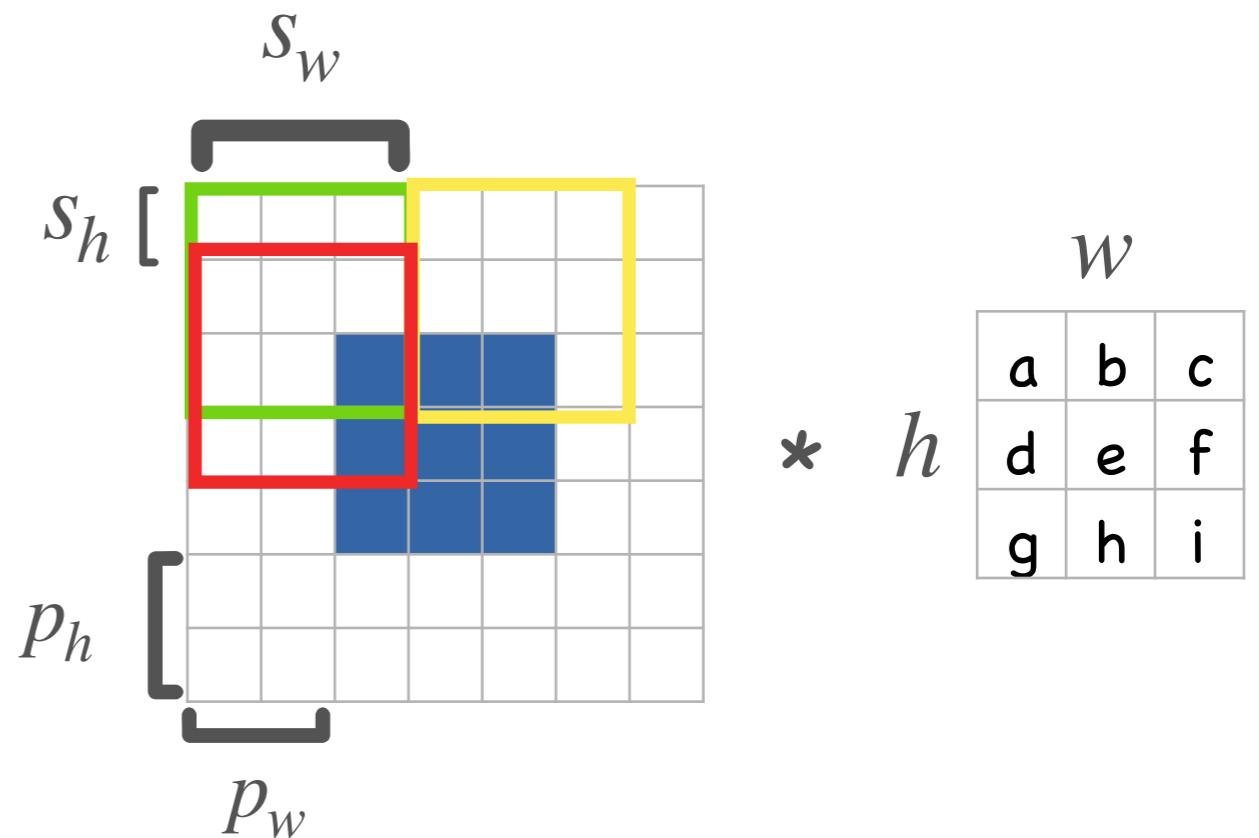


$$C_1 = 3$$

$$C_2 = 3$$

Hyper-parameters of convolutions

- Kernel size: $w \times h$
- Padding: p_w, p_h
- Stride: s_w, s_h



Convolutional operators

- Run arbitrary operation $f(\mathbf{x})$ "over" image

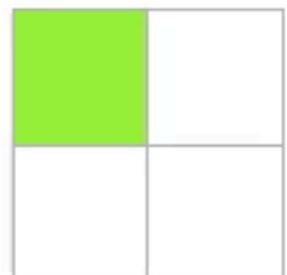
$f(\mathbf{x})$



*

a	b	c
d	e	f
g	h	i

=

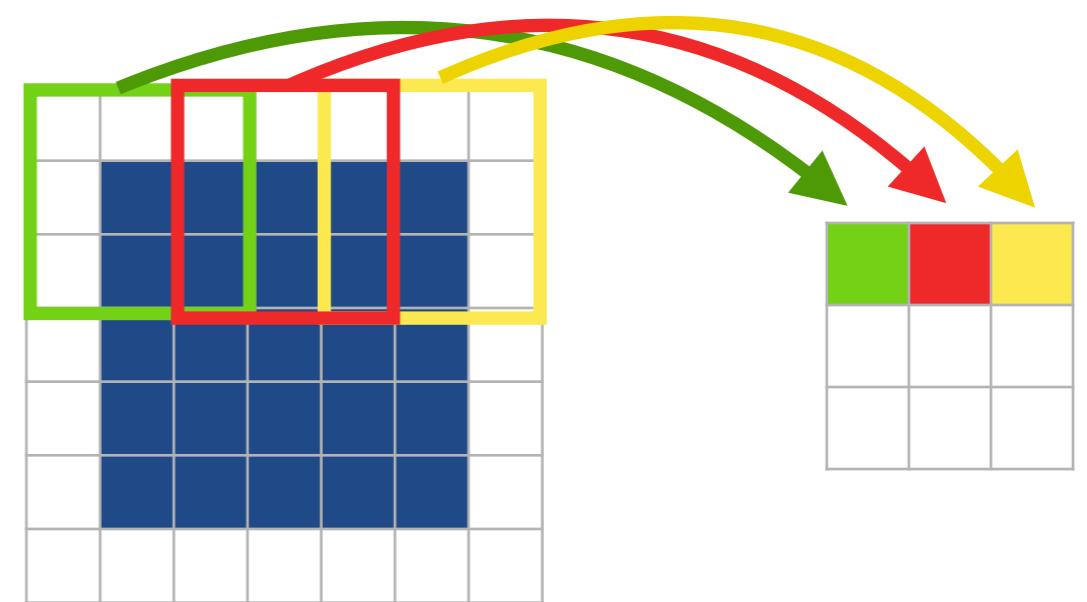


Average pooling

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

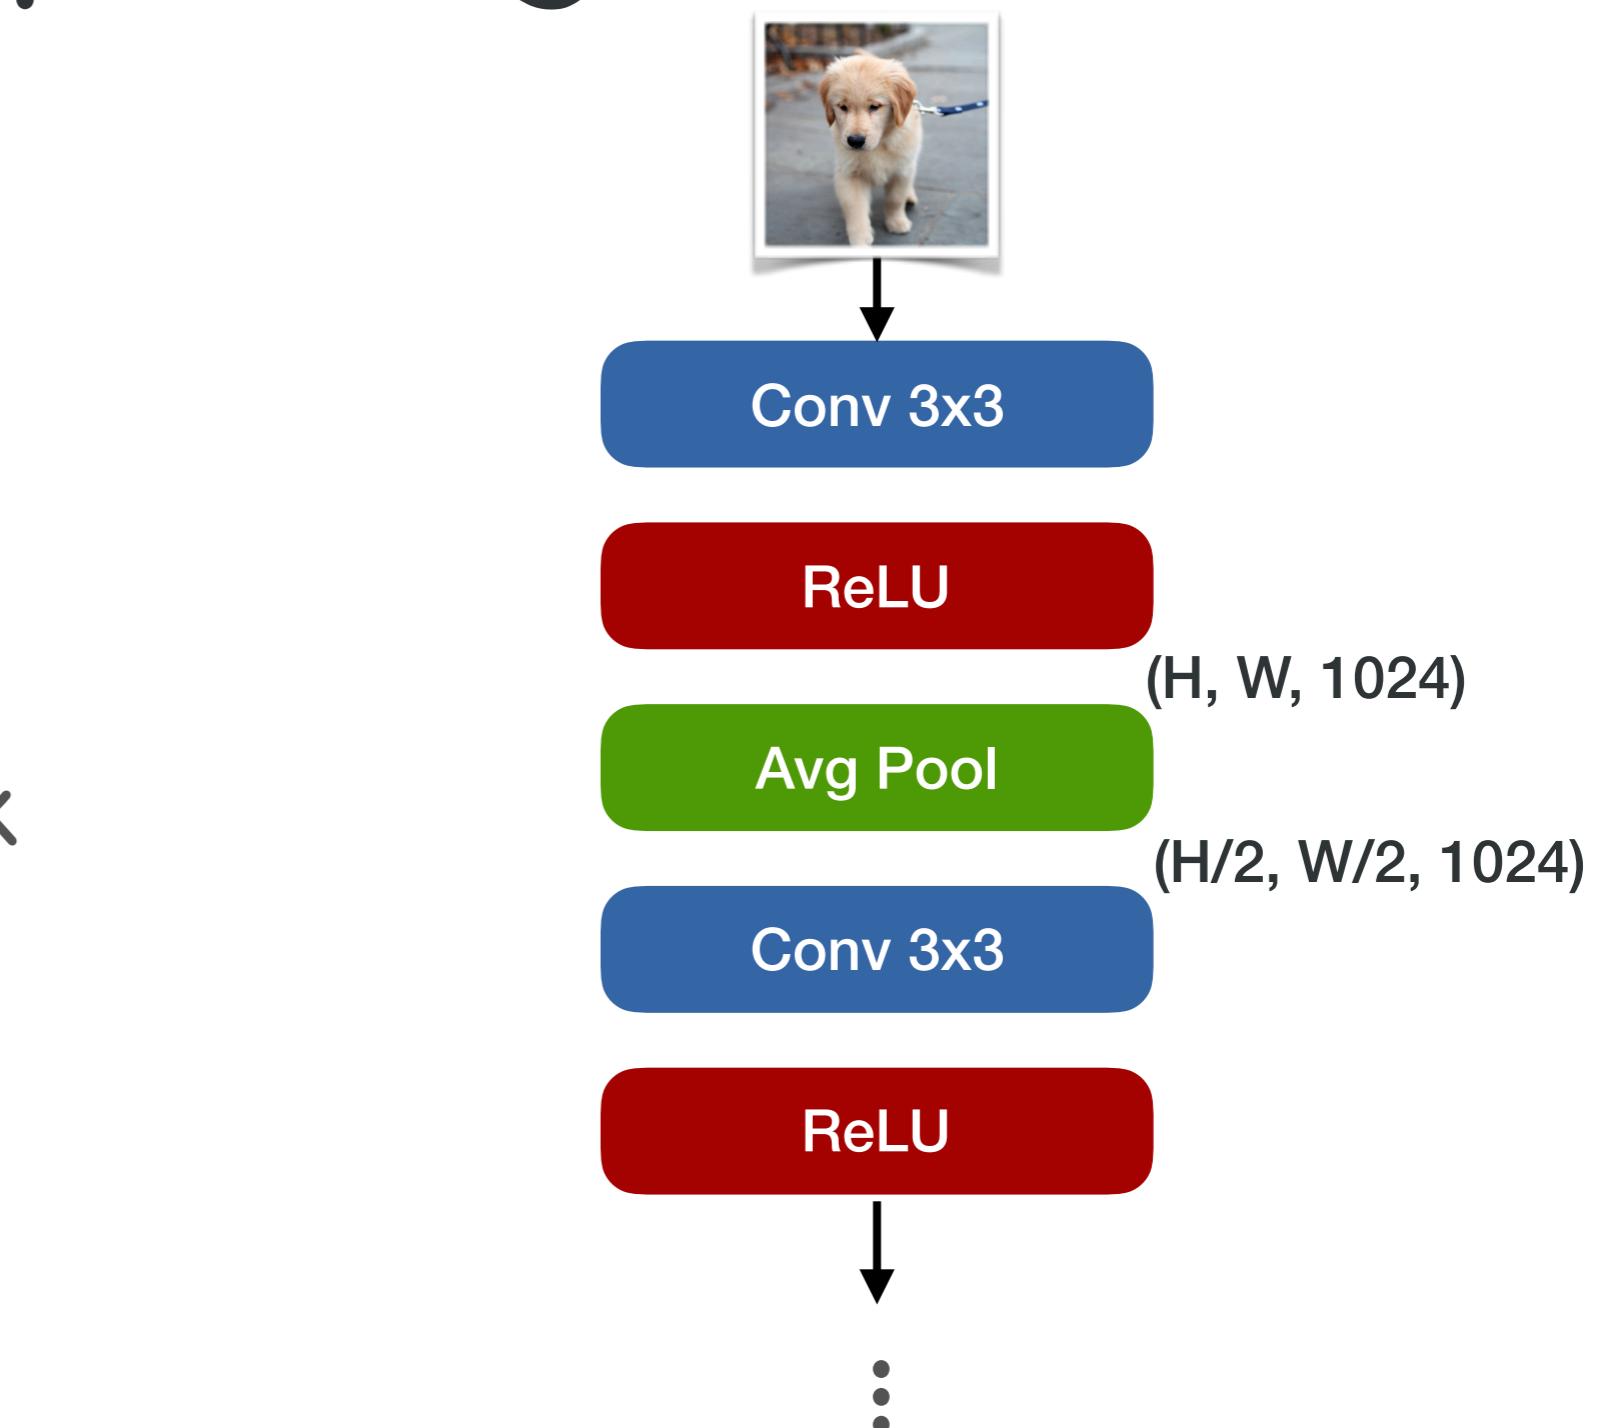
Average pooling

- Convolutional operator
 - $f_c(\mathbf{x}) = \text{mean}_{i,j}(\mathbf{x}_{i,j,c})$



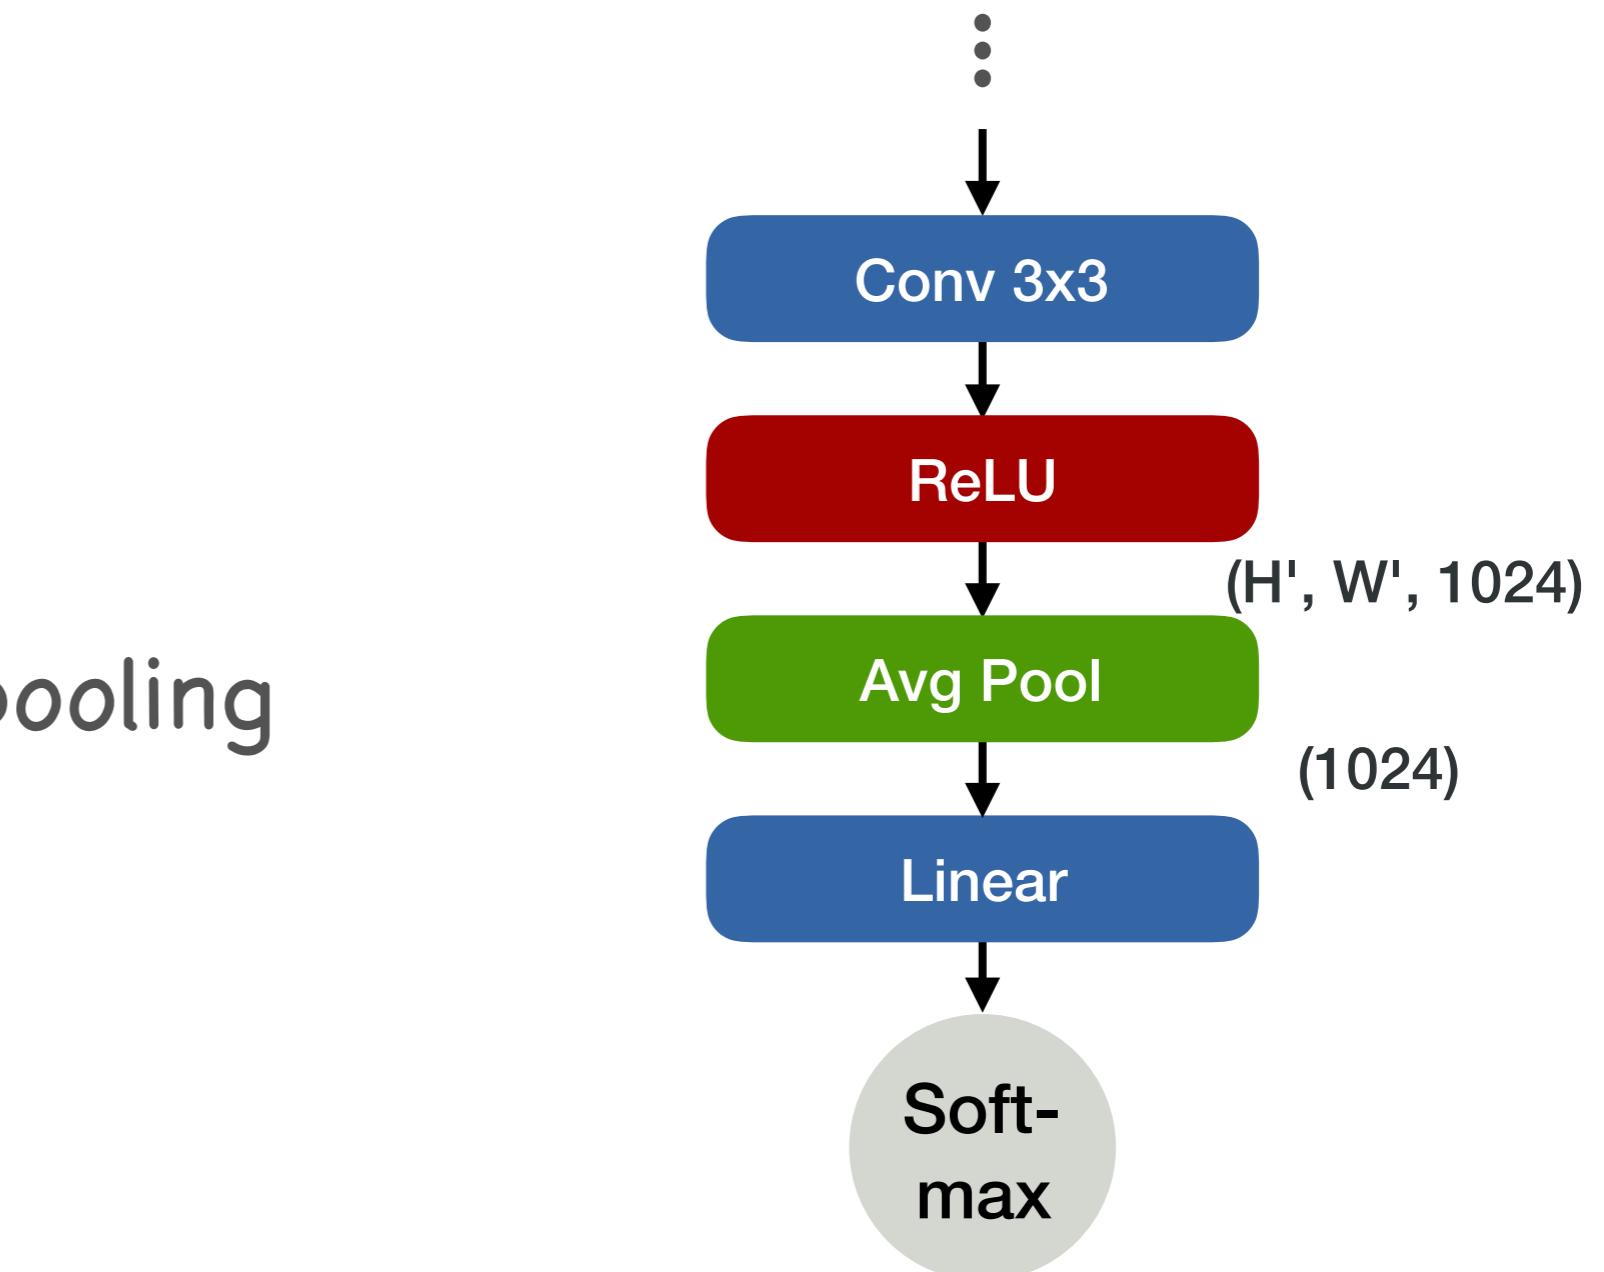
Where to use average pooling?

- Older networks:
 - Inside a network



Where to use average pooling?

- Modern networks
 - Global average pooling

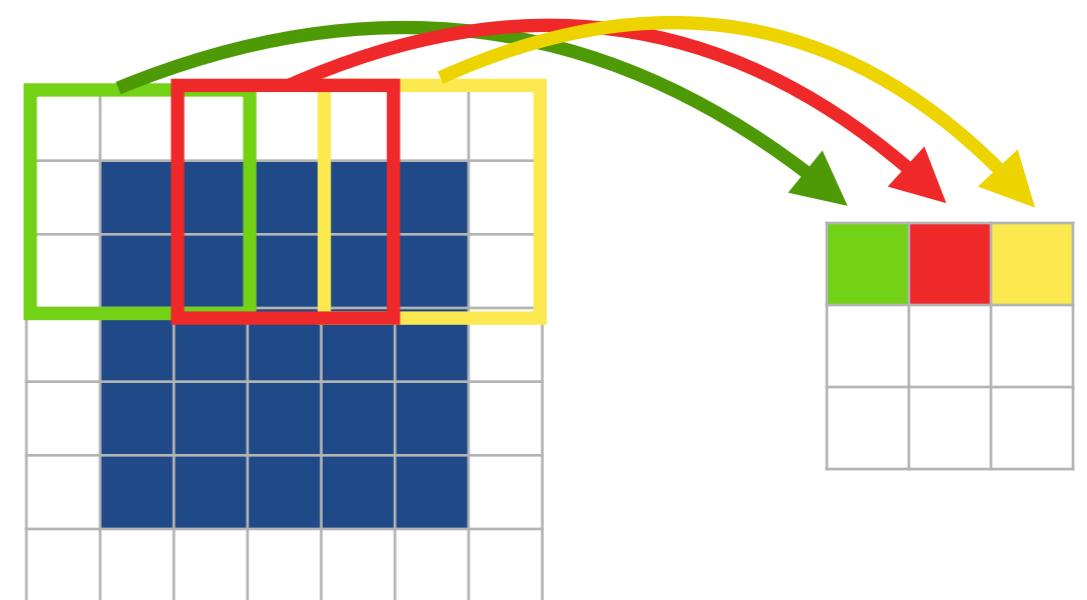


Max pooling

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

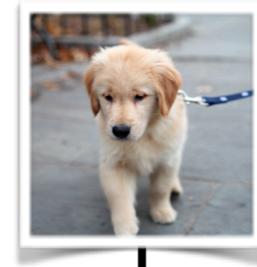
Max pooling

- Convolutional operator
 - $f_c(\mathbf{x}) = \max_{i,j}(\mathbf{x}_{i,j,c})$



Where to use max pooling?

- Inside a network
 - With strides, as down-sampling



Conv 3x3

ReLU

(H, W, 1024)

Max Pool

(H/2, W/2, 1024)

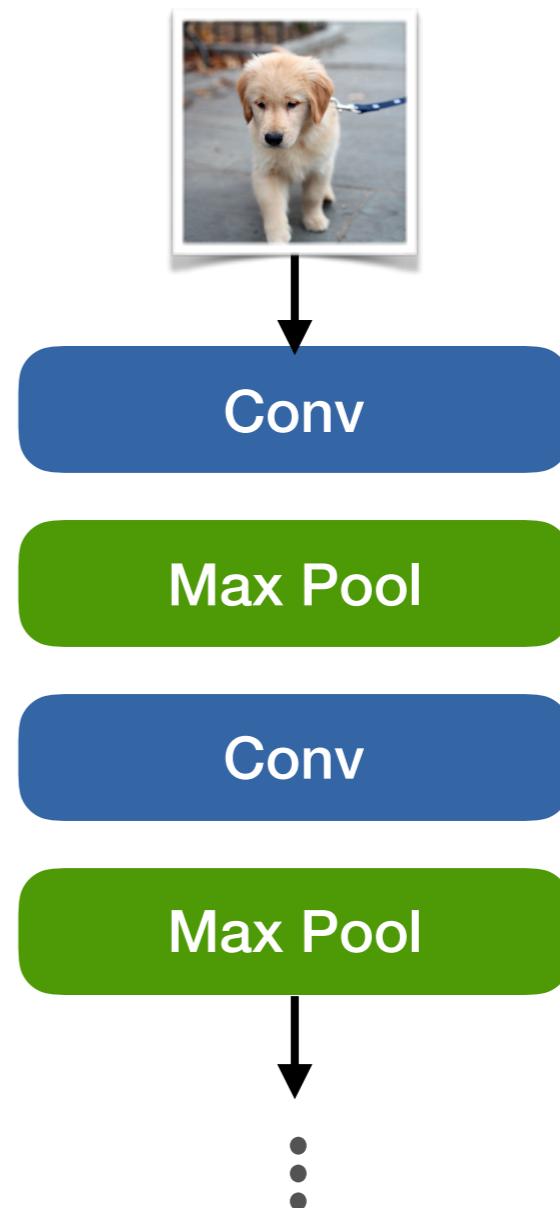
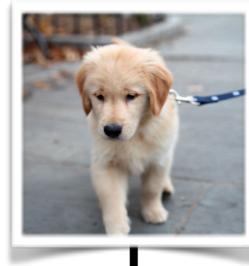
Conv 3x3

ReLU

⋮

Max pooling as a non-linearity

- Similar to maxout



January 23, 2024

```
[1]: %pylab inline
import torch
import sys
sys.path.append('..')
sys.path.append('../..')
from data import load
train_data, train_label = load.get_dogs_and_cats_data(resize=(32,32),
n_images=10)
device = torch.device('cuda') if torch.cuda.is_available() else torch.
device('cpu')
print('device = ', device)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
Populating the interactive namespace from numpy and matplotlib
device = cuda

```
[2]: class ConvNet(torch.nn.Module):
    def __init__(self, layers=[], n_input_channels=3, kernel_size=3, stride=2):
        super().__init__()
        L = []
        c = n_input_channels
        for l in layers:
            L.append(torch.nn.Conv2d(c, l, kernel_size, padding=(kernel_size-1)/
2, stride=stride))
            L.append(torch.nn.ReLU())
            c = l
        L.append(torch.nn.Conv2d(c, 1, kernel_size=1))
        self.layers = torch.nn.Sequential(*L)

    def forward(self, x):
        return self.layers(x).mean([1,2,3])

net = ConvNet([32,64])
```

```
[3]: print( train_data[:1].shape )
print( net(train_data[:1]).shape )
```

torch.Size([1, 3, 32, 32])

```

torch.Size([1])

[4]: net2 = ConvNet([32,64,128])
print( net2(train_data[:1]).shape )

torch.Size([1])

[5]: class ConvNet2(torch.nn.Module):
    def __init__(self, layers=[], n_input_channels=3, kernel_size=3, stride=2):
        super().__init__()
        L = []
        c = n_input_channels
        for l in layers:
            L.append(torch.nn.Conv2d(c, l, kernel_size, padding=(kernel_size-1)/
                                   2))
            L.append(torch.nn.ReLU())
            L.append(torch.nn.MaxPool2d(3, padding=1, stride=stride))
            c = l
        L.append(torch.nn.Conv2d(c, 1, kernel_size=1))
        self.layers = torch.nn.Sequential(*L)

    def forward(self, x):
        return self.layers(x).mean([1,2,3])

net3 = ConvNet2([32,64,128])

```

```

[6]: print( net3(train_data[:1]).shape )

torch.Size([1])

[7]: %load_ext tensorboard
import tempfile
log_dir = tempfile.mkdtemp()
%tensorboard --logdir {log_dir} --reload_interval 1

```

```

<IPython.core.display.HTML object>

```

```

[8]: from util import train
train.train(net2, batch_size=128, resize=(32,32), log_dir=log_dir+ '/net2', ▾
           device=device, n_epochs=100)

```

```

WARNING:root:loading dataset
WARNING:root:loading done
0%|          | 0/100 [00:00<?, ?it/s]

```

```

[9]: from util import train

```

```
train.train(net3, batch_size=128, resize=(32,32), log_dir=log_dir+'/net3',  
device=device, n_epochs=100)
```

```
WARNING:root:loading dataset  
WARNING:root:loading done  
0%|          | 0/100 [00:00<?, ?it/s]
```

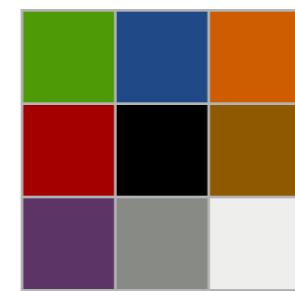
```
[ ]:
```

Receptive fields

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Receptive fields

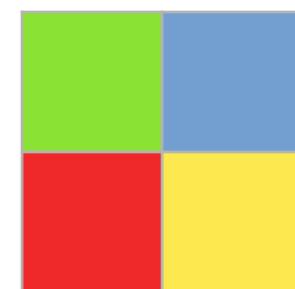
input: $3 \times 3 \times 1$



Conv 2x2



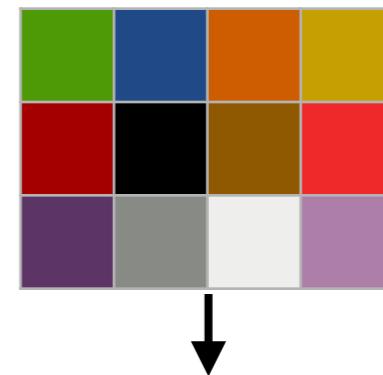
output: $2 \times 2 \times 1$



- Can input x_{abc} affect output z_{ijk} ?

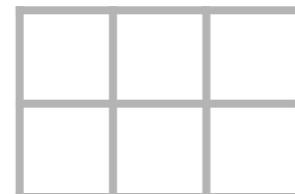
Receptive fields

input: $3 \times 4 \times 1$



- Can input x_{abc} affect output z_{ijk} ?

Conv 2x2



Conv 2x2

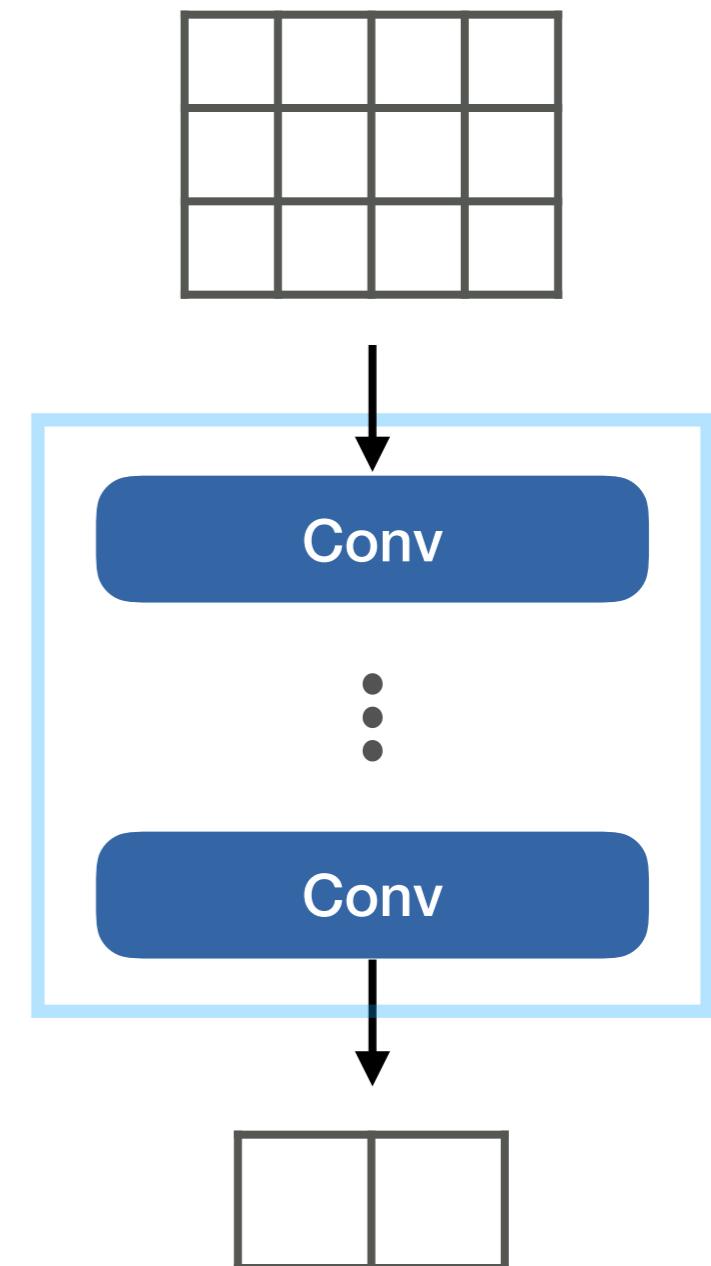
output: $2 \times 2 \times 1$



How do we compute the receptive field

- Option 1: lots of math
- Option 2: Computationally

- Feed a image of 0s to the network
- Change a single element to NaN
- See output changes



Structure of receptive field?

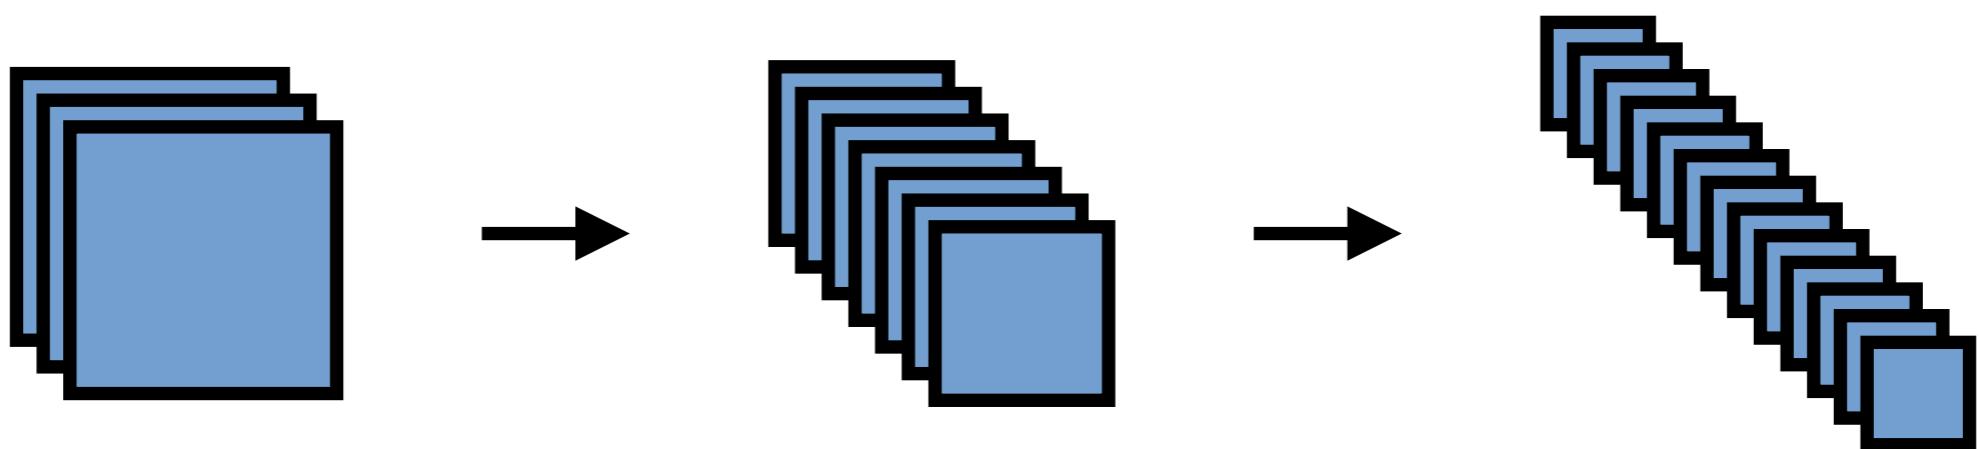


Design principles of convolutional networks

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

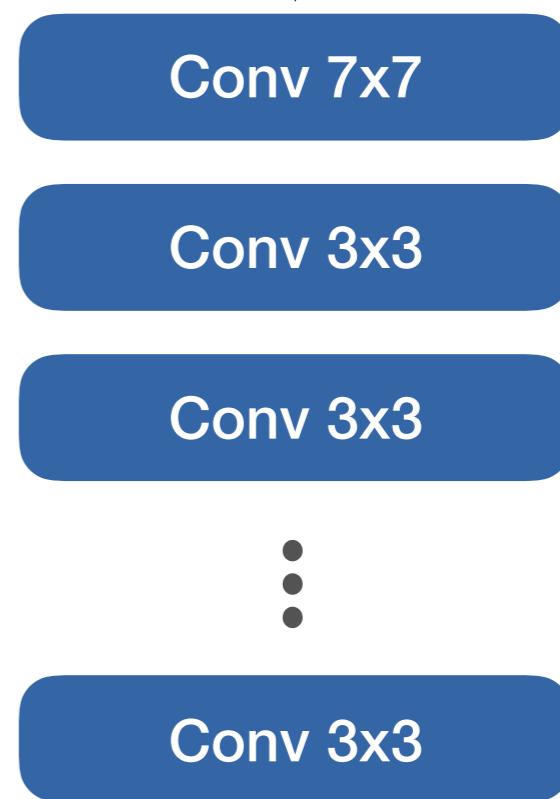
Use striding, increase channels

- Trade spatial resolution for channels
- Balance computation



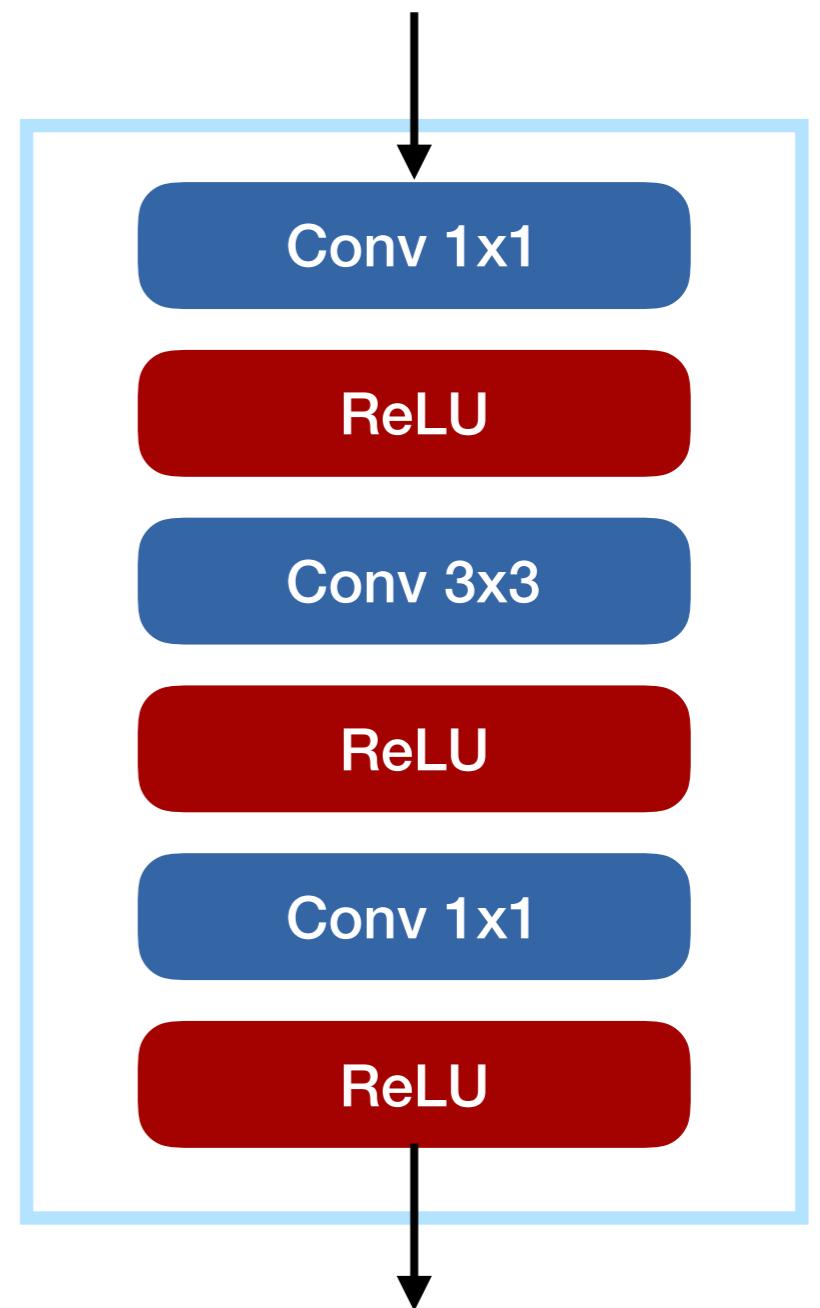
Keep kernels small

- 3x3 kernels almost everywhere
- exception:
 - first layer up to 7x7



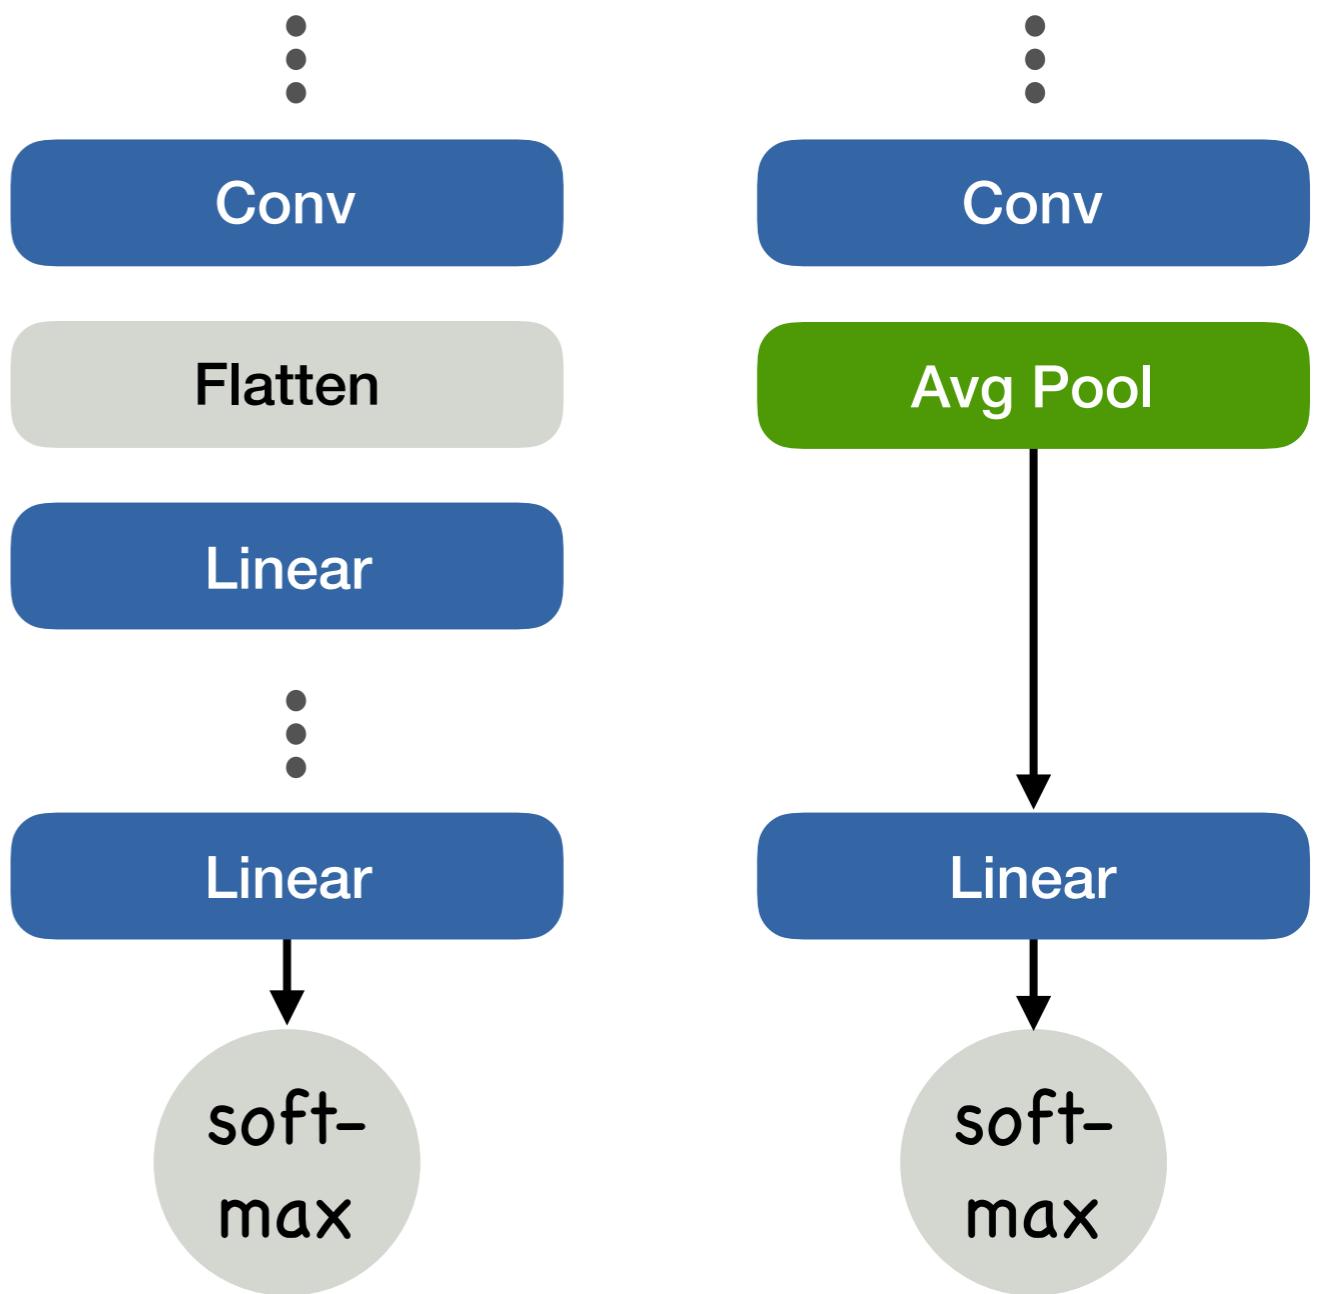
Repeat patterns

- First layer or two are special and not repeated
- All others usually follow a fixed pattern



All-convolutional

- Average in the end
- Fewer parameters
- Better training signal
- “Ensemble”/voting effect for testing



January 23, 2024

```
[1]: %pylab inline
import torch
import sys
sys.path.append('..')
sys.path.append('../..')
from data import load
train_data, train_label = load.get_dogs_and_cats_data(resize=(128,128), n_images=10)
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
print('device = ', device)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
Populating the interactive namespace from numpy and matplotlib
device = cuda

```
[2]: class ConvNet(torch.nn.Module):
    class Block(torch.nn.Module):
        def __init__(self, n_input, n_output, stride=1):
            super().__init__()
            self.net = torch.nn.Sequential(
                torch.nn.Conv2d(n_input, n_output, kernel_size=3, padding=1, stride=stride),
                torch.nn.ReLU(),
                torch.nn.Conv2d(n_output, n_output, kernel_size=3, padding=1),
                torch.nn.ReLU()
            )

        def forward(self, x):
            return self.net(x)

    def __init__(self, layers=[32,64,128], n_input_channels=3):
        super().__init__()
        L = [torch.nn.Conv2d(n_input_channels, 32, kernel_size=7, padding=3, stride=2),
             torch.nn.ReLU(),
             torch.nn.MaxPool2d(kernel_size=3, stride=2, padding=1)]
```

```
c = 32
for l in layers:
    L.append(self.Block(c, l, stride=2))
    c = l
self.network = torch.nn.Sequential(*L)
self.classifier = torch.nn.Linear(c, 1)

def forward(self, x):
    # Compute the features
    z = self.network(x)
    # Global average pooling
    z = z.mean(dim=[2,3])
    # Classify
    return self.classifier(z)[:,0]

net = ConvNet()
```

```
[3]: z = net(train_data[:1])
print(z)
```

```
tensor([0.0218], grad_fn=<SelectBackward0>)
```

```
[4]: %load_ext tensorboard
import tempfile
log_dir = tempfile.mkdtemp()
%tensorboard --logdir {log_dir} --reload_interval 1
```

```
<IPython.core.display.HTML object>
```

```
[5]: from util import train
train.train(net, batch_size=128, resize=(128,128), log_dir=log_dir, ↴
device=device, n_epochs=100)
```

```
WARNING:root:loading dataset
WARNING:root:loading done
```

```
0% | 0/100 [00:00<?, ?it/s]
```

```
[ ]:
```

Deep representations and exploiting the structure of the data

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Structure of input data

- Images
 - Repeating patterns
 - at various scales



Structure of convolutional networks

- Exploit repeating structure of images

inception_4e unit 789



inception_4e unit 750



inception_5b unit 626



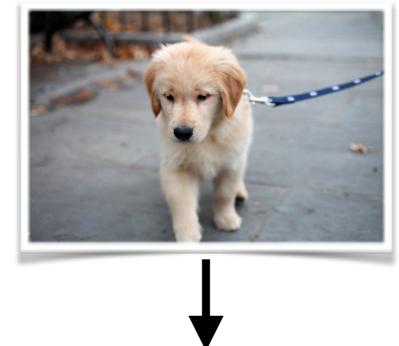
inception_4e unit 175



inception_4e unit 225



inception_5b unit 415

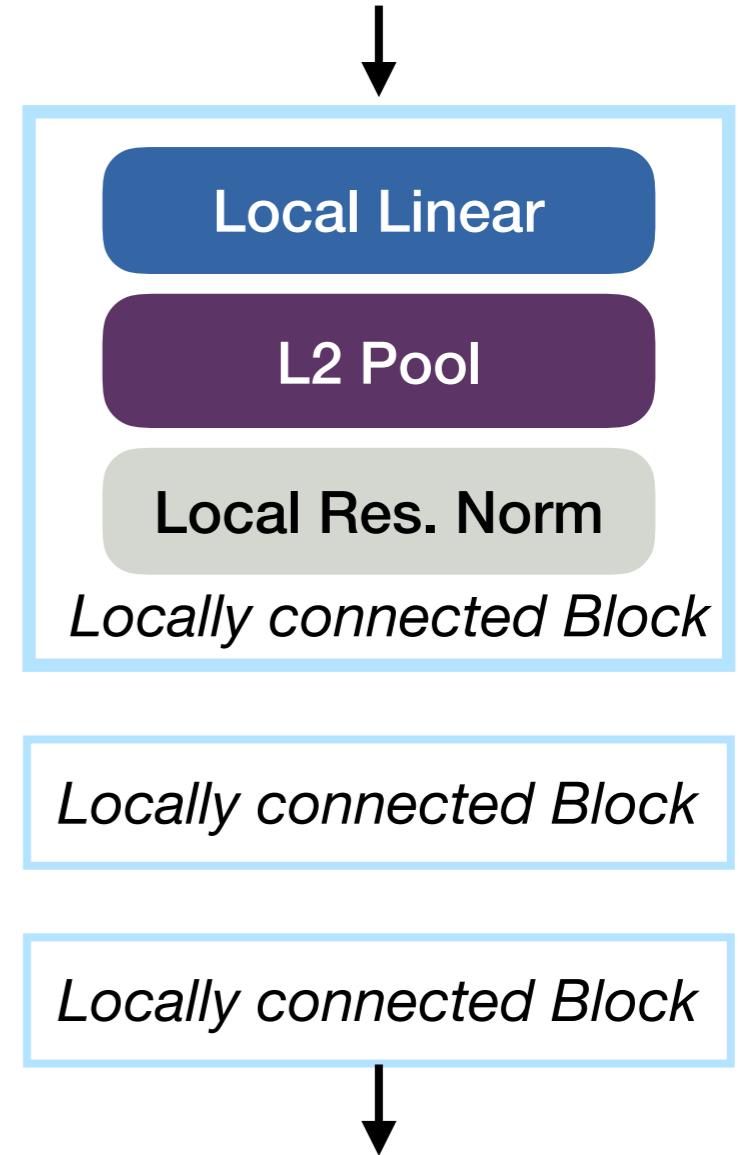


What do networks learn?



Linear layered networks do not work well on images

- Largest linear network for computer vision
 - Locally connected



January 23, 2024

```
[1]: %pylab inline

import json
import numpy as np
from PIL import Image

import torch
import torchvision.models as models
import torchvision.transforms as transforms
device = torch.device('cuda') if torch.cuda.is_available() else torch.
    device('cpu')
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
Populating the interactive namespace from numpy and matplotlib

```
[ ]: model = models.resnet18(pretrained=True)
model.eval()
class_idx = json.load(open("network_vis/imagenet_class_index.json"))
```

```
[4]: transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))
```

```
[5]: image = transform(Image.open('network_vis/dog.jpg'))
plt.imshow(plt.imread('network_vis/dog.jpg'))

predictions = model(image.unsqueeze(0)).detach().numpy()[0]

print('Ranking: class (score)')
for i, idx in enumerate(np.argsort(predictions)[-10:][::-1]):
    print('%d: %s (%.04f)' % (
        i,
        class_idx[str(idx)][1],
        predictions[idx]))
```

```
Ranking: class (score)
0: wire-haired_fox_terrier (10.5043)
1: toy_terrier (9.4139)
2: English_foxhound (9.1540)
3: Ibizan_hound (8.7189)
4: Chihuahua (8.0401)
5: Lakeland_terrier (7.9056)
6: Border_terrier (7.8533)
7: Brittany_spaniel (7.7991)
8: Walker_hound (7.7619)
9: basenji (7.7018)
```

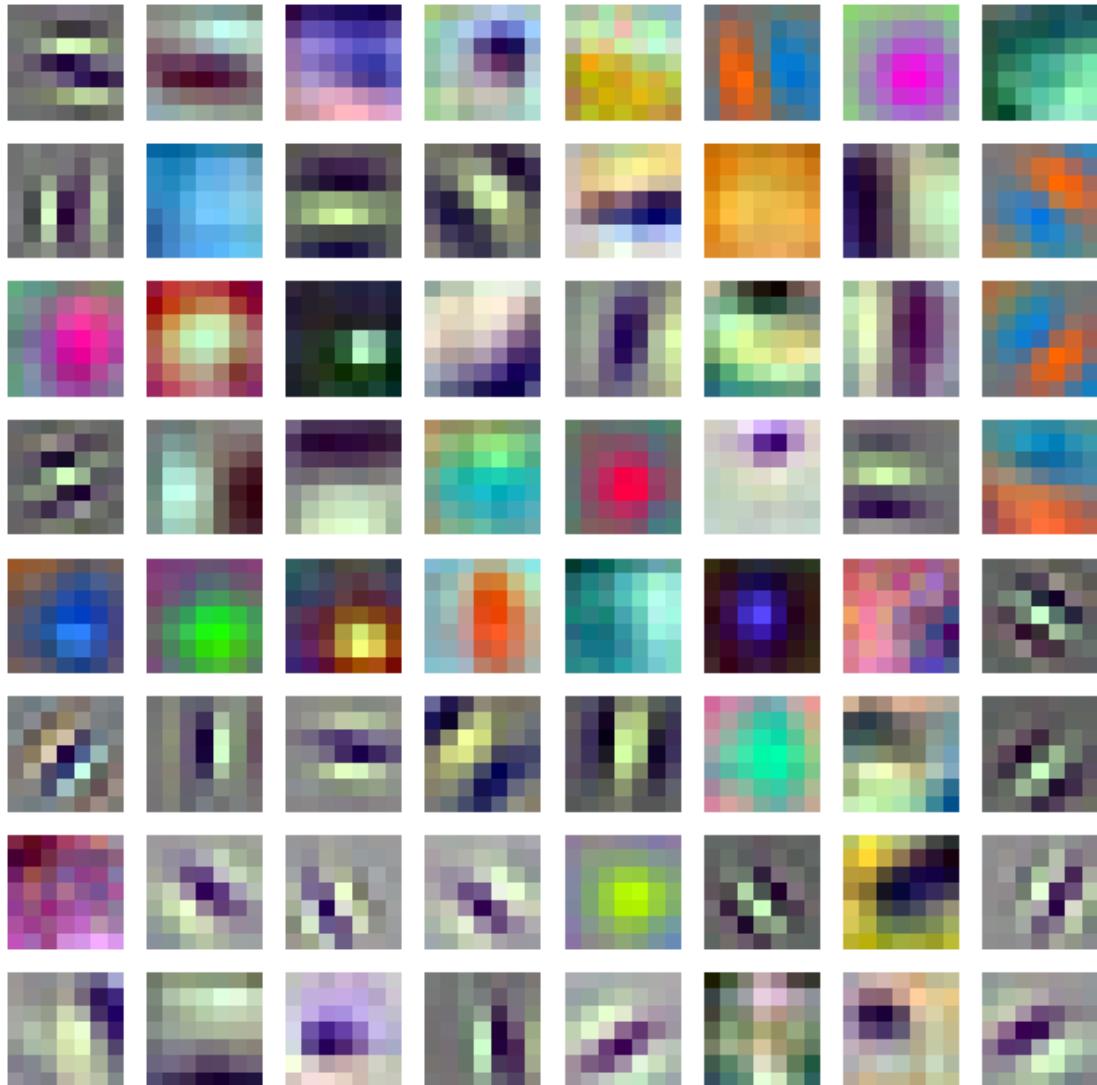


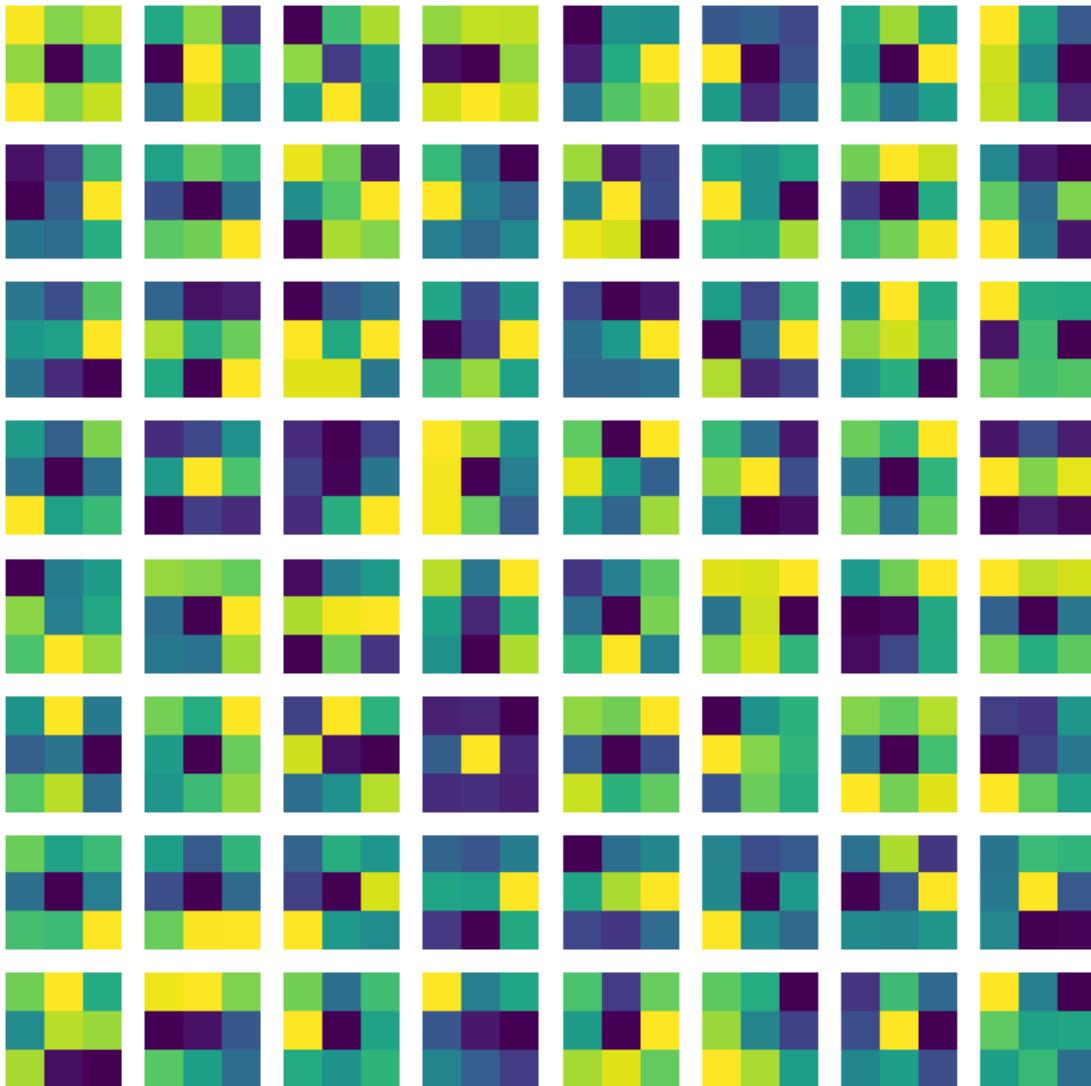
```
[6]: def visualize_layer(weight):
    fig=plt.figure(figsize=(8, 8))

    for i in range(64):
        x = weight[i, ...].transpose([1, 2, 0])
        x = (x - np.min(x))/np.ptp(x)

        fig.add_subplot(8, 8, i + 1)
        if x.shape[2] == 3:
            imshow(x)
```

```
else:  
    imshow(x[:, :, 0])  
    axis('off')  
  
show()  
  
visualize_layer(model.conv1.weight.data.cpu().numpy())  
visualize_layer(model.layer1[0].conv1.weight.data.cpu().numpy())
```





```
[7]: import sys
sys.path.append('..')
sys.path.append('../..')
from data import load
valid_data = load.get_dogs_and_cats('valid', resize=None, batch_size=128, ↴
is_resnet=True)
```

```
[8]: def undo_transform(x):
    mean = np.array([0.485, 0.456, 0.406])[np.newaxis, :, np.newaxis, np.
    ↴newaxis]
    std = np.array([0.229, 0.224, 0.225])[np.newaxis, :, np.newaxis, np.newaxis]
    return (x * std) + mean
```

```
[9]: all_data = []
C1 = []
C2 = []
C3 = []

def store_activation(L, m, grad_in, grad_out):
    L.append(grad_out.detach().cpu().numpy())

model.to(device)

try:
    h1 = model.conv1.register_forward_hook(lambda *args: store_activation(C1, *args))
    h2 = model.layer1[0].conv1.register_forward_hook(lambda *args: store_activation(C2, *args))
    h3 = model.layer2[0].conv1.register_forward_hook(lambda *args: store_activation(C3, *args))

    for it, (data, label) in enumerate(valid_data):
        all_data.append(undo_transform(data.numpy()))
        if device is not None:
            data, label = data.to(device), label.to(device)
            result = model(data)
finally:
    h1.remove()
    h2.remove()
    h3.remove()

all_data = np.vstack(all_data)
C1 = np.vstack(C1)
C2 = np.vstack(C2)
C3 = np.vstack(C3)
```

```
[10]: def to_pil(x):
        x = x.transpose((1, 2, 0))
        x -= x.min(0).min(0)[None, None]
        x /= x.max(0).max(0)[None, None]
        return Image.fromarray((x * 255).astype('uint8'))

def viz_act(img, act):

    num_acts = len(act)
    img = to_pil(img)

    fig=plt.figure(figsize=(20, 20))

    fig.add_subplot(1, 1+num_acts, 1)
```

```

imshow(img)
axis('off')

for i, c in enumerate(act):
    fig.add_subplot(1, 1+num_acts, i + 2)
    imshow(c)
    axis('off')

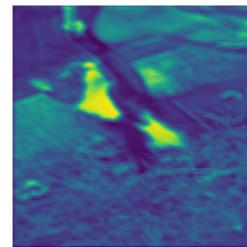
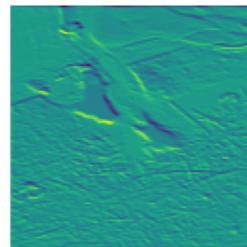
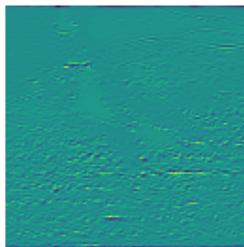
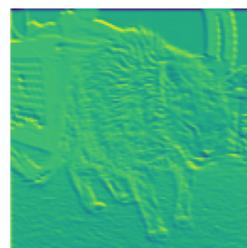
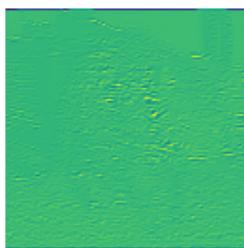
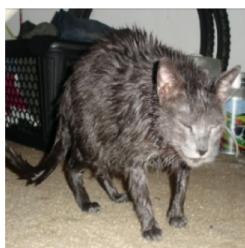
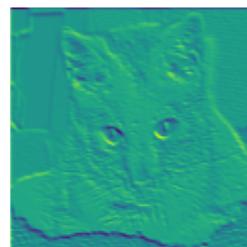
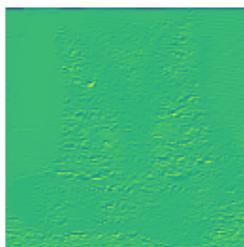
show()

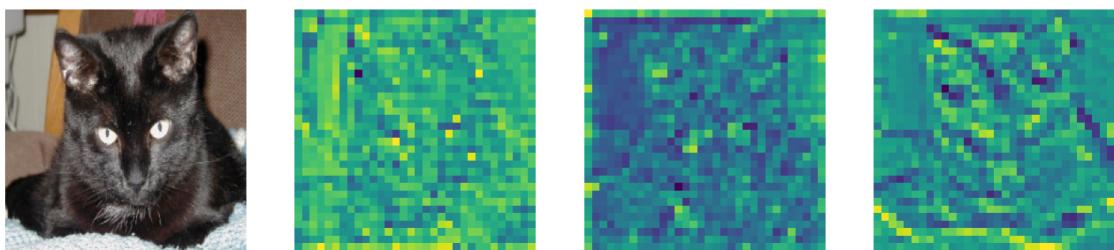
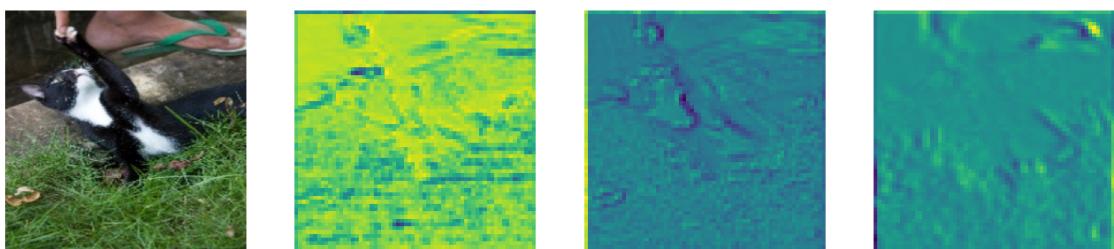
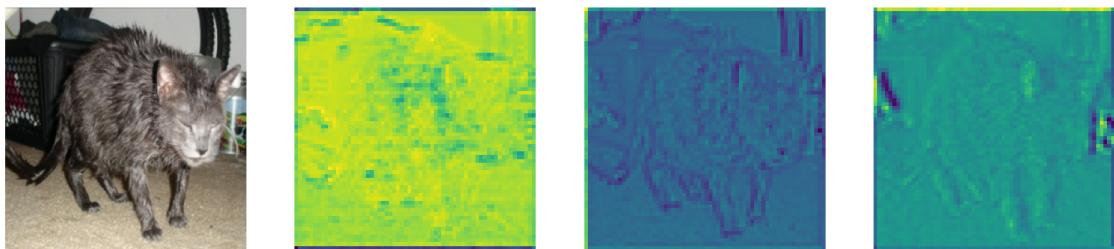
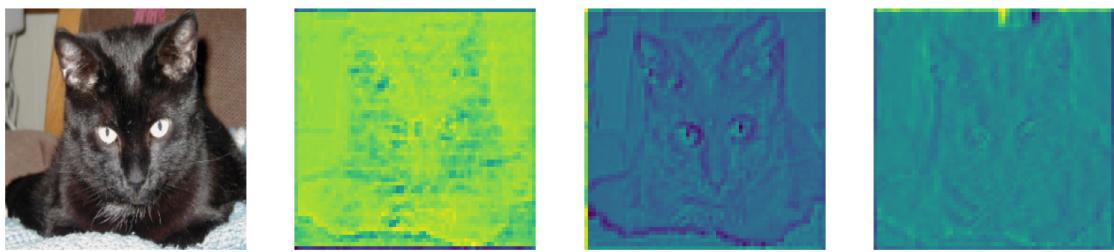
for i in range(3):
    viz_act(all_data[i], C1[i, :3])

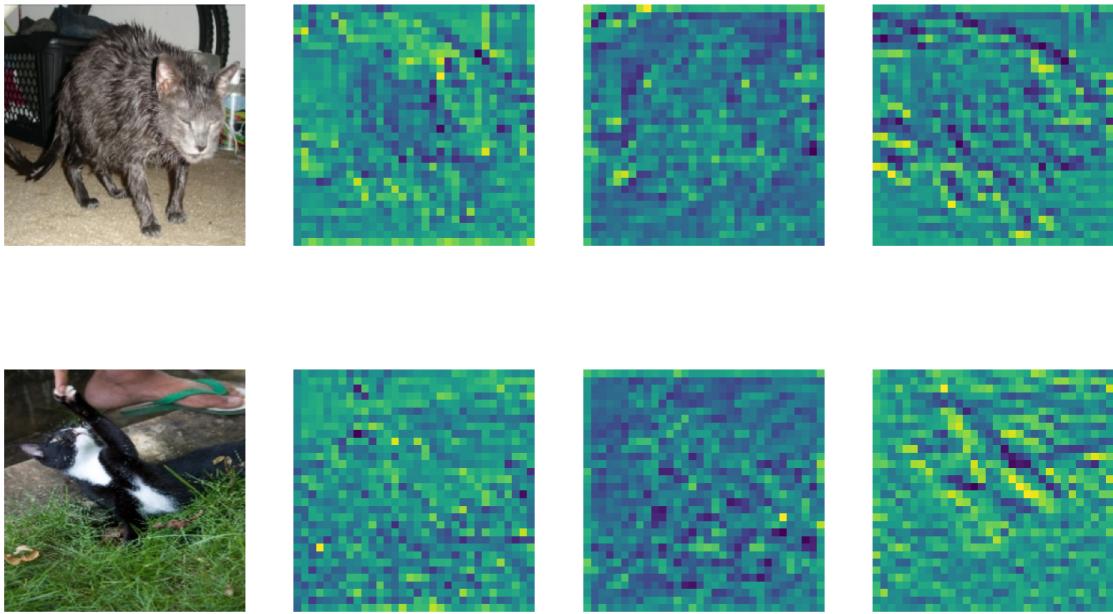
for i in range(3):
    viz_act(all_data[i], C2[i, :3])

for i in range(3):
    viz_act(all_data[i], C3[i, :3])

```







```
[11]: def find_max_act(C, filter, W=10, img_size=[224, 224]):
    fig=plt.figure(figsize=(8, 8))

    filter_act = C[:, filter]
    for i, idx in enumerate(np.argsort(-filter_act, axis=None)[:64]):
        max_pos = list(np.unravel_index(idx, filter_act.shape))

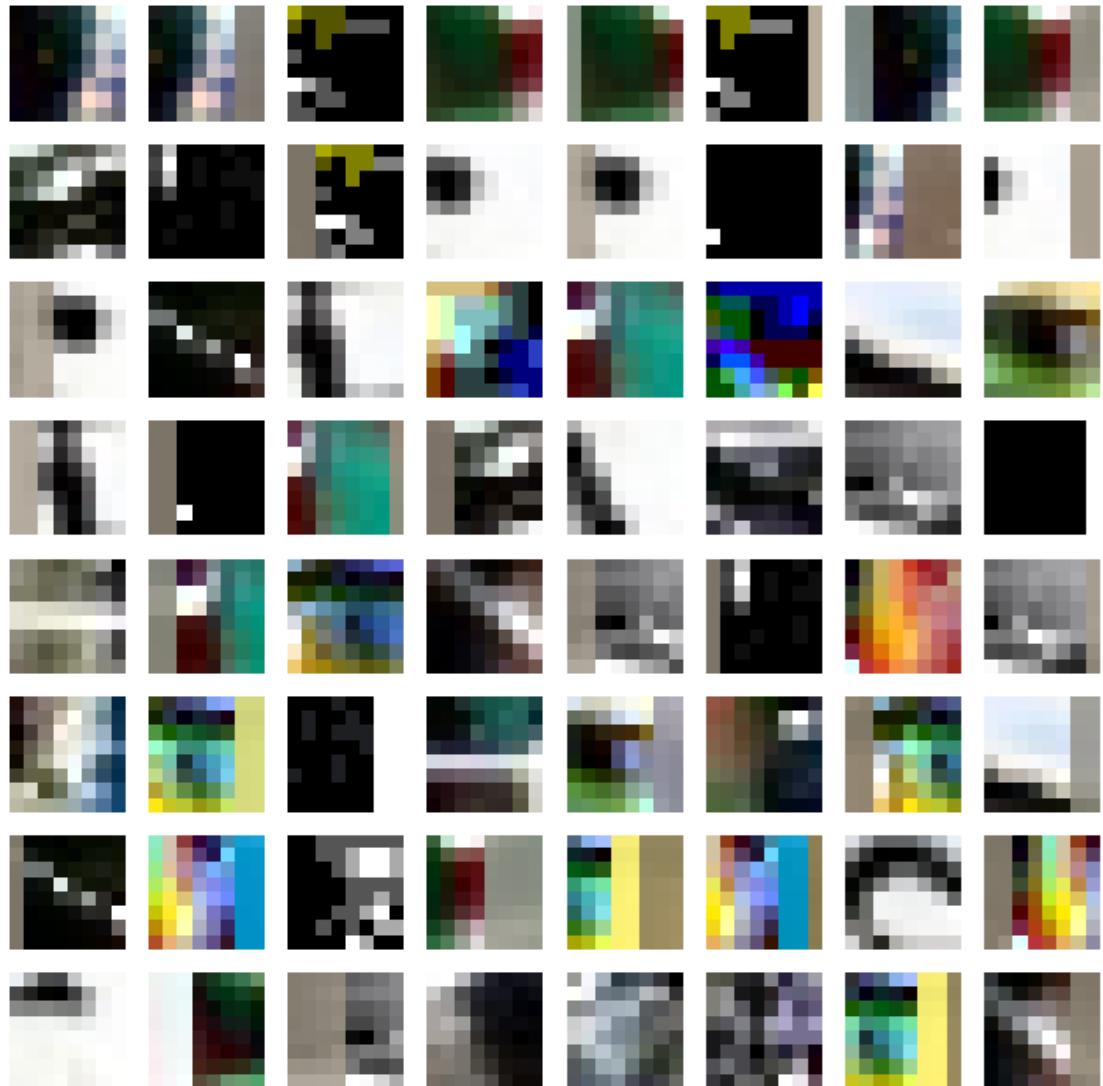
        max_pos[1] = int(max_pos[1] * float(img_size[0]) / filter_act.shape[1])
        max_pos[2] = int(max_pos[2] * float(img_size[1]) / filter_act.shape[2])

        img = to_pil(all_data[max_pos[0]],
                    :,
                    max(max_pos[1] - W, 0) : max_pos[1] + W,
                    max(max_pos[2] - W, 0) : max_pos[2] + W)

        fig.add_subplot(8, 8, i + 1)
        imshow(img)
        axis('off')

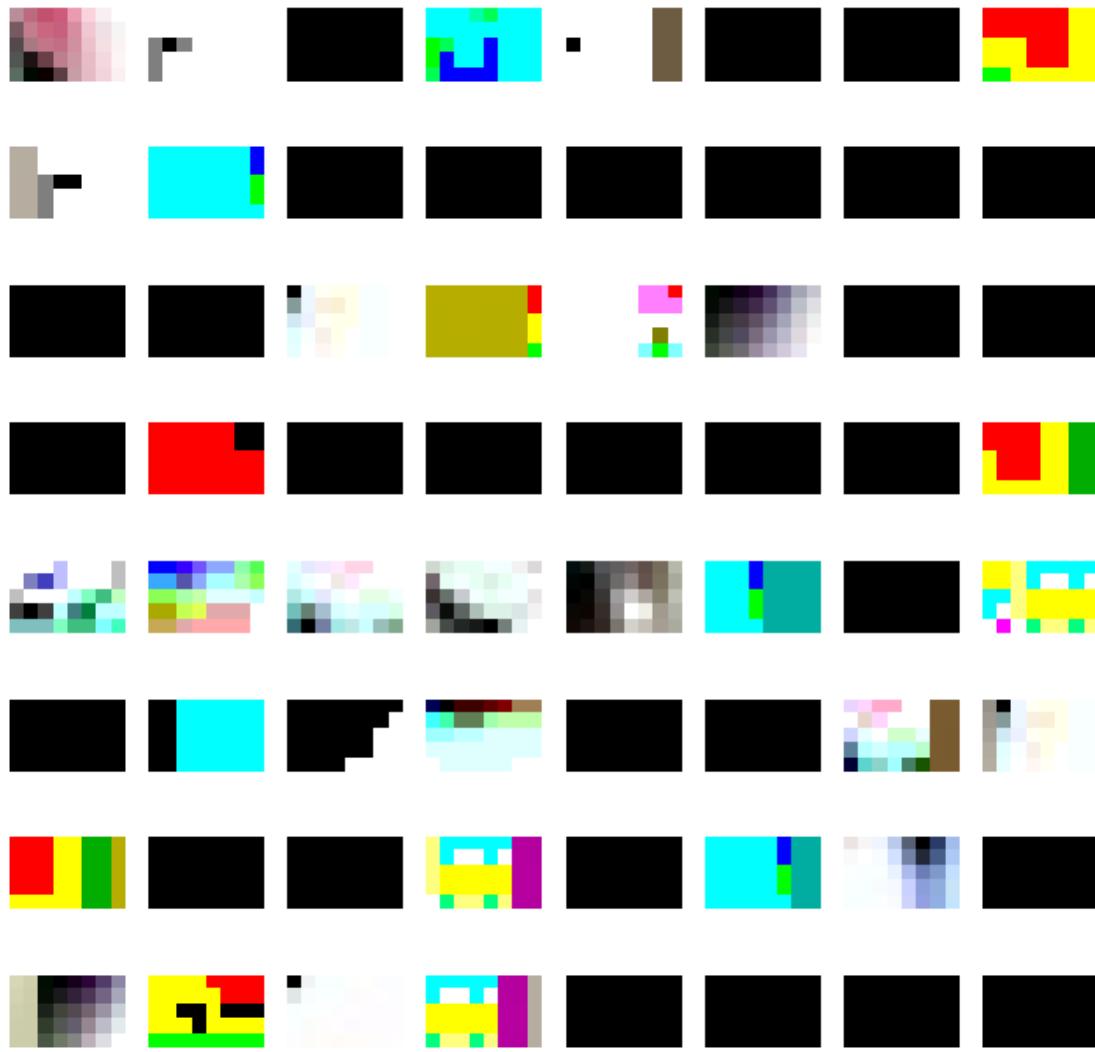
    show()

find_max_act(C1, filter=0, W=4)#
find_max_act(C1, filter=1, W=4)#
find_max_act(C1, filter=2, W=4)#
img_size=img.size)
```



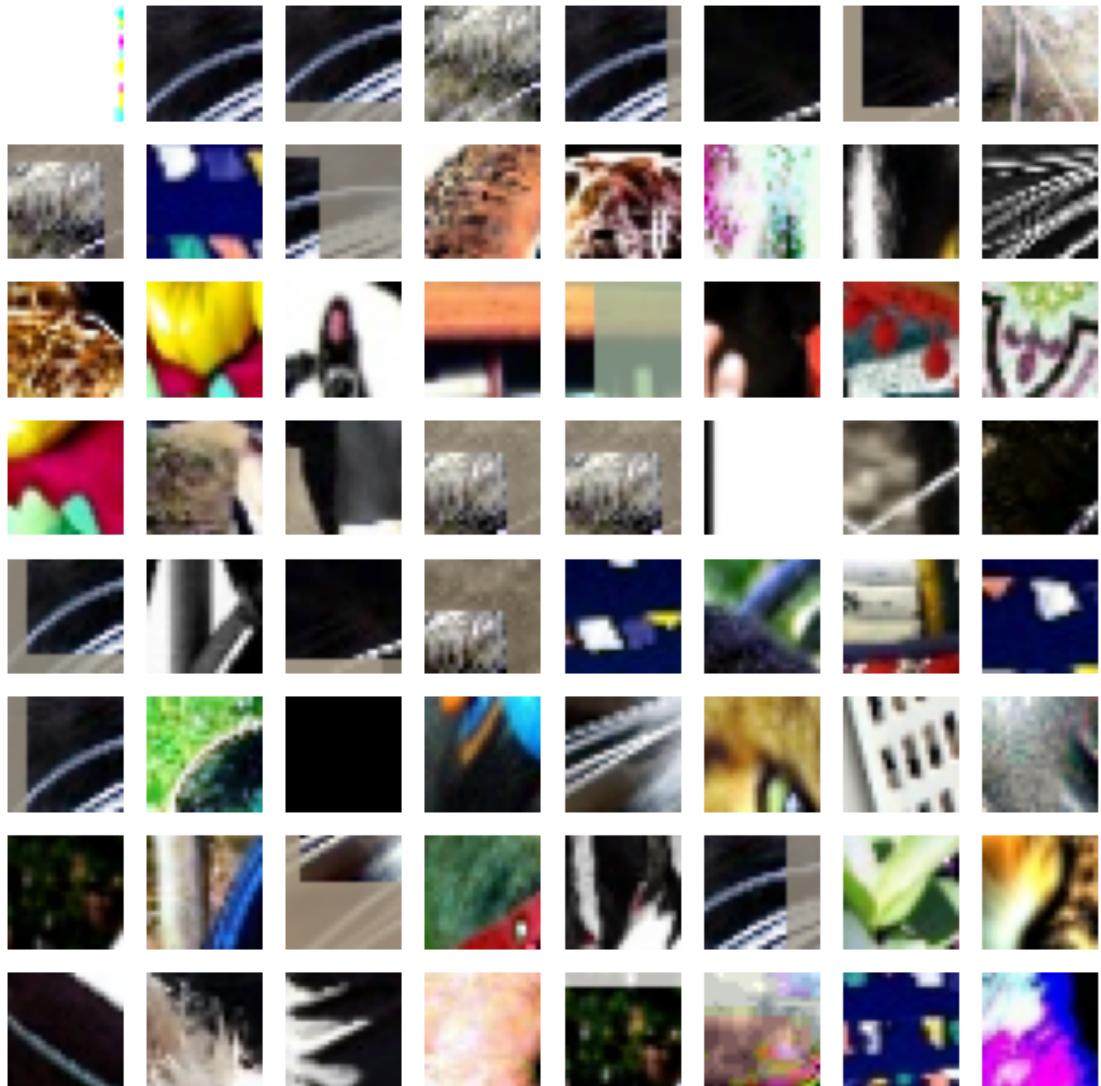


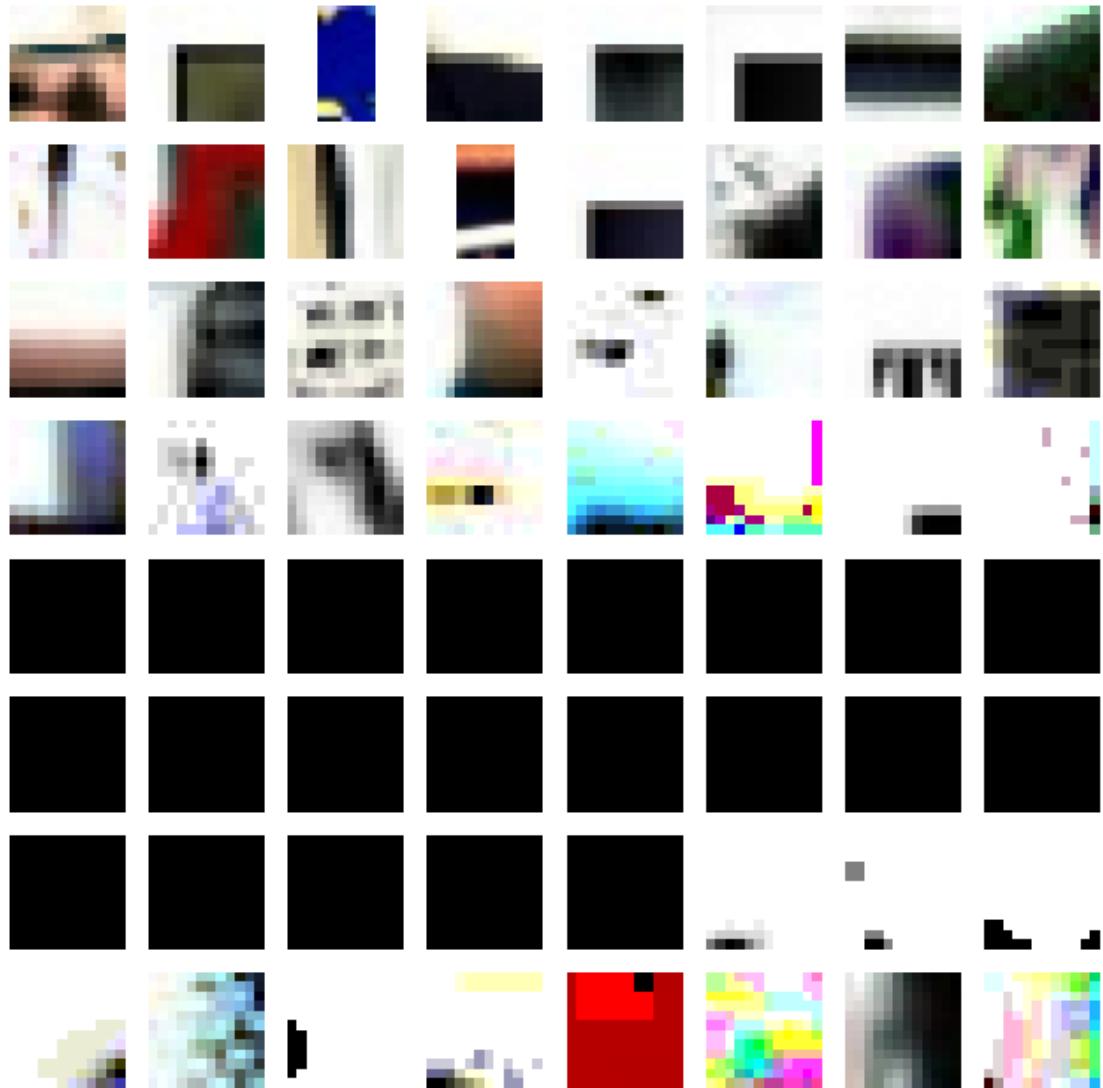
```
/tmp/ipykernel_2833/1084252010.py:4: RuntimeWarning: invalid value encountered  
in divide  
    x /= x.max(0).max(0)[None,None]  
/tmp/ipykernel_2833/1084252010.py:5: RuntimeWarning: invalid value encountered  
in cast  
    return Image.fromarray((x * 255).astype('uint8'))
```

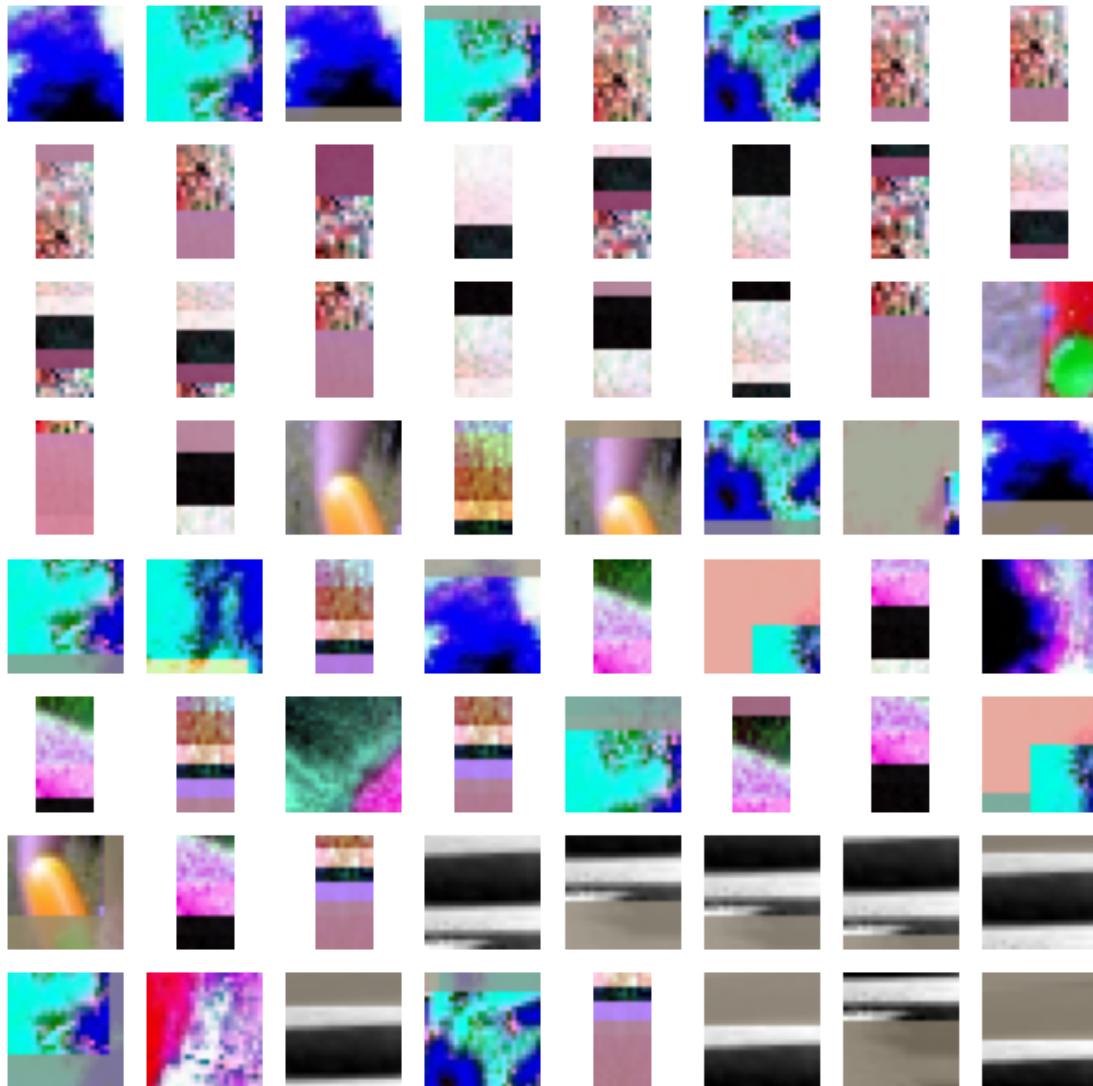


```
[12]: find_max_act(C2, filter=0, W=12) #, img_size=img.size)
find_max_act(C2, filter=1, W=12) #, img_size=img.size)
find_max_act(C2, filter=2, W=12) #, img_size=img.size)
```

```
/tmp/ipykernel_2833/1084252010.py:4: RuntimeWarning: invalid value encountered
in divide
    x /= x.max(0).max(0)[None,None]
/tmp/ipykernel_2833/1084252010.py:5: RuntimeWarning: invalid value encountered
in cast
    return Image.fromarray((x * 255).astype('uint8'))
```

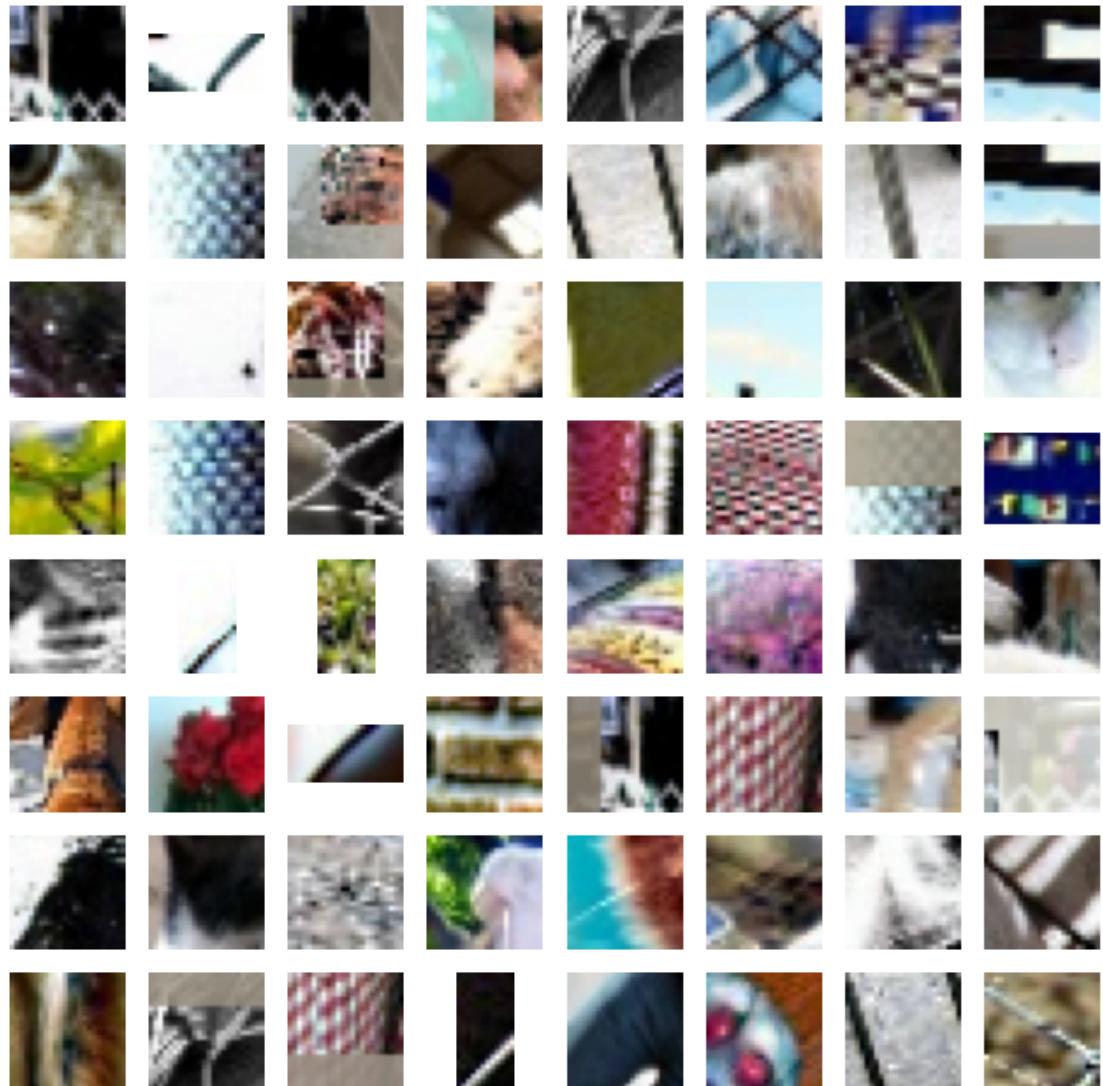


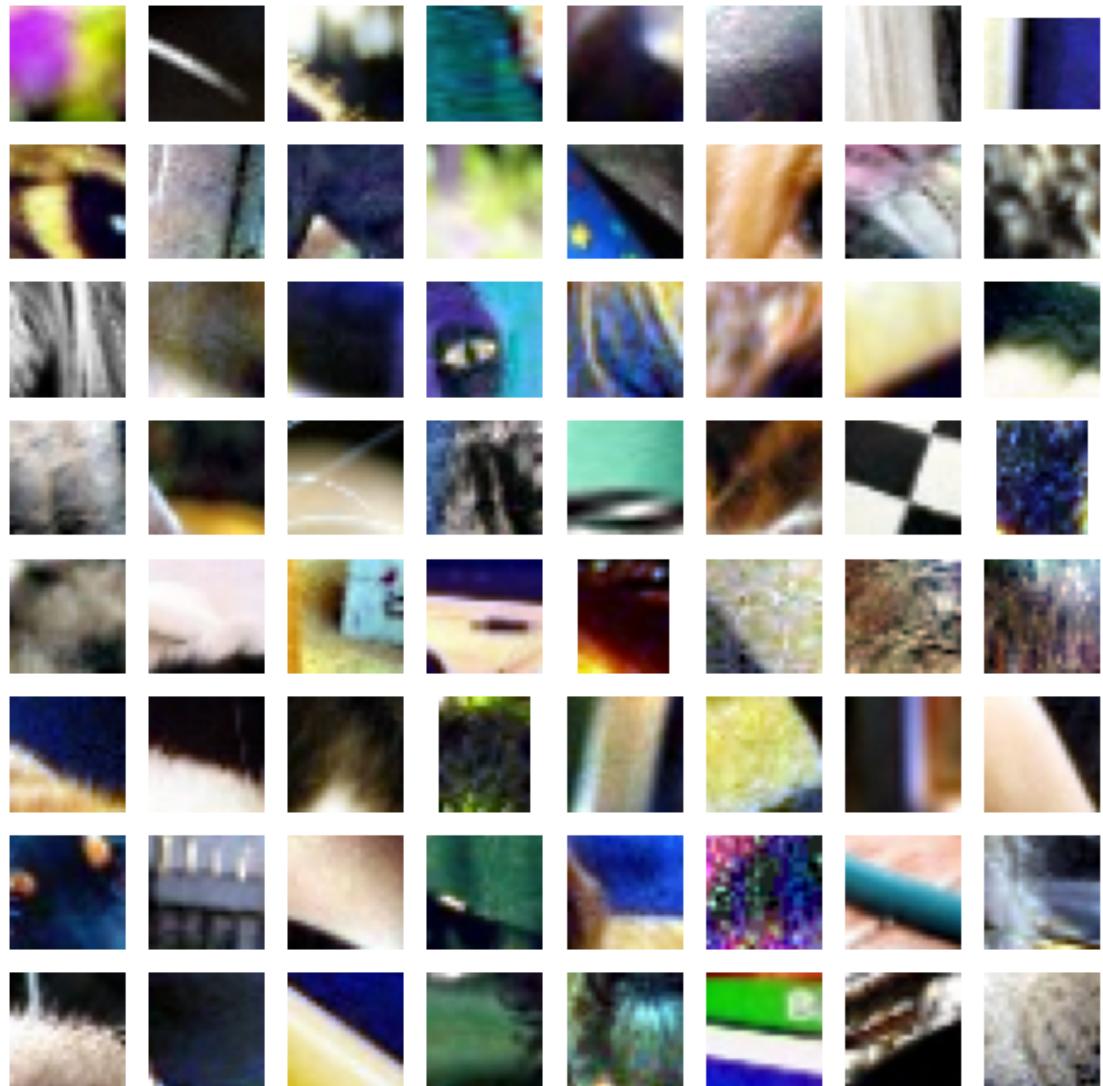




```
[13]: find_max_act(C3, filter=0, W=12)#
find_max_act(C3, filter=1, W=12)#
find_max_act(C3, filter=2, W=12)#
```







[]:

Dilation

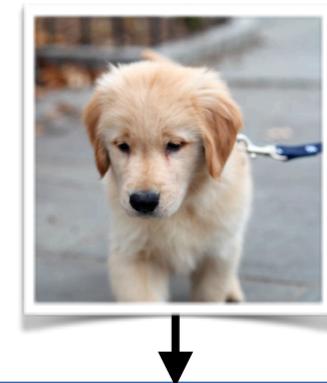
© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Segmentation



Receptive field

- How to increase receptive field?
 - Large kernel size
 - Striding



Conv 3x3

ReLU

Conv 3x3

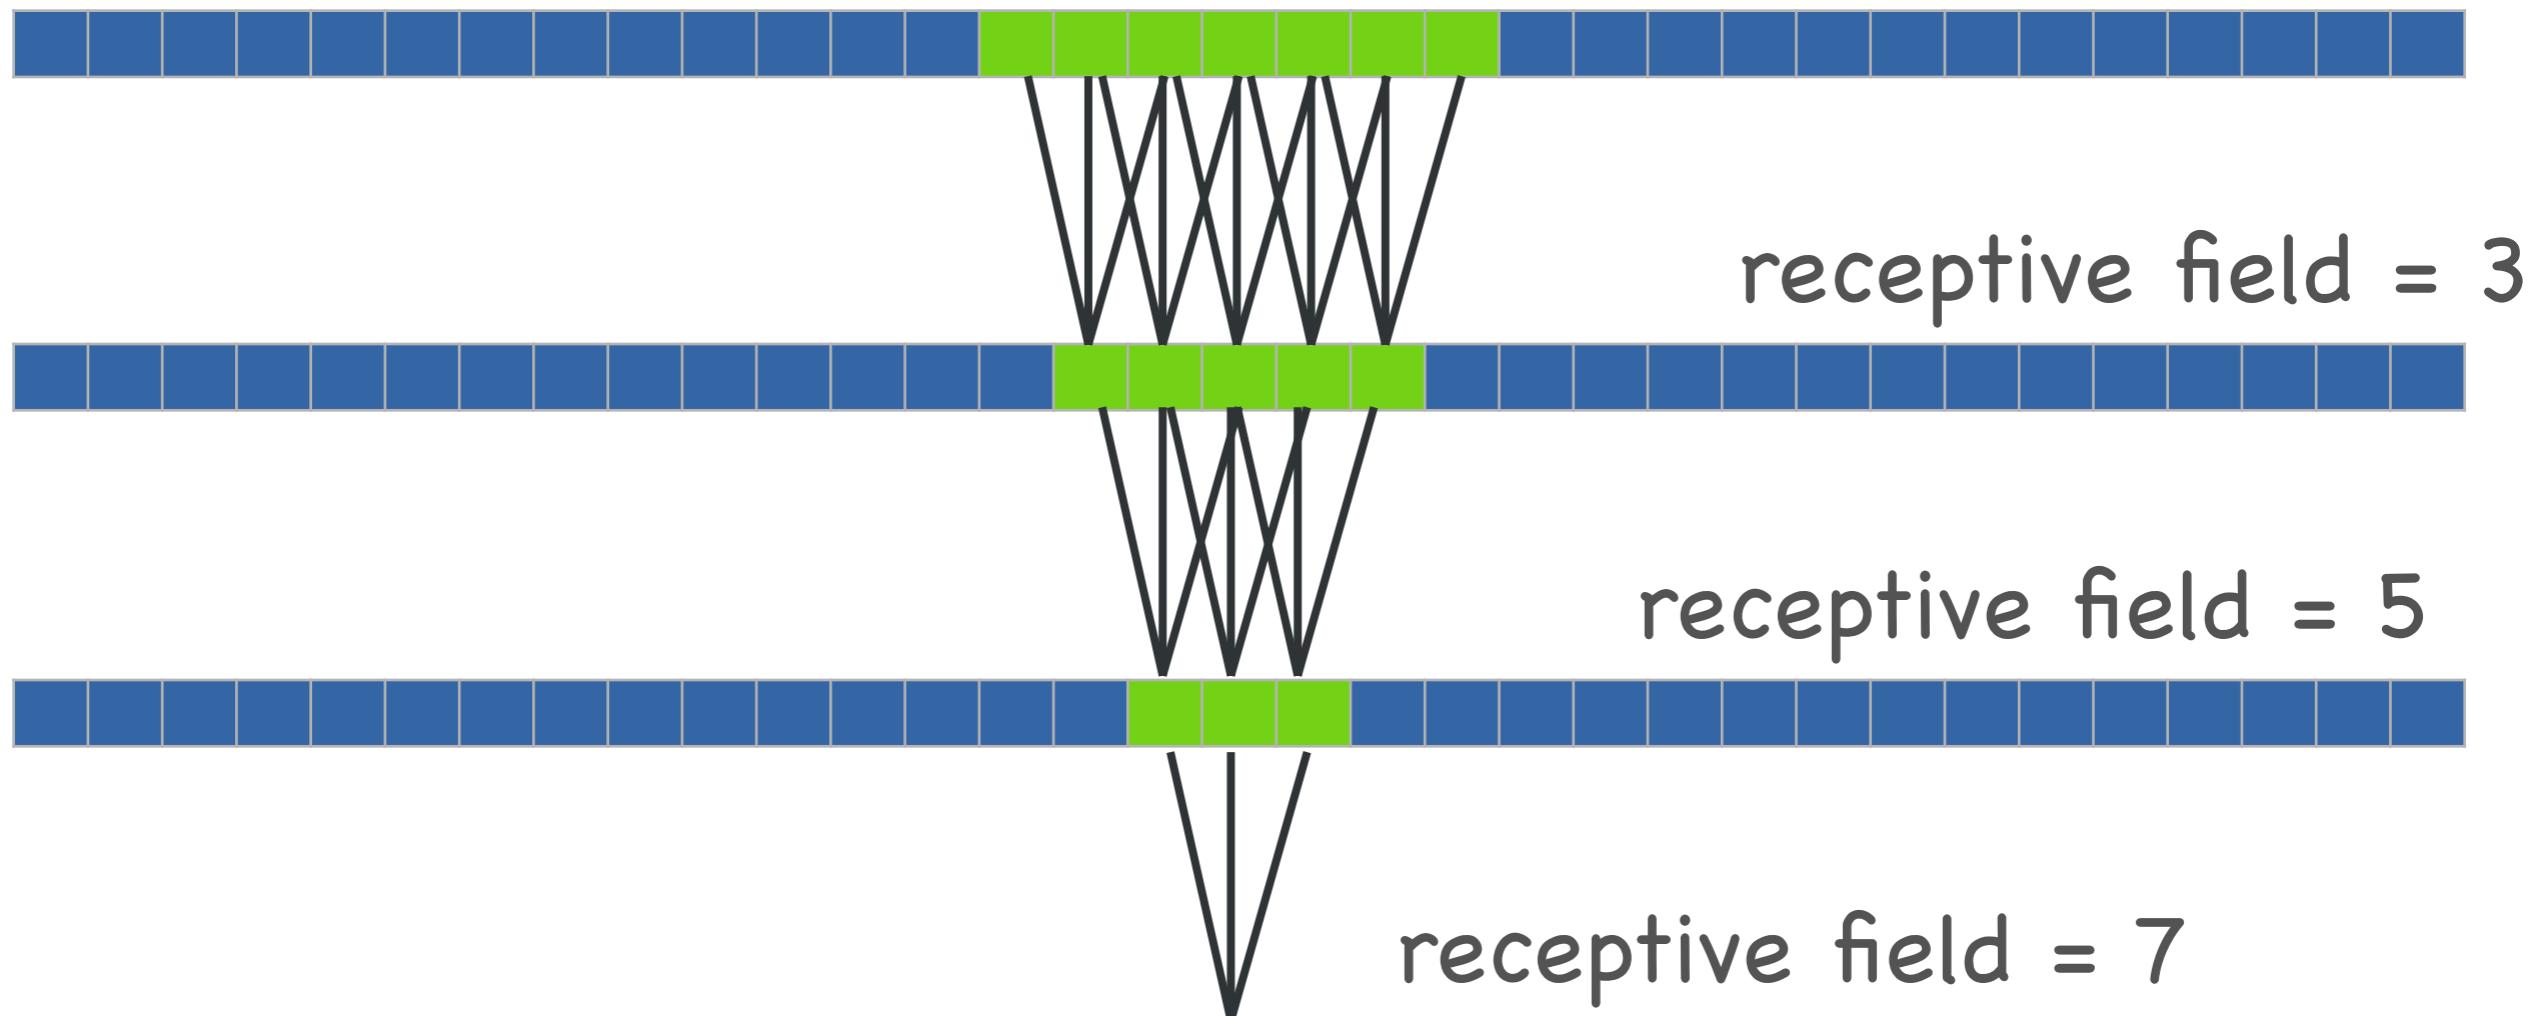
ReLU

Conv 3x3

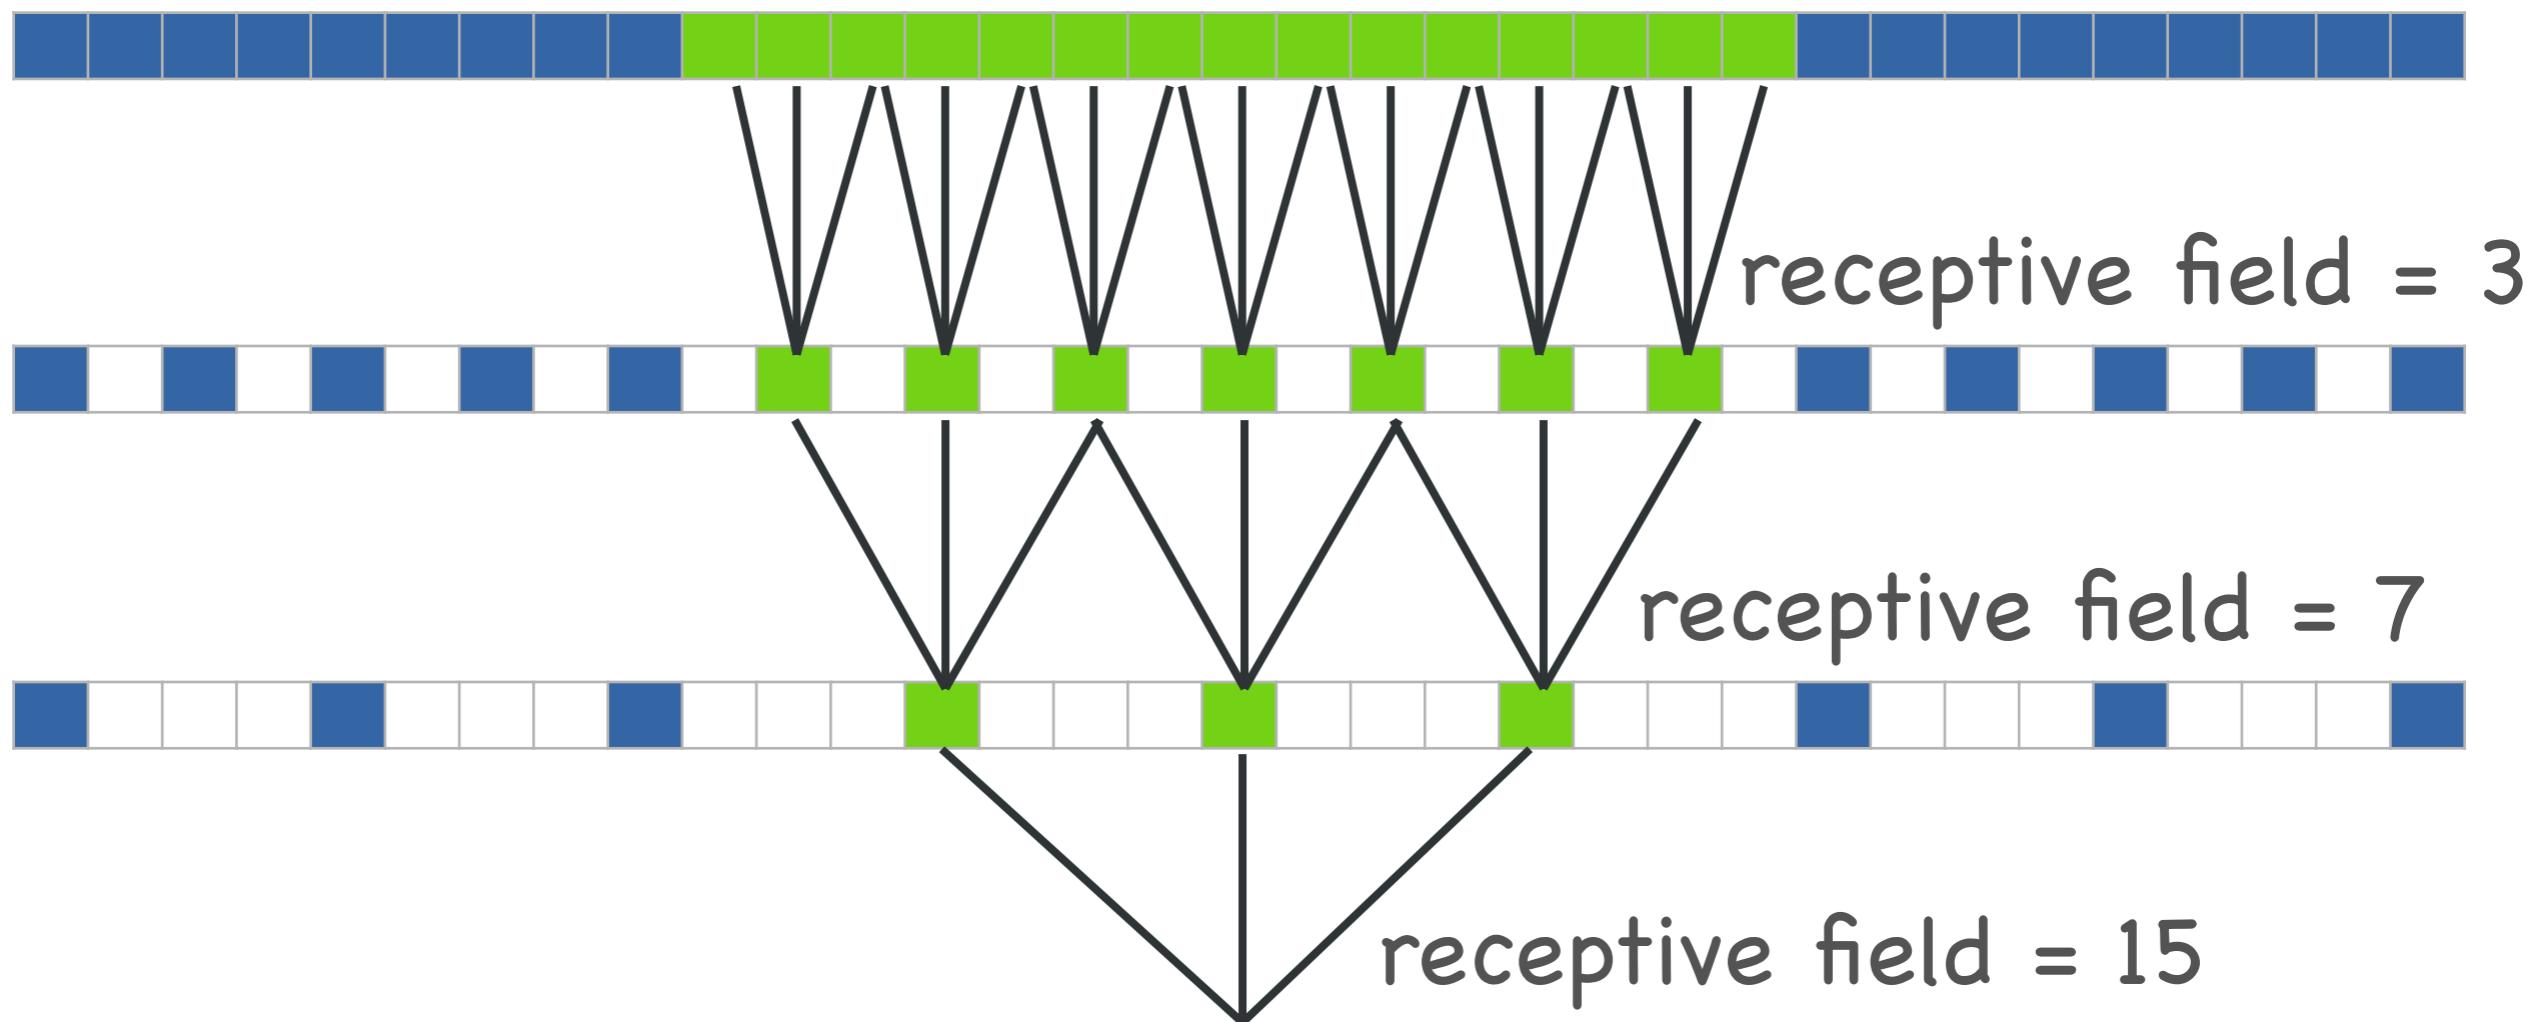
ReLU

⋮

Receptive field

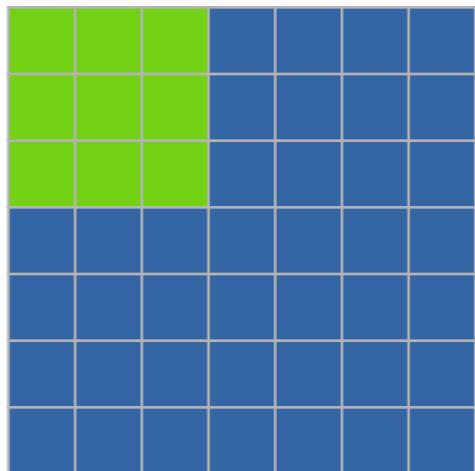


Receptive field and striding



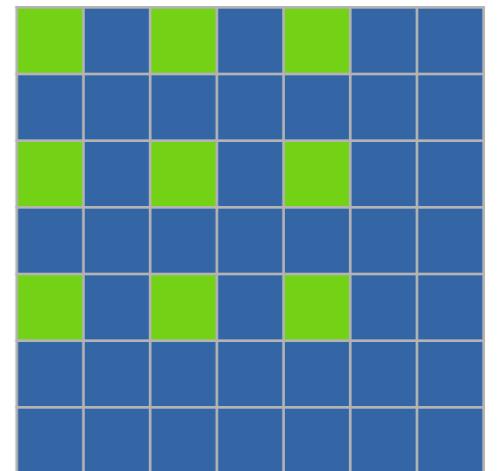
Dilation

- Add 0-padding between values in convolutional kernel



*

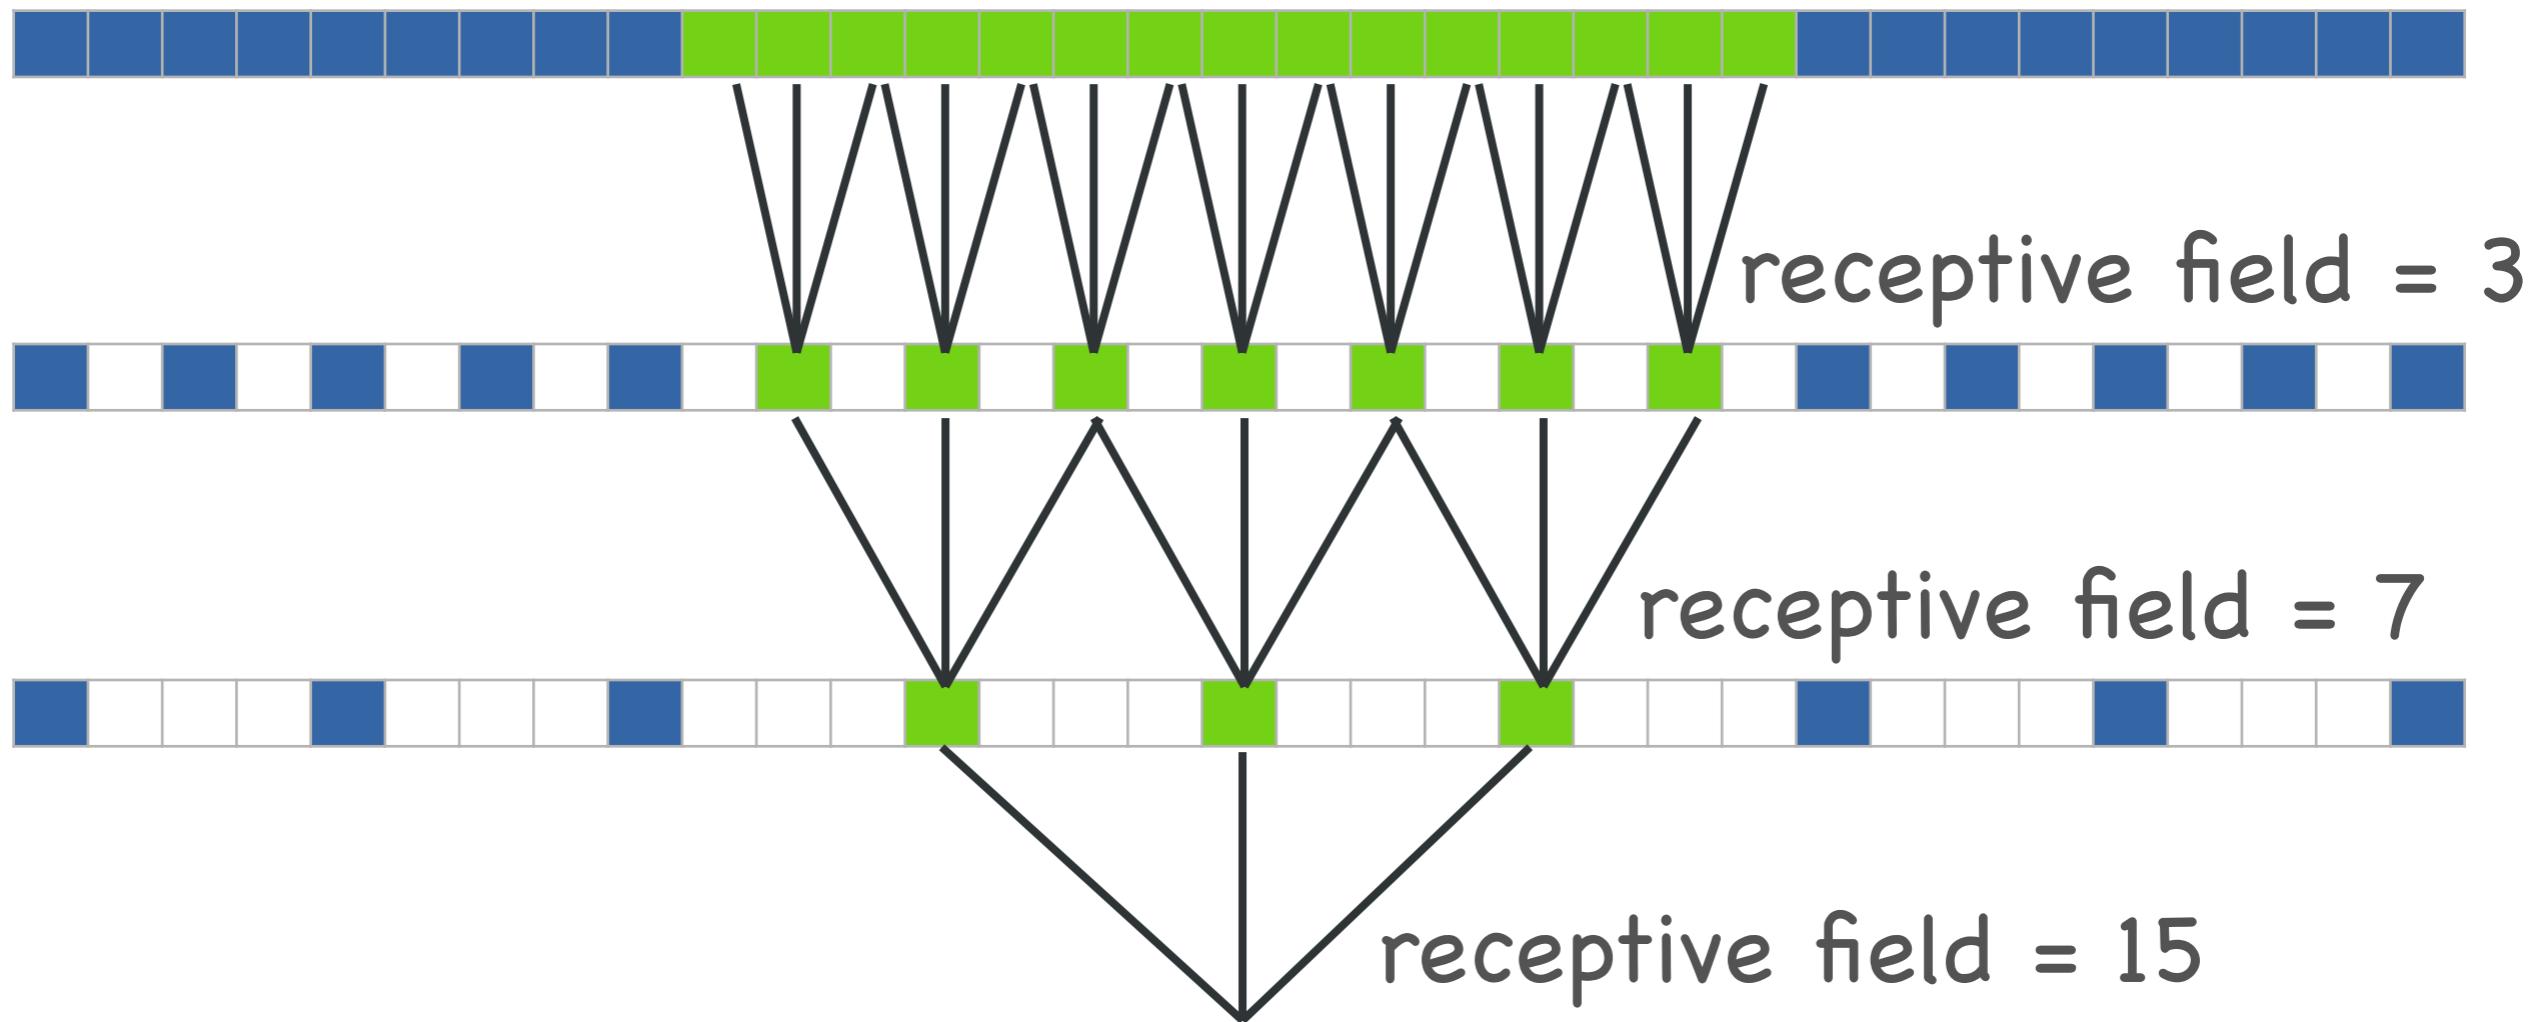
a	b	c
d	e	f
g	h	i



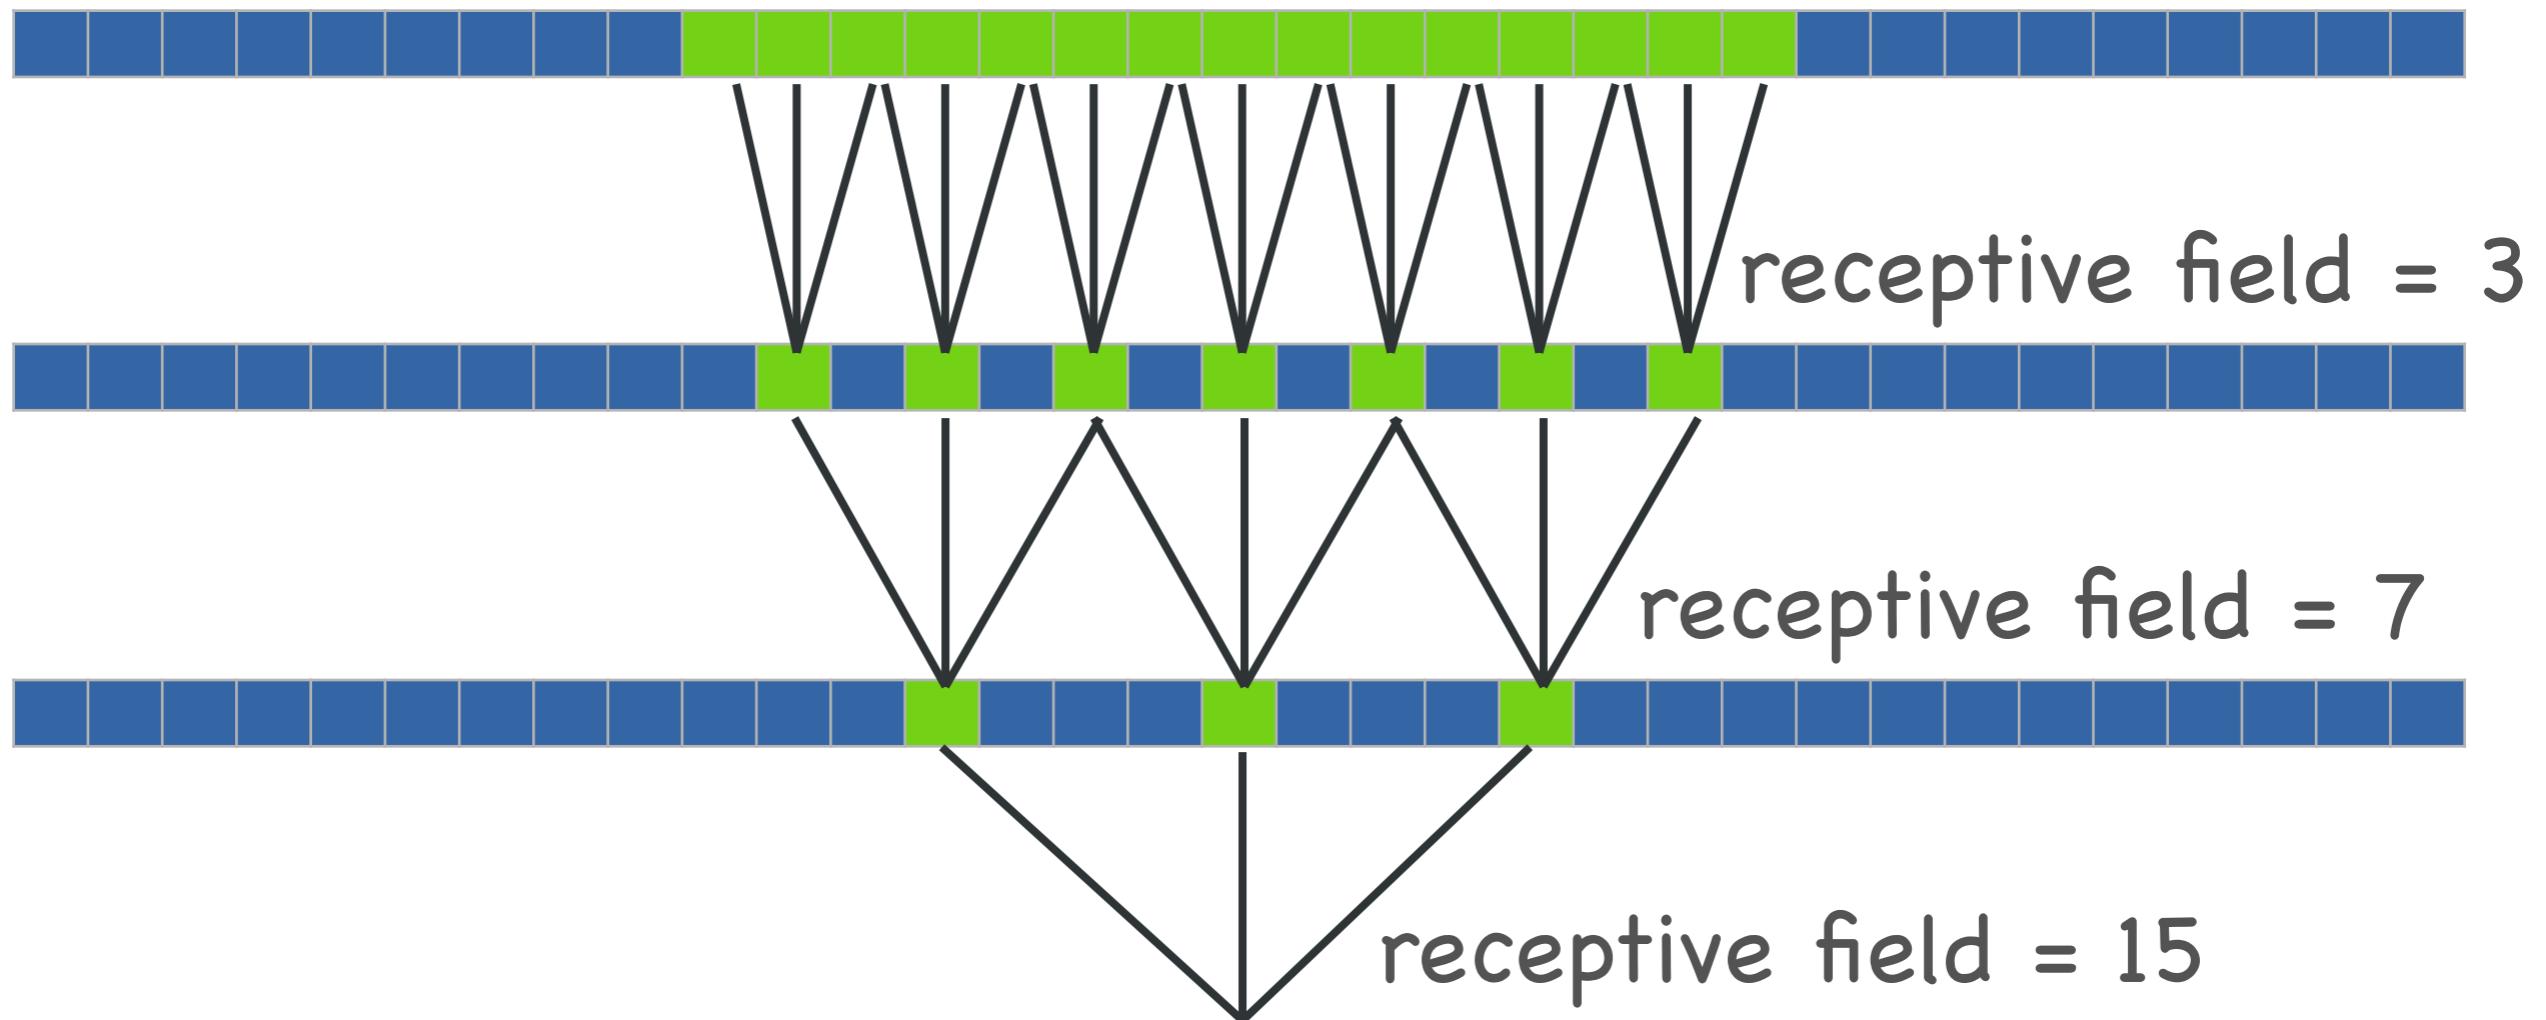
*

a		b	c
d		e	f
g		h	i

Dilation vs striding



Dilation vs striding



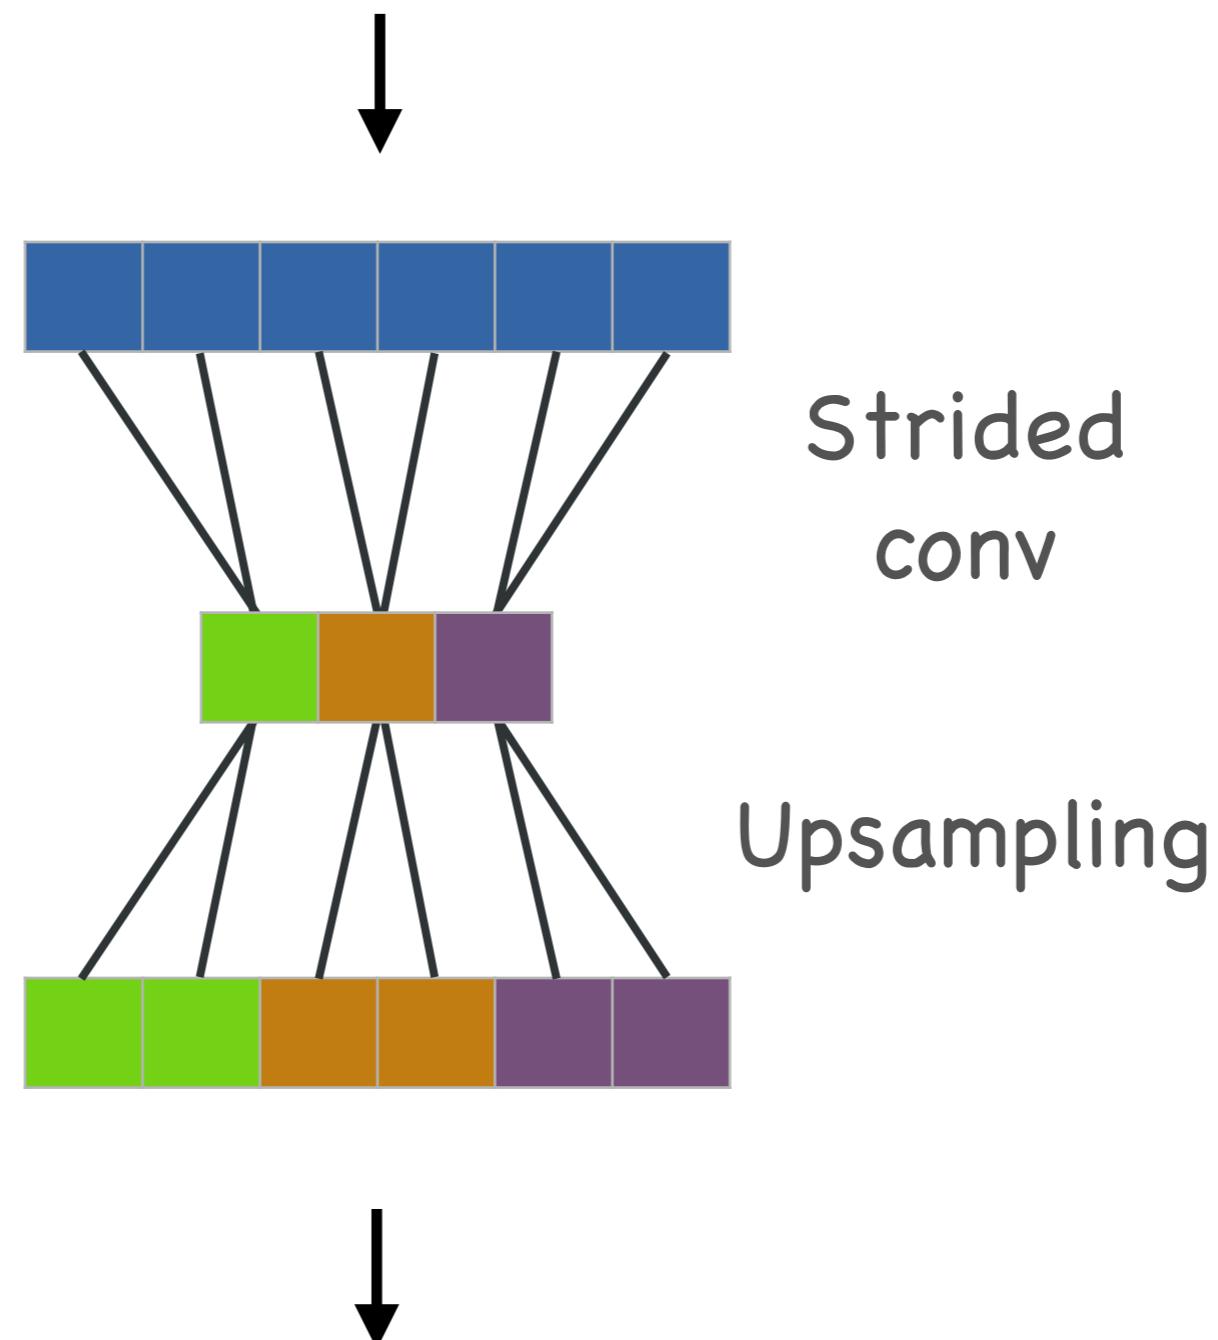
The many names of dilation

- hole
- a trous

Up-convolution

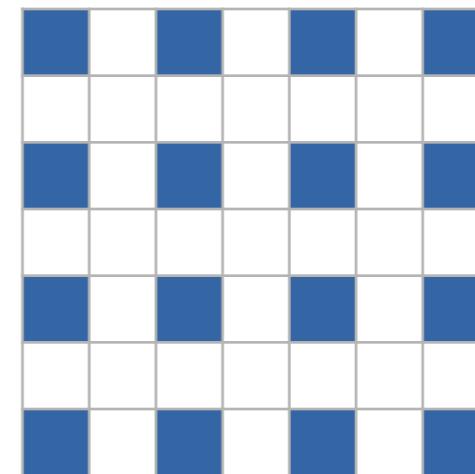
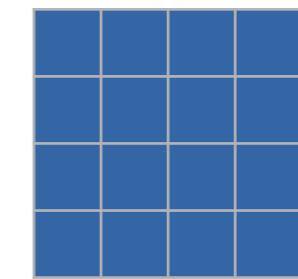
© 2019 Philipp Krähenbühl and Chao-Yuan Wu

Inverse of strided convolution



Up-convolution

- Dilation of the input

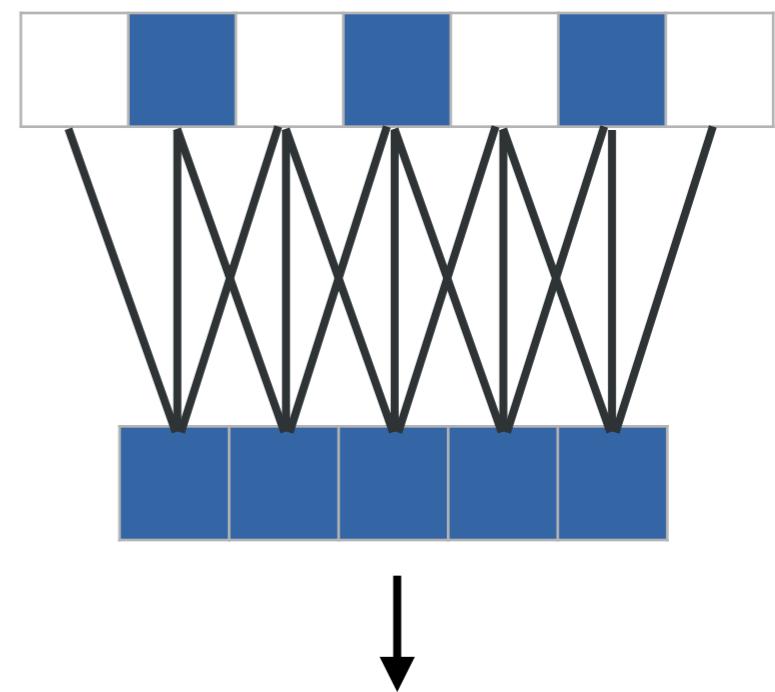
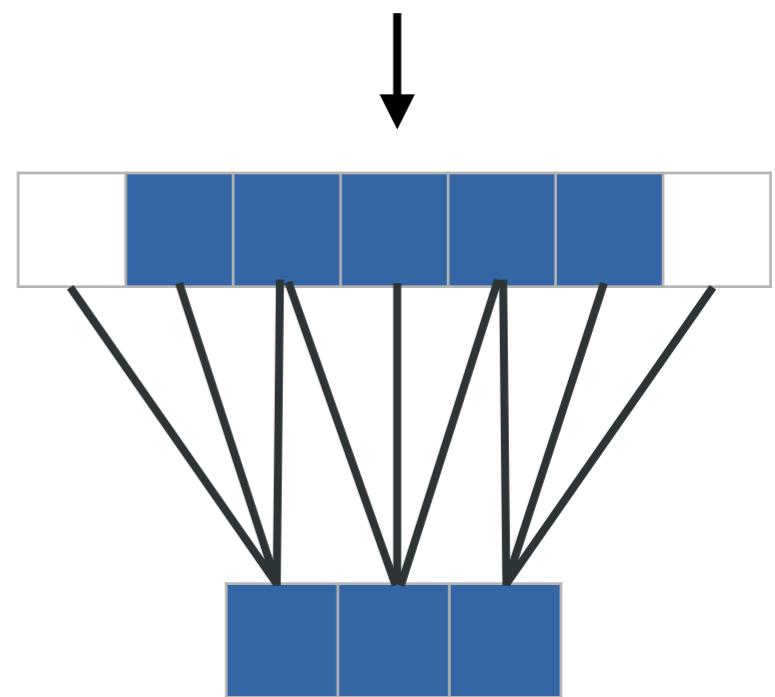


*

a	b	c
d	e	f
g	h	i

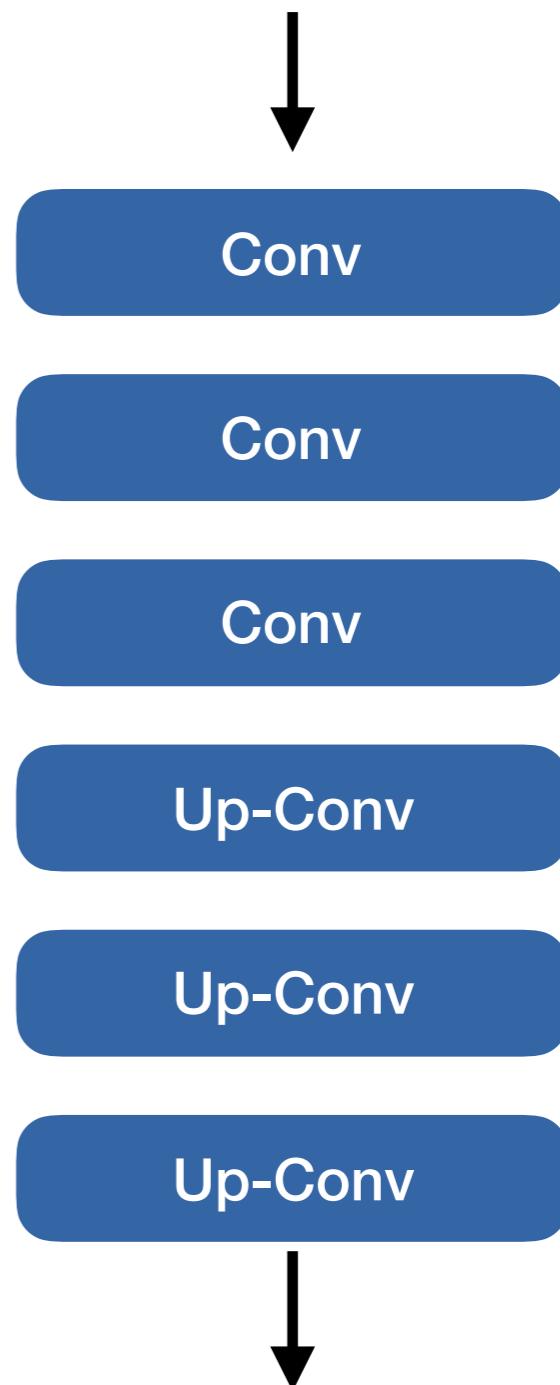
Rounding

- Strided convolution rounds down
- How to correct for this?



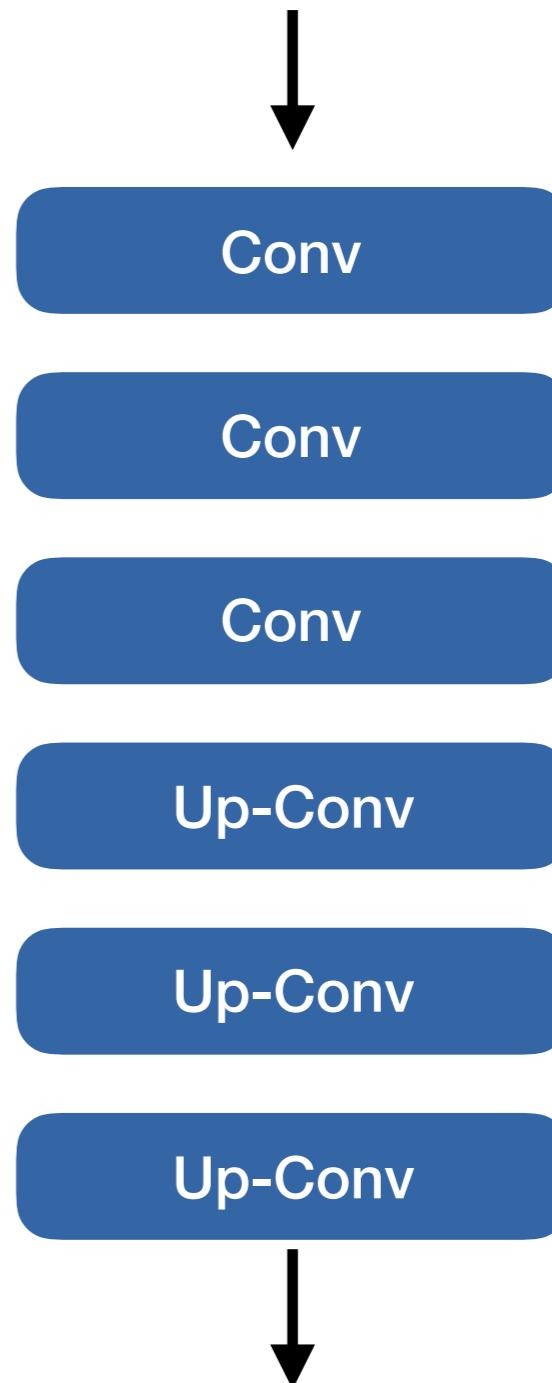
Up-convolution in action

- Used closer to output layers



Up-convolutions and skip connections

- Provides lower-level high-resolution features to output



The many names of up-convolution

- Transpose convolution
- “Deconvolution”
- fractionally strided convolutions

Summary

© 2019 Philipp Krähenbühl and Chao-Yuan Wu

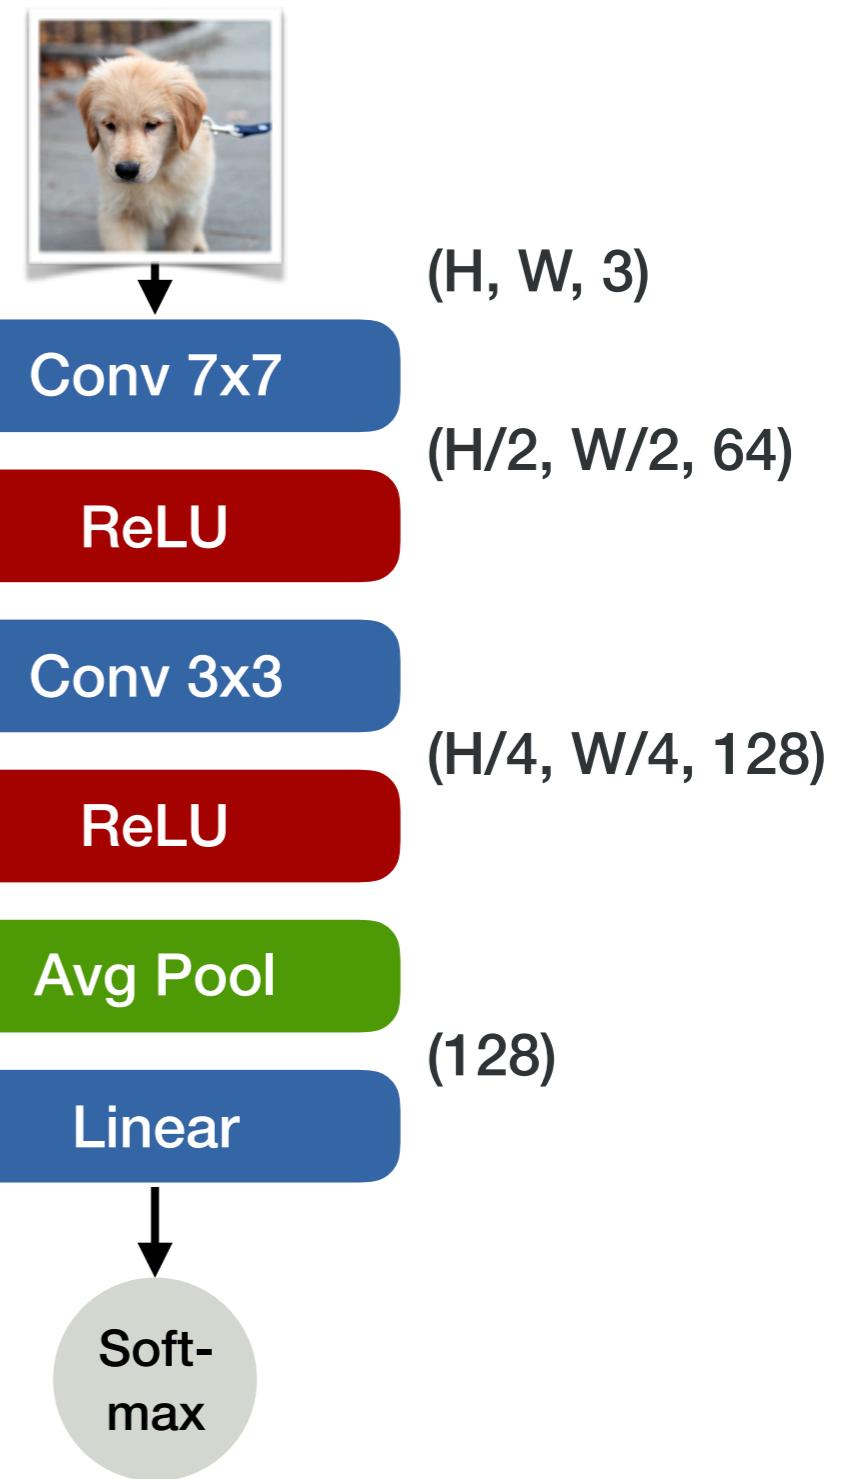
Convolutional networks

- Linear transformations
 - Convolutions
- Convolutional non-linearities
 - Pooling



ConvNet Design

- Stride and increase channels
- Small kernel size
- All convolutional
- Up-convolution close to output (optional)



Applications of ConvNets

- Autonomous vehicles
- Analyze medical images
- Geoscience: Analyze scans of rock formations to find oil

