

January 23, 2024

```
[1]: %pylab inline
import torch
import sys
sys.path.append('.')
sys.path.append('../..')
from data import load
device = torch.device('cuda') if torch.cuda.is_available() else torch.
    ↪device('cpu')
print('device = ', device)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
 Populating the interactive namespace from numpy and matplotlib
 device = cuda

```
[2]: class ConvNet(torch.nn.Module):
    class Block(torch.nn.Module):
        def __init__(self, n_input, n_output, stride=1):
            super().__init__()
            self.net = torch.nn.Sequential(
                torch.nn.Conv2d(n_input, n_output, kernel_size=3, padding=1, ↪
            ↪stride=stride, bias=False),
                torch.nn.BatchNorm2d(n_output),
                torch.nn.ReLU(),
                torch.nn.Conv2d(n_output, n_output, kernel_size=3, padding=1, ↪
            ↪bias=False),
                torch.nn.BatchNorm2d(n_output),
                torch.nn.ReLU()
            )
            self.downsample = None
            if stride != 1 or n_input != n_output:
                self.downsample = torch.nn.Sequential(torch.nn.Conv2d(n_input, ↪
            ↪n_output, 1, stride=stride),
                                                         torch.nn.
            ↪BatchNorm2d(n_output))

        def forward(self, x):
            identity = x
```

```

        if self.downsample is not None:
            identity = self.downsample(x)
            return self.net(x) + identity

    def __init__(self, layers=[32,64,128], n_input_channels=3):
        super().__init__()
        L = [torch.nn.Conv2d(n_input_channels, 32, kernel_size=7, padding=3,
↪stride=2, bias=False),
            torch.nn.BatchNorm2d(32),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
        ]
        c = 32
        for l in layers:
            L.append(self.Block(c, l, stride=2))
            c = l
        self.network = torch.nn.Sequential(*L)
        self.classifier = torch.nn.Linear(c, 1)

    def forward(self, x):
        # Compute the features
        z = self.network(x)
        # Global average pooling
        z = z.mean(dim=[2,3])
        # Classify
        return self.classifier(z)[:0]

model = ConvNet()
model.train()
print( model.training )
model.eval()
print( model.training )

```

True
False

```

[4]: n_epochs = 100

train_data = load.get_dogs_and_cats(resize=(128,128), batch_size=32,
↪is_resnet=True)
optimizer = torch.optim.Adam(model.parameters())
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=30, gamma=0.1)
scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, [50,75], gamma=0.1)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'max')

# Transfer the data to a GPU (optional)
model = model.to(device)

```

```

# Construct the loss and accuracy functions
loss = torch.nn.BCEWithLogitsLoss()
accuracy = lambda o, l: ((o > 0).long() == l.long()).float()

# Train the network
for epoch in range(n_epochs):
    model.train()
    accuracies = []
    for it, (data, label) in enumerate(train_data):
        # Transfer the data to a GPU (optional)
        data, label = data.to(device), label.to(device)

        # Produce the output
        o = model(data)

        # Compute the loss and accuracy
        loss_val = loss(o, label.float())
        accuracies.extend(accuracy(o, label).detach().cpu().numpy())

        # Take a gradient step
        optimizer.zero_grad()
        loss_val.backward()
        optimizer.step()
        break
    # scheduler.step()
    scheduler.step(np.mean(accuracies))
    print( 'epoch = ', epoch, 'optimizer_lr', optimizer.param_groups[0]['lr'], 'accuracy', np.mean(accuracies))

```

```

epoch = 0 optimizer_lr 0.001 accuracy 0.46875
epoch = 1 optimizer_lr 0.001 accuracy 0.65625
epoch = 2 optimizer_lr 0.001 accuracy 0.96875
epoch = 3 optimizer_lr 0.001 accuracy 0.96875
epoch = 4 optimizer_lr 0.001 accuracy 1.0
epoch = 5 optimizer_lr 0.001 accuracy 1.0
epoch = 6 optimizer_lr 0.001 accuracy 1.0
epoch = 7 optimizer_lr 0.001 accuracy 1.0
epoch = 8 optimizer_lr 0.001 accuracy 1.0
epoch = 9 optimizer_lr 0.001 accuracy 1.0
epoch = 10 optimizer_lr 0.001 accuracy 1.0
epoch = 11 optimizer_lr 0.001 accuracy 1.0
epoch = 12 optimizer_lr 0.001 accuracy 1.0
epoch = 13 optimizer_lr 0.001 accuracy 1.0
epoch = 14 optimizer_lr 0.001 accuracy 1.0
epoch = 15 optimizer_lr 0.0001 accuracy 1.0
epoch = 16 optimizer_lr 0.0001 accuracy 1.0

```


[]: